

AI-Based Paddy Seed Quality Assessment And Farmer Feedback System

*A thesis submitted in partial fulfilment of the requirements
for the award of the degree of*

Bachelor of Technology

By

Joymangol Chingangbam (2201CS0004)

Nilambar Elangbam (2201CS0005)

Khangembam Yaiphaba Singh (2201CS0102)

Dayananda Laishram (2201CS0113)

Under the Guidance of

Dr. Roshni Rajkumari

Assistant Professor



**Department of
Computer Science & Engineering**

Manipur Technical University

May, 2026

DEDICATION

This thesis is dedicated to all those who have inspired, supported and stood by us throughout the journey of this project.

To our families: Your unwavering love, patience, and encouragement have been the bedrock of our emotional strength during the most challenging phases. *Your belief in our potential has fortified our resilience and determination.*

To our mentors and instructors: Your expertise, thoughtful guidance, and commitment to our academic growth have been instrumental in shaping both this project and our understanding of the field. *Your encouragement and constructive feedback have been invaluable at each stage.*

To our friends and peers: Your camaraderie, moral support, and motivation have illuminated our path. *Your presence transformed moments of stress into opportunities for shared growth and learning.*

To the global community of researchers, engineers, and innovators in Artificial Intelligence and Computer Vision: Your pioneering work has laid the foundation for this thesis. *Your dedication to advancing knowledge continues to inspire and propel us to contribute meaningfully to this dynamic field.*

This thesis stands as a testament to the collective support, shared vision, and unwavering encouragement of all who have been part of our academic and personal journey.

AI-Based Paddy Seed Quality Assessment And Farmer Feedback System

*A thesis submitted in partial fulfilment of the requirements
for the award of the degree of*

Bachelor of Technology

By

Joymangol Chingangbam (2201CS0004)

Nilambar Elangbam (2201CS0005)

Khangembam Yaiphaba Singh (2201CS0102)

Dayananda Laishram (2201CS0113)

Under the Guidance of

Dr. Roshni Rajkumari

Assistant Professor



**Department of
Computer Science & Engineering**

Manipur Technical University

May, 2026



MANIPUR TECHNICAL UNIVERSITY

(A University established under the **Manipur Technical University Act, 2016**)

Imphal, Manipur – 795004

BONAFIDE CERTIFICATE

This is to certify that the thesis entitled “**AI-Based Paddy Seed Quality Assessment and Farmer Feedback System**” submitted by:

Joymangol Chingangbam	(2201CS0004)
Nilambar Elangbam	(2201CS0005)
Khangembam Yaiphaba Singh	(2201CS0102)
Dayananda Laishram	(2201CS0113)

to the **Department of Computer Science & Engineering**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree in **Computer Science and Engineering**, is a genuine and original work carried out by them under my supervision during the academic year **2025–2026**.

The project embodies the results of their own research and development work and has not been submitted, either in part or in full, for the award of any other degree or diploma of this university or any other institution. The work has been carried out with sincerity, academic integrity, and adherence to research ethics.

We consider this work worthy of consideration for the partial fulfillment of the requirements for the said degree.

Supervisor

Dr. Roshni Rajkumari
Assistant Professor,
Computer Science & Engineering

Head of Department

Mrs. Tayenjam Aarena
Assistant Professor & Head,
Computer Science & Engineering



MANIPUR TECHNICAL UNIVERSITY

(A University established under the **Manipur Technical University Act, 2016**)

Imphal, Manipur – 795004

DECLARATION

I hereby declare that the thesis entitled “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**”, submitted to the **Department of Computer Science and Engineering, Manipur Technical University**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree, is a result of my own work and effort. Under the guidance of **Dr. Roshni Rajkumari**, Assistant Professor, Department of Computer Science and Engineering, Manipur Technical University.

I affirm that this work has been carried out with academic honesty and integrity. All sources of information have been properly cited and acknowledged. This thesis has not been submitted earlier, either in part or in full, for the award of any degree or diploma at this or any other institution.

Date:

Place:

Joymangol Chingangbam

Reg. No. - 2201CS0004

Computer Science & Engineering



MANIPUR TECHNICAL UNIVERSITY

(A University established under the **Manipur Technical University Act, 2016**)

Imphal, Manipur – 795004

DECLARATION

I hereby declare that the thesis entitled “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**”, submitted to the **Department of Computer Science and Engineering, Manipur Technical University**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree, is a result of my own work and effort. Under the guidance of **Dr. Roshni Rajkumari**, Assistant Professor, Department of Computer Science and Engineering, Manipur Technical University.

I affirm that this work has been carried out with academic honesty and integrity. All sources of information have been properly cited and acknowledged. This thesis has not been submitted earlier, either in part or in full, for the award of any degree or diploma at this or any other institution.

Date:

Place:

Nilambar Elangbam

Reg. No. - 2201CS0005

Computer Science & Engineering



MANIPUR TECHNICAL UNIVERSITY

(A University established under the **Manipur Technical University Act, 2016**)

Imphal, Manipur – 795004

DECLARATION

I hereby declare that the thesis entitled “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**”, submitted to the **Department of Computer Science and Engineering, Manipur Technical University**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree, is a result of my own work and effort. Under the guidance of **Dr. Roshni Rajkumari**, Assistant Professor, Department of Computer Science and Engineering, Manipur Technical University.

I affirm that this work has been carried out with academic honesty and integrity. All sources of information have been properly cited and acknowledged. This thesis has not been submitted earlier, either in part or in full, for the award of any degree or diploma at this or any other institution.

Date:

Place:

Khangembam Yaiphaba Singh

Reg. No. - 2201CS0102

Computer Science & Engineering



MANIPUR TECHNICAL UNIVERSITY

(A University established under the **Manipur Technical University Act, 2016**)

Imphal, Manipur – 795004

DECLARATION

I hereby declare that the thesis entitled “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**”, submitted to the **Department of Computer Science and Engineering, Manipur Technical University**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree, is a result of my own work and effort. Under the guidance of **Dr. Roshni Rajkumari**, Assistant Professor, Department of Computer Science and Engineering, Manipur Technical University.

I affirm that this work has been carried out with academic honesty and integrity. All sources of information have been properly cited and acknowledged. This thesis has not been submitted earlier, either in part or in full, for the award of any degree or diploma at this or any other institution.

Date:

Place:

Dayananda Laishram

Reg. No. - 2201CS0113

Computer Science & Engineering



MANIPUR TECHNICAL UNIVERSITY

(A University established under the **Manipur Technical University Act, 2016**)

Imphal, Manipur – 795004

ACKNOWLEDGEMENT

We sincerely acknowledge the collective support and guidance received throughout the successful completion of our thesis project “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**”.

We take this opportunity to thank our Vice Chancellor, **Prof. W. Chandbabu Singh**, whose aid and encouragement provided invaluable support and inspiration throughout our academic journey. We are deeply grateful to the **Department of Computer Science and Engineering**, Manipur Technical University, for providing the academic environment, infrastructure and encouragement that fostered innovation and technical exploration.

Our heartfelt thanks go to our project supervisor, **Dr. Roshni Rajkumari**, Assistant Professor, whose consistent guidance, insightful feedback and unwavering support played a pivotal role in shaping this project. We also extend our sincere appreciation to **Mrs. Tayenjam Aerena**, Head of Department, for her inspiring leadership and motivation.

We are also thankful to all the faculty members, technical staff and friends who provided their valuable input and moral support throughout the journey.

Date:

Place:

Sincerely,

Joymangol Chingangbam

Nilambar Elangbam

Dayananda Laishram

Khangembam Yaiphaba Singh

ABSTRACT

The quality of paddy seed plays a vital role in crop yield, germination rate and overall agricultural productivity. Many small and marginal farmers suffer crop loss due to substandard or counterfeit seeds that closely resembles certified products. Due to lack of affordable and rapid verification methods, farmers rely on packaging labels and basic manual inspection, which is time consuming, inconsistent and error prone. Poor quality seeds are generally known after the seeds are sown, and thus, results in low germination, reduce yield and incur a financial loss.

To address these challenges, our project entitled "**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**" is designed to combat the spread of counterfeit and substandard paddy seeds. This project is a web-based platform that utilizes artificial intelligence and deep learning model to automate the classification of quality seeds into healthy and suspicious ones based on characteristics like size, shape uniformity, colour, surface texture and visible defects. Farmers can scan seed samples using a simple image-based interface to receive an estimated quality assessment of paddy seeds, presented in clear, farmer friendly language.

In addition to seed quality assessment, our system also includes a Farmer Complaint and Feedback module that enables reporting of field level issues such as poor germination or crop failure, cultivation experience, seed performance, etc.

The system avoids chemical analysis, specialized hardware or complex supply chain systems, making it low-cost, scalable and suitable for rural deployment. It also helps farmers select better quality seeds, reduce crop failure, and improve productivity. The integration of artificial intelligence with farmer feedback system makes the proposed system a useful tool for modern agricultural practices and sustainable farming development.

Table of Contents

BONAFIDE CERTIFICATE	i
DECLARATION	ii
ACKNOWLEDGEMENT	vi
ABSTRACT	vii
List of Tables	xiv
List of Figures	xvi
Symbols and Acronyms	1
1 Introduction	1
1.1 Background and Context	1
1.2 Motivation	1
1.3 Problem Statement	2
1.4 Objectives	2
1.5 Scope of the Project	3
2 Literature Review	4
2.1 Seed Quality Assessment: Traditional Methods	4
2.2 Deep Learning for Agricultural Image Classification	4
2.3 Transfer Learning and ResNet Architectures	5
2.4 Farmer Feedback Systems	6

2.5	Web-Based Agricultural Platforms	6
2.6	Review of Related Work	7
2.7	Summary and Research Gaps	7
3	System Design and Architecture	9
3.1	Overall System Architecture and Technology Stack	9
3.2	Functional and Non-Functional Requirements	11
3.3	High-Level System Architecture	11
3.4	System Data Flow Pipeline	12
3.5	Database Entity-Relationship Model	14
4	Deep Learning Model for Paddy Seed Quality Assessment	17
4.1	Dataset Description	17
4.1.1	Data Sources and Collection (Kaggle)	17
4.1.1.1	Public Source	17
4.1.1.2	Custom Source	18
4.1.2	Dataset Statistics and Class Distribution	18
4.1.3	Train/Validation/Test Split Strategy	19
4.1.4	Data Preprocessing and Augmentation	20
4.2	Model Architecture	21
4.2.1	ResNet50 Backbone (25.6M Parameters)	21
4.2.2	Custom Classification Head	21
4.2.3	Anti-Overfitting Techniques	21
4.3	DUS Parameter Integration	22
4.3.1	Feature Extraction—17 Physical Parameters	22
4.3.2	Pixel-to-Physical Unit Conversion	23
4.3.3	Variety Matching (12 Reference Varieties)	23
4.4	Training Pipeline	24
4.4.1	Phase 1—Classification Head Training (Epochs 1–8)	24
4.4.2	Phase 2—Full Fine-Tuning (Epochs 9–15)	24

4.4.3	Hyperparameter Configuration and Optimiser	24
4.4.4	Model Export Formats	25
4.5	Model Comparison and Selection	26
4.6	Final Model Performance Summary	26
5	Model Serving API	28
5.1	API Architecture and Request Orchestration	28
5.1.1	FastAPI and ASGI Serving Stack	28
5.1.2	End-to-End Model Request Flow	29
5.2	API Specification and Endpoint Contracts	30
5.3	Containerisation and Deployment Architecture	31
5.3.1	CI/CD Deployment Pipeline	32
5.3.2	Live Production API Endpoint	32
6	Backend System Architecture and Implementation	33
6.1	Architectural Design	33
6.2	Technology Stack	34
6.3	Project Structure	35
6.4	Domain Layer	36
6.4.1	Core Entities	36
6.4.2	Interfaces	37
6.5	Application Layer	37
6.5.1	Application Services	37
6.5.2	DTOs	38
6.5.3	Validation	38
6.6	Infrastructure Layer	38
6.6.1	Database Management	39
6.6.2	Repository Pattern	39
6.6.3	AI and Cloud Integrations	39
6.7	Database Schema	40

6.8	Presentation Layer	41
6.8.1	API Controllers	41
6.8.2	Middleware	41
6.9	Authentication and Security	42
6.10	Verification Workflow	43
6.11	Configuration and Deployment	44
6.12	Summary	44
7	Frontend System Architecture and Implementation	46
7.1	Frontend Architecture and Design Patterns	46
7.1.1	Architectural Layers	46
7.1.2	Secure Proxy Architecture	47
7.2	Technology Stack Analysis	47
7.3	Modular Project Structure	48
7.4	User Role-Based Access Control (RBAC)	48
7.4.1	Role Partitioning	48
7.4.2	Route Protection Middleware	49
7.5	Authentication and Secure Persistence	49
7.5.1	Token Storage Architecture	49
7.5.2	Automated Token Rotation	49
7.6	API Integration and Type Safety	49
7.7	State Management	49
7.7.1	Asynchronous Server State (TanStack Query)	50
7.7.2	Synchronous Client State (Zustand)	50
7.8	Responsive UI Design and Field Usability	50
7.8.1	Standardized UI Library	50
7.8.2	Accessibility Enhancements	51
8	System Integration	52
8.1	System Integration Matrix	52

8.2	End-to-End Architectural Workflow	52
9	Results and Evaluation	55
9.1	Model Performance	55
9.1.1	Accuracy, Precision, Recall, and F1 Score	55
9.2	Seed Quality Results: Pure vs. Impure Seeds	57
9.3	Confusion Matrix Analysis	58
9.4	Regularization Effectiveness	59
9.5	Model Comparison and Architecture Selection	60
9.6	System Performance and Production Latency	61
9.6.1	Inference Latency Evaluation	61
9.6.2	End-to-End System Performance Under Load	62
9.7	Functional Verification and Test Execution	62
9.7.1	Seed Classification Subsystem Validation	63
9.7.2	Grievance Management Subsystem Validation	63
9.7.3	Role-Based Access Control and Security Verification	63
9.8	System User Interfaces and Operational Workflows	63
9.8.1	Public Portal, Onboarding, and Multilingual Support	65
9.8.2	Farmer Portal and AI Seed Quality Assessment Workflow	66
9.8.3	Grievance Redressal and Officer Investigation Interfaces	71
9.8.4	Platform Administration, User Governance, and Telemetry	73
9.9	Challenges Encountered and Mitigation Strategies	74
9.10	Summary and System Limitations	76
10	Conclusion And Future Work	77
10.1	Summary of Work	77
10.2	Key Contributions and Findings	78
10.2.1	AI-Based Automated Paddy Seed Classification	78
10.2.2	Integration of Multiple AI Services	78
10.2.3	DUS Parameter Evaluation	78

10.2.4	Real-Time Farmer Assistance	79
10.2.5	Secure and Scalable System Architecture	79
10.2.6	Complaint Redressal and Governance Workflow	79
10.2.7	Reduction of Manual Dependency	79
10.2.8	Cloud-Native AI Deployment	79
10.3	Future Enhancements	80
10.3.1	Offline-First Progressive Web Application (PWA)	80
10.3.2	Enhanced Security Mechanisms	80
10.3.3	Blockchain-Based Verification Ledger	80
10.3.4	IoT and Smart Agriculture Integration	80
10.3.5	Expansion to Multiple Crop Varieties	80
10.3.6	Mobile Application Development	81
10.3.7	Multilingual Farmer Assistance	81
10.3.8	Advanced Analytics Dashboard	81
10.3.9	Federated Learning for Privacy-Preserving AI	81
10.3.10	Explainable AI (XAI)	81
10.4	Conclusion	81

References	83
-------------------	-----------

List of Tables

3.1	Backend Architectural Layers and Responsibilities	9
3.2	Consolidated Technology Stack of AgriVerify	10
3.3	System Functional Requirements	11
3.4	System Non-Functional Requirements	12
4.1	Dataset Source Summary	18
4.2	Class-wise Distribution of the Combined Dataset	19
4.3	Dataset Split Statistics	19
4.4	Hyperparameter Configuration and Training Settings	25
4.5	Model Export Formats	25
4.6	Architecture Comparison on Test Set	26
4.7	Final Test Set Performance—ResNet50	27
4.8	Training Behaviour Summary	27
5.1	API Specification for the POST /predict Model Inference Endpoint	30
5.2	API Specification for the GET /health Liveness Endpoint	31
6.1	Backend Technology Stack	35
6.2	Backend Project Structure	36
6.3	Important DTOs	38
6.4	External Service Integrations	39
6.5	API Controllers	41
7.1	Frontend Technology Stack	48
7.2	Standardized UI Components	50

8.1	How Different Parts of AgriVerify Communicate	53
9.1	Overall Model Evaluation Results	56
9.2	Performance on Each Seed Category	57
9.3	Detailed Look at Every Prediction	58
9.4	How Our Anti-Overfitting Techniques Helped	59
9.5	Comparing Four Different Neural Network Architectures	60
9.6	Speed on Different Hardware	62
9.7	Real-World System Performance Under Load	62
9.8	Seed Verification Test Cases	63
9.9	Complaint System Test Cases	64
9.10	Security and Access Control Test Cases	64

List of Figures

3.1	High-level system architecture of the AgriVerify platform.	13
3.2	Data flow diagram for the seed verification pipeline	15
3.3	Simplified entity-relationship diagram for the AgriVerify domain model. Cardinality notation follows UML conventions.	16
5.1	Asynchronous model request processing flow.	30
5.2	Automated CI/CD pipeline and Serverless GCP Deployment Architecture.	31
6.1	Backend Clean Architecture Layers	34
6.2	Entity Relationship Diagram – FakeSeedDetection System	40
6.3	JWT Authentication Flow	42
6.4	Seed Verification Workflow	43
7.1	Frontend Layered Architecture and Request Execution Flow	47
8.1	How AgriVerify processes seed verification requests and automatically deploys updates. Top section shows what happens when a farmer submits images; bottom section shows how we keep the system updated.	53
9.1	How the model improved over 15 training cycles. The top line shows accuracy getting better; the bottom line shows the error getting smaller.	56
9.2	Visual representation of how often the model got predictions right or wrong.	58

9.3	Visual comparison of how each model performed and how long training took.	61
9.4	Platform landing page displaying value proposition and core features.	65
9.5	Multilingual selection interface supporting regional languages. . . .	65
9.6	Role selection interface for login.	66
9.7	Farmer registration form collecting demographic and location information.	67
9.8	Secure admin login requiring multi-factor authentication.	67
9.9	Farmer dashboard interface displaying verification statistics and history.	68
9.10	Farmer profile configuration interface with verification achievements.	68
9.11	User profile modification interface.	69
9.12	Mobile camera interface for uploading seed packaging images. . . .	69
9.13	Verification result indicating pure seed batch with morphological measurements.	70
9.14	Verification alert for counterfeit or low-quality seed batch.	70
9.15	Grievance registration interface for lodging seed quality complaints.	71
9.16	Farmer tracking interface for monitoring grievance status.	71
9.17	Agricultural officer dashboard presenting regional complaint metrics.	72
9.18	Filtered and prioritized grievance investigation queue for officers. . .	72
9.19	Officer grievance detail pane featuring user evidence and resolution controls.	73
9.20	Administrator real-time system performance telemetry dashboard. . .	73
9.21	User administration and access control interface.	74
9.22	Analytical reporting dashboard showing quality trends and system statistics.	75
9.23	District-wise geospatial hot-spot mapping of counterfeit occurrences.	75

Chapter 1

Introduction

1.1 Background and Context

Agriculture is the foundational backbone of many global economies, heavily influencing food security, employment and national income. At the core of all agricultural productivity depends on the quality of the seed planted, which is the most critical and non-negotiable factor determining crop yield, resilience to disease and overall harvest quality. However, the agricultural sector is currently facing a severe crisis due to the increasing problem of counterfeit or fake seeds in the market.

Counterfeit seeds are often packaged to resemble premium brands or are visually indistinguishable from high-yield varieties, yet they exhibit drastically lower germination rates and poor genetic traits. Traditional methods of verifying seed authenticity rely on destructive laboratory testing or subjective physical inspection, both of which are inaccessible, expensive and far too slow for the an average farmer. Consequently, farmers often discover that they have been deceived only after planting, leading to catastrophic losses in time, resources and income.

1.2 Motivation

The primary motivation for this project stems from the devastating socio-economic impact that low-quality seeds have on farming communities. When farmers unknowingly purchase and plant low-quality seeds, they face severe financial distress, which cascades into broader issues of local food insecurity and loss of trust in agricultural supply chains. There is an urgent, unmet need for accessible technological tools that empower grassroots agricultural workers to protect

themselves against fraud.

Simultaneously, the rapid advancement of deep learning, computer vision, and cloud computing presents an unprecedented opportunity to address this crisis. By leveraging deep learning and computer vision techniques it is now possible to achieve automated, laboratory-grade visual analysis of seeds. The motivation is to harness these cutting-edge AI technologies and package them into a simple, web-based tool, putting the power of automated, real-time seed verification directly into the hands of farmers.

1.3 Problem Statement

Despite the devastating effects of agricultural fraud, farmers currently lack a reliable, accessible, and rapid method to accurately distinguish between genuine and counterfeit seeds before the planting season begins. Visual similarities make manual detection nearly impossible and existing institutional supply chains lack a transparent, real-time verification system.

The core problem is the absence of an integrated, technology-driven solution that can automatically analyze seed imagery, identify fraudulent anomalies and capture feedback from farmers to map and identify the spread of counterfeit seeds in real-time.

1.4 Objectives

The main goal of this project is to create an intelligent ecosystem for seed verification. To achieve this, the project outlines the following specific objectives:

1. To design and train a deep learning model capable of accurately classifying seed images and identifying anomalies associated with counterfeit seeds.
2. To develop and deploy a user-friendly web-based platform that allows farmers to upload images of their seeds for AI inference and verification.
3. To integrate a "Farmer Feedback System" within the platform, enabling users to report their post-purchase and post-planting experiences (such as germination failure rates).
4. To utilize the feedback and newly uploaded images to continuously refine the database, creating an early warning system against emerging counterfeit supply chains.

1.5 Scope of the Project

The scope of this project encompasses the end-to-end development of an AI-driven web application tailored for agricultural quality control. This includes the collection, curation, and preprocessing of a targeted seed image dataset, followed by the training, validation, and optimization of a Convolutional Neural Network (CNN).

Additionally, the scope covers the full-stack development of the web platform, including an intuitive frontend interface for image uploads and a robust backend to handle model inference and database management. The project is strictly limited to the non-destructive, visual assessment of seeds via digital imagery; it does not extend to biochemical, genetic, or DNA-based testing methods. The final deliverable will be a functional prototype that demonstrates the feasibility of combining deep learning with crowdsourced feedback to combat agricultural fraud.

Chapter 2

Literature Review

2.1 Seed Quality Assessment: Traditional Methods

Conventional seed quality assessment rests on two broad approaches: manual visual inspection and laboratory-based biochemical testing. In field practice, inspectors evaluate seeds by examining size, shape, colour and surface texture; judgements that are inherently subjective, vary between observers and resist standardisation across districts or seasons. Laboratory tests such as the Tetrazolium (TZ) test and standard germination assays offer more reliable viability estimates, but both are destructive and typically require four to seven days under controlled conditions. That turnaround makes them impractical for screening large consignments at intake points, or for farmers who need a decision before the planting window closes.

Near-infrared (NIR) spectroscopy and X-ray transmission imaging have been investigated as non-destructive alternatives. NIR detects internal moisture gradients and biochemical composition without damaging the sample; X-ray imaging can reveal embryo defects invisible on the seed surface. Both perform well in laboratory conditions, but equipment costs typically USD 15,000-80,000 per unit and the requirement for trained operators confine their use to national seed-testing stations. Smallholder farmers and district-level inspectors in regions such as northeast India have no practical access to either technology, which is precisely where the need for affordable, rapid screening is most acute.

2.2 Deep Learning for Agricultural Image Classification

Convolutional Neural Networks (CNNs) have been applied to a range of agricultural classification problems over the past decade, including leaf disease iden-

tification [1], weed recognition [2], and post-harvest quality grading [3]. Their principal advantage over classical machine-learning pipelines is that spatial features are learned directly from raw pixel data during training, removing the need to manually engineer feature extractors for each new crop or defect type.

Applying CNNs to seed classification follows naturally from this body of work. Paddy seeds exhibit subtle differences in surface texture and colour between certified and counterfeit lots; differences that resist quantification through handcrafted descriptors but emerge reliably from sufficiently large training sets. The computational cost of running a trained CNN at inference time is also modest enough to support server-side web deployment, which matters when target users have intermittent electricity but reasonable mobile-data coverage.

2.3 Transfer Learning and ResNet Architectures

Training a CNN from random initialisation requires on the order of tens of thousands of labelled examples per class to achieve stable convergence. Annotated seed image datasets of that scale do not yet exist for most crop varieties, and certainly not for the specific task of detecting counterfeits. Transfer learning is the standard response: a model pre-trained on ImageNet (1.28 million labelled images across 1,000 categories) is fine-tuned on the smaller domain-specific dataset. The early convolutional layers, which encode low-level features such as edges and textures, are retained; only the later task-specific layers are updated. In practice, this approach consistently reaches accuracy close to what a from-scratch model would achieve on a dataset roughly ten times larger.

ResNet architectures are well-suited to this fine-tuning scenario. The residual skip connections introduced by He et al. [4] allow gradients to propagate through very deep networks without vanishing, making it feasible to fine-tune models with 50 or more layers on datasets that would cause shallower architectures to overfit. ResNet-50 has become a common baseline in agricultural image classification, balancing parameter count (≈ 25 million) against accuracy on standard benchmarks, with pre-trained weights publicly available. For this project, it provides a well-validated starting point that reduces both training time and the risk of overfitting on the available paddy seed dataset.

2.4 Farmer Feedback Systems

A detection model trained on a fixed dataset will degrade over time if counterfeiting practices evolve and no new training data are collected. Field-level feedback from farmers is one practical mechanism for capturing these shifts. When a farmer submits a seed batch for image-based verification and subsequently reports germination failure or abnormal yield, that report constitutes a labelled example, a model prediction paired with a ground-truth outcome, that can be incorporated into periodic retraining cycles.

Feedback also serves a more immediate diagnostic function. Aggregating reports across users can flag a suspect lot or supplier before the model itself is updated: three independent failure reports from the same district within a fortnight is a meaningful signal, regardless of what the image classifier predicts. This rule-based alerting complements the model rather than replacing it. Sustaining farmer participation remains a practical difficulty; deployments in low-resource contexts consistently find that single-tap rating interfaces produce far higher response rates than free-text fields or structured forms [5].

2.5 Web-Based Agricultural Platforms

Deploying a trained model as a web service separates the computational workload from the end user’s device. A farmer uploads a seed image through a browser; the server runs inference and returns a result. This architecture offers two concrete advantages over distributing a standalone mobile application: the model can be updated centrally without requiring client-side installations, and every prediction is logged in a structured format that supports monitoring and retraining.

For farmers in Manipur and comparable regions, browser-based access via a mid-range Android handset is more realistic than assuming a dedicated application installation. Keeping the frontend lightweight, avoiding heavy JavaScript bundles and large media assets reduces page-load times on 4G connections with variable signal quality. Containerised inference endpoints handle computationally intensive tasks without requiring a permanently running server, reducing operating costs during off-season periods when query volume is low.

2.6 Review of Related Work

Early automated seed classification employed Support Vector Machines (SVMs) and Random Forests operating on handcrafted features derived from shape measurements, colour histograms, and texture descriptors. Ni et al. [6] classified rice seed varieties using morphological parameters extracted from binary silhouettes, achieving 94% accuracy across five varieties under controlled lighting. These methods perform adequately when intra-class variation is low and imaging conditions are uniform, but generalise poorly to photographs taken in field settings.

Deep learning approaches have largely displaced handcrafted methods in more recent work. Liu et al. [7] applied VGG16 transfer learning to maize seed defect detection, reporting 97.3% accuracy on a 12-class dataset. Xu et al. [8] used ResNet-50 for soybean quality grading and found that fine-tuning from ImageNet weights outperformed training from scratch by approximately 8 percentage points when training data were limited to fewer than 2,000 images per class. Both studies, however, address naturally occurring defects, such as cracking, mould and insect damage, rather than intentional substitution. A counterfeit seed is engineered to resemble the genuine article; distinguishing the two is a harder classification problem than separating a damaged kernel from an intact one. The visual gap between classes is narrower, and the adversarial nature of counterfeiting means the distribution of counterfeits may shift as producers adapt. No published study directly evaluates CNN performance on verified counterfeit-versus-genuine seed pairs, and none integrates a structured feedback collection mechanism into the evaluation pipeline.

2.7 Summary and Research Gaps

The reviewed literature establishes that transfer-learned CNNs are effective for agricultural image classification, including seed quality assessment. Three gaps are directly relevant to this work. First, no published study has targeted intentional counterfeiting as distinct from natural quality defects; the two problems differ in that counterfeits may pass casual visual inspection and require the model to detect subtle discriminative cues. Second, existing models are evaluated on laboratory-quality images collected under controlled conditions; performance on photographs taken by farmers using mobile devices in variable lighting has not been reported. Third, no end-to-end system has been published that connects model inference to a structured feedback collection mechanism and uses that feedback iteratively for retraining.

This project addresses all three gaps. The classifier is trained and evaluated specifically on a fake-versus-genuine paddy seed dataset; the web platform is tested with mobile-captured images; and the farmer feedback module is built directly into the application, with logged responses stored in a format suitable for future retraining rounds.

Chapter 3

System Design and Architecture

3.1 Overall System Architecture and Technology Stack

The AgriVerify platform is designed using a multi-tiered architecture based on the Clean Architecture model to ensure strict separation of concerns, decoupling of business logic from external frameworks and high maintainability [9]. The system comprises three primary layers: a Next.js frontend client web application, an ASP.NET Core backend Web API and a FastAPI deep learning inference microservice. These tiers interact to process seed validation requests, manage user authentication and persist data.

The backend API follows Onion Architecture principles, dividing responsibilities into four distinct concentric layers as detailed in Table 3.1.

Table 3.1: Backend Architectural Layers and Responsibilities

Architectural Layer	Core Responsibilities and Encapsulated Components
Domain Layer	Core domain entities (User, VerificationHistory), value objects, and domain exceptions. Completely independent of external frameworks.
Application Layer	Orchestration logic (VerificationService, UserService), Data Transfer Objects (DTOs), mapping configurations and request validation rules.
Infrastructure Layer	External integration implementations, including database persistence via Entity Framework Core, Azure Cognitive Services and Blob storage managers.
API Layer	REST API controllers, request-response handling, JWT authentication middleware and unified exception filters.

Table 3.2 details the software development kits (SDKs), frameworks, libraries, and cloud platforms employed across the frontend, backend, and machine learning components.

Table 3.2: Consolidated Technology Stack of AgriVerify

Component	Framework / Library	Functionality / Purpose
Frontend	Next.js 14 (TS)	Client interface, server-side rendering, and API proxy routing [10].
	TanStack Query	Server state synchronization and asynchronous data caching [11].
	Zustand	Client-side authentication state management [12].
	Tailwind CSS	Styling system and component customization [13].
Backend	ASP.NET Core 8	Orchestration Web API, routing, and controller layer [14].
	Entity Framework Core	Object-Relational Mapping (ORM) for SQL Server database access.
	Azure Cognitive Services	Optical Character Recognition (OCR) and Custom Vision packaging verification.
	Azure OpenAI Service	Natural language synthesis for seed verification reports.
ML	PyTorch 2.1	Deep learning framework for neural network training [15].
	ResNet-50	Pre-trained CNN backbone for seed variety classification [4].
	OpenCV	Preprocessing and morphometric DUS parameter analysis [16].
	FastAPI	Serverless microservice API on GCP Cloud Run [17].

The PyTorch ResNet-50 model is containerized and deployed as a FastAPI microservice on Google Cloud Platform (GCP) Cloud Run to enable automatic, serverless scaling [17]. During inference, input images undergo size standardization to 224×224 pixels and normalization using ImageNet statistics ($\mu = [0.485, 0.456, 0.406]$, $\sigma = [0.229, 0.224, 0.225]$). OpenCV extracts 17 Distinctness, Uniformity, and Stability (DUS) morphometric parameters, which are compared against reference variety profiles to compute similarity scores [18].

3.2 Functional and Non-Functional Requirements

System requirements are categorized into functional specifications, defining the core operations and roles within the application, and non-functional specifications, outlining quality attributes and system constraints. These requirements ensure that the platform meets both agricultural user needs and architectural standards. The functional requirements (Table 3.3) define the core capabilities across different system actors (farmers, agricultural officers, and administrators). Concurrently, the non-functional requirements (Table 3.4) dictate operational performance, security, and portability parameters, ensuring standard enterprise scalability and reliability.

Table 3.3: System Functional Requirements

ID	Module / Role	Functional Specification
FR-01	Authentication	User registration and authentication via JWT for farmers, officers, and administrators.
FR-02	Image Submission	Multi-modal submission (packaging and seed image) for real-time verification.
FR-03	AI Inference	Execution of Custom Vision, OCR text extraction, and ResNet-50 DUS profiling.
FR-04	Feedback System	Reporting mechanism allowing farmers to submit grievances on counterfeit seed batches.
FR-05	Case Management	Dashboards for officers to track, update, and resolve seed complaints.
FR-06	Risk Analytics	Automated aggregation of regional complaints in a Batch Risk Registry for analytics.

3.3 High-Level System Architecture

The system architecture is structured as a three-tier model: the Client Layer, the Backend Orchestration Tier, and the Cloud AI and Data Infrastructure Tier. As illustrated in Figure 3.1, client applications do not interact directly with the backend API; rather, all requests are routed through a server-side proxy in the Next.js

Table 3.4: System Non-Functional Requirements

ID	Quality Attribute	Specification
NFR-01	Performance	Target end-to-end inference latency under 3.0 seconds through parallel service calls.
NFR-02	Scalability	Stateless API design facilitating horizontal scaling and microservice containment.
NFR-03	Security	HTTPS communication, JWT session control, role-based authorization (RBAC), and hashed secrets [19].
NFR-04	Reliability	Atomic database operations and soft-delete implementation for data auditability.
NFR-05	Usability	Responsive UI designed for mobile and desktop screens, optimized for low-bandwidth environments.

frontend web tier. This proxy injects JSON Web Tokens (JWT) for authentication and forwards requests to the ASP.NET Core 8 Web API. The API layer orchestrates calls to the persistent data store (SQL Server), cloud storage, and AI inference services.

3.4 System Data Flow Pipeline

The step-by-step verification pipeline and data transmission flow are depicted in Figure 3.2. The execution sequence is structured as follows:

1. **Submission:** The client uploads multi-part image data representing seed packaging and physical seeds.
2. **Routing and Proxying:** The Next.js API proxy interceptor validates user sessions, appends a secure bearer token, and executes an HTTP POST request to the backend.
3. **Parallel Execution:** The ASP.NET Core orchestration engine coordinates concurrent asynchronous requests using asynchronous task-based execution.
4. **Feature Extraction:**
 - Custom Vision classifies the packaging (Genuine, Counterfeit, Suspicious).
 - Azure AI Vision extracts text strings via OCR to identify seed batch codes.




Figure 3.1: High-level system architecture of the AgriVerify platform.

- The GCP Cloud Run ResNet-50 service determines 17 morphological parameters.
5. **Synthesis and Persistence:** The OpenAI integration consolidates these individual model predictions into a coherent diagnostic summary, which is stored in the relational database via Entity Framework Core before the final JSON payload is returned to the user interface.

3.5 Database Entity-Relationship Model

The relational database schema is modeled to ensure data consistency, minimize redundancy, and support real-time querying. Figure 3.3 defines the entity relations, key fields, and cardinalities within the AgriVerify database. The core architectural design decisions are detailed as follows:

- **User and Security Tracking:** The `User` entity supports role-based access control and maintains a 1-to-many relationship with `RefreshToken` to secure active sessions.
- **Verification History:** The `VerificationHistory` table preserves all client evaluations and maintains a soft-delete status (`IsDeleted`) to enforce an immutable audit trail.
- **Feedback and Case Resolution:** When a seed dispute is identified, a `ProductComplaint` record is initiated. This entity links to the executing officer, records ongoing investigations via `ComplaintAction`, and syncs batch metrics with the `BatchRiskRegistry`.
- **Batch Analytics:** The `BatchRiskRegistry` dynamically compiles aggregation data (complaint counts, severity indices, and regional distributions) to optimize query performance for system dashboards.




Figure 3.2: Data flow diagram for the seed verification pipeline




Figure 3.3: Simplified entity-relationship diagram for the AgriVerify domain model. Cardinality notation follows UML conventions.

Chapter 4

Deep Learning Model for Paddy Seed Quality Assessment

4.1 Dataset Description

4.1.1 Data Sources and Collection (Kaggle)

We built our training dataset by combining two sources from Kaggle. The first is a publicly available dataset, and the second is a custom collection we put together in partnership with India’s Council of Agricultural Research (ICAR) facility in Manipur. By bringing these two sources together, we got a larger and more realistic dataset—one that includes seed images captured in different ways and from different parts of India. This diversity is crucial because it helps the model learn to recognize seeds regardless of lighting conditions or camera setup.

4.1.1.1 Public Source

The first source is the Paddy Seeds Quality Classification dataset published on Kaggle by Ayyan Arkadalkani:

`https://www.kaggle.com/datasets/ayyanarkadalkani/paddy-seeds-quality-classification-dataset`

This dataset contains 1,213 labeled images split into three categories: Healthy, Damaged, and Fake/Low Quality seeds. The images take up about 183 MB of storage. These images were captured under many different lighting conditions and backgrounds, which made them a solid foundation for our model to learn general patterns about what different seed qualities look like.

4.1.1.2 Custom Source

The second source is a dataset we collected directly from ICAR’s Manipur laboratory. ICAR researchers took photographs of paddy seeds under controlled lab conditions, focusing on the specific varieties grown in northeastern India. This custom collection includes 350 images occupying about 24.7 MB of storage. We’ve published this dataset on Kaggle as an open resource so other researchers can use it:

<https://www.kaggle.com/datasets/nilambarelangbam/paddy-seed-images-collected-from-icar-manipur>

When we combined both datasets, we ended up with 1,563 images that represent a wider variety of seed conditions than either dataset alone. This combined dataset better reflects what farmers and agricultural officers actually encounter in the field across different regions of India.

4.1.2 Dataset Statistics and Class Distribution

After merging the two datasets, we had a total of 1,563 paddy seed images divided into three quality categories. A detailed breakdown of the datasets, their image counts, storage sizes, and public Kaggle repositories is summarized in Table 4.1, while the class-wise distribution of these images is provided in Table 4.2.

Table 4.1: Dataset Source Summary

Source	Images	Size	Kaggle Identifier
Paddy Seeds Quality Classification (Ayyan Arkadalkani)	1,213	~183 MB	ayyanarkadalkani/paddy-seeds-quality-classification-dataset
ICAR Manipur Custom Dataset (Our Collection)	350	~23.7 MB	nilambarelangbam/paddy-seed-images-collected-from-icar-manipur
Combined Total	1,563	~208 MB	—

As you can see, we have more healthy seed images than damaged or fake ones—which makes sense because in real life, most seeds are healthy. However, this imbalance could cause the model to become biased toward recognizing healthy seeds and miss the problematic ones. To fix this, we used a technique called weighted sampling during training: when building each training batch, we ensured

Table 4.2: Class-wise Distribution of the Combined Dataset

Class	Number of Images	Percentage (%)
Healthy Paddy Seeds	~750	~48.0
Damaged Paddy Seeds	~510	~32.6
Fake/Low Quality Seeds	~303	~19.4
Total	1,563	100.0

each quality category contributed fairly, regardless of how many examples we had of each type.

4.1.3 Train/Validation/Test Split Strategy

We divided the 1,563 images into three separate groups: one for training, one for validation, and one for testing. A key principle we followed was stratified splitting, which means we made sure each group had roughly the same mix of healthy, damaged, and fake seeds as the overall dataset. The proportions, approximate image counts, and primary objectives of each split are summarized in Table 4.3. This partitioning prevents any group from accidentally having too many of one type, which could give us misleading results.

Table 4.3: Dataset Split Statistics

Split	Proportion	Approx. Images	Purpose
Training Set	80%	~1,250	Used to teach the model by adjusting its internal weights
Validation Set	10%	~157	Used to tune settings and watch for the model learning too much from training data
Test Set	10%	~156	Held completely separate to give a final, unbiased measure of how well the model works

Here’s how we used each group: The training set is where the model learns—we show it images and let it adjust its internal parameters to get better at classification. The validation set is our quality control checkpoint—after each training cycle, we test on these images to see if the model is still improving or starting to overfit (basically, memorizing the training data instead of learning general patterns). The test set is completely hidden until the very end. We only use it once to get a fair, final assessment of the model’s performance on images it has never seen before.

4.1.4 Data Preprocessing and Augmentation

Before showing images to the neural network, we prepared them using a standard pipeline. All images were resized to 224×224 pixels—this is the standard size that ResNet50 (our neural network architecture) expects. We then adjusted the color values using ImageNet statistics (numbers from a large, well-known image dataset):

$$\mu = [0.485, 0.456, 0.406], \quad \sigma = [0.229, 0.224, 0.225] \quad (4.1)$$

This normalization step is important because the pre-trained model we started with was trained on images normalized this way, so we keep the same approach for consistency.

To make the model more robust and help it learn from our limited training data, we applied six data augmentation techniques during training. Think of augmentation as artificially creating new training examples by tweaking the original images in realistic ways:

1. **Random Resized Crop**—We randomly crop a portion of the image and resize it to 224×224 pixels. This simulates seeds being photographed from different distances and positions within the frame.
2. **Horizontal Flip**—We randomly flip images left-to-right about half the time. Seeds have a natural bilateral symmetry, so this helps the model learn that a flipped seed is still the same seed.
3. **Vertical Flip**—We randomly flip images top-to-bottom about half the time, adding more variation in how the seeds appear.
4. **Random Rotation**—We rotate images randomly by up to $\pm 30^\circ$ (30 degrees in either direction). Real seed photos might not be perfectly aligned, so this prepares the model for that reality.
5. **Translation/Random Shift**—We randomly shift the image content left, right, up, or down. This prevents the model from relying on the seed always appearing in the center of the photo.
6. **Colour Jitter**—We randomly adjust brightness, contrast, color saturation, and hue. Since real agricultural photos are taken under different lighting conditions, this helps the model handle that variation.

4.2 Model Architecture

4.2.1 ResNet50 Backbone (25.6M Parameters)

We chose ResNet50 as the foundation of our classification model. This neural network has 50 layers and about 25.6 million parameters (adjustable settings inside the network). What makes ResNet special is its use of "skip connections"—shortcuts that let information flow directly around groups of layers. This design makes it much easier to train very deep networks because it prevents the gradient (the training signal) from fading as it travels backward through the network.

Rather than training from scratch, we started with weights that were pre-trained on ImageNet, a massive dataset of millions of everyday images. This pre-trained starting point is like giving the model a head start: it already knows how to recognize edges, textures, and shapes, so we only need to teach it the specific patterns that distinguish different seed qualities.

4.2.2 Custom Classification Head

On top of the ResNet50 backbone, we added a custom section designed specifically for seed classification. This section takes the features learned by ResNet and processes them into a final decision about whether a seed is Healthy, Damaged, or Fake/Low Quality. Here's what this custom section does:

1. Compresses the spatial information into a single summary vector (Global Average Pooling)
2. Passes it through a fully connected layer with 512 hidden units
3. Applies Batch Normalisation to stabilize the values
4. Uses ReLU activation to introduce non-linearity
5. Applies Dropout (randomly ignoring 50% of values) to reduce overfitting
6. Outputs 3 final scores (one for each seed quality class) with softmax to convert them to probabilities

4.2.3 Anti-Overfitting Techniques

Overfitting is the enemy of machine learning—it's when a model memorizes the training data instead of learning generalizable patterns. We used several techniques

to prevent this:

- **Dropout**—During training, we randomly "turn off" 50% of neurons in the classification head. This forces the network to learn redundant representations, so it can't rely on a few specific neurons. Think of it like training a team where you randomly swap out players—the team learns to work together, not depend on one star player.
- **Batch Normalisation**—After each layer, we normalize the values to keep them in a healthy range. This stabilizes training and acts as a natural regularizer.
- **Label Smoothing**—Instead of forcing the network to be 100% confident about each prediction, we "soften" the targets slightly (using $\epsilon = 0.1$). This prevents the model from becoming overconfident and improves generalization.
- **Weight Decay**—We penalize large weights during training, keeping the model's parameters from growing too large. A smaller model is typically more generalizable.
- **Data Augmentation**—As described earlier, we artificially expand the training data with slight variations, giving the model more diverse examples to learn from.
- **Early Stopping**—We continuously monitor performance on the validation set. If it starts getting worse, we stop training immediately and use the best weights we've found so far.

4.3 DUS Parameter Integration

4.3.1 Feature Extraction—17 Physical Parameters

Beyond just classifying seeds as Healthy, Damaged, or Fake, we wanted our model to extract physical measurements that agriculture experts understand. The standard set of measurements in agriculture is called DUS—Distinctiveness, Uniformity, and Stability. We programmed our model to extract 17 of these physical parameters from each seed image:

- How long and how wide each seed is
- The ratio of length to width

- The perimeter and total area of the seed
- How circular the seed is (circularity) and how solid it is (solidity)
- How elongated the seed is (eccentricity)
- Mathematical descriptions of how the seed’s mass is distributed
- Hu moments—a special mathematical descriptor of seed shape that works the same way whether the seed is rotated, scaled, or shifted in the image

These 17 parameters bridge the gap between what the neural network learns (visual patterns) and what agricultural experts care about (standard physical measurements). This makes our system more transparent and trustworthy to end users.

4.3.2 Pixel-to-Physical Unit Conversion

When we measure seeds in a photograph, the measurements are in pixels, which depends on the camera and how far away it was. To convert pixel measurements to real-world millimetres, we include a reference object (a ruler with known dimensions) in each photograph. We then calculate:

$$\text{scale} = \frac{\text{reference_length_mm}}{\text{reference_length_pixels}} \quad (4.2)$$

This scale factor is applied to all measurements, so when we report that a seed is 6.2 mm long, it’s truly 6.2 millimetres—not dependent on camera distance or image resolution.

4.3.3 Variety Matching (12 Reference Varieties)

We created a reference database of 12 standard paddy seed varieties used in India, based on official ICAR documentation. For each variety, we computed the typical values of all 17 physical parameters from multiple certified seed samples. When a new seed image is analyzed, we compare its 17 parameters against these 12 reference profiles using a simple distance calculation:

$$d_i = \sqrt{\sum_{j=1}^{17} (p_j - r_{i,j})^2} \quad (4.3)$$

Here, p_j is one of the 17 measurements from the new seed, $r_{i,j}$ is the standard value for variety i , and d_i tells us how different the new seed is from variety i . The variety with the smallest distance is the best match—this is the variety the seed most closely resembles. We report the top 3 matches so users can see alternatives if needed.

4.4 Training Pipeline

4.4.1 Phase 1—Classification Head Training (Epochs 1–8)

We trained the model in two phases. In Phase 1 (the first 8 training cycles), we froze the ResNet50 backbone—meaning its weights didn’t change—and only trained the custom classification head we added on top. Think of this like keeping a well-established expert’s knowledge fixed and just training a new student to interpret their findings. We used a learning rate of 0.001, which controls how big the weight updates are. A learning rate that’s too high causes instability; too low and training is slow. The batch size (number of images processed before updating weights) was set to 32.

At the end of each of these 8 epochs, we checked performance on the validation set to make sure the head was learning properly and not starting to overfit.

4.4.2 Phase 2—Full Fine-Tuning (Epochs 9–15)

In Phase 2 (epochs 9–15), we unfroze the entire network and trained everything end-to-end. We reduced the learning rate by 10 times to 0.0001—a smaller rate because we’re now adjusting a fully-trained backbone and don’t want to destroy what it’s already learned. This phase refined the backbone’s features to specialize better for seed classification.

We monitored validation performance carefully and stopped training at epoch 15 when we noticed the validation loss starting to increase (a sign the model was beginning to overfit). We then loaded the best weights we’d found during the entire training run—these became our final model.

4.4.3 Hyperparameter Configuration and Optimiser

A comprehensive summary of the optimizer, learning rates, regularizers, and general training configurations is provided in Table 4.4:

Table 4.4: Hyperparameter Configuration and Training Settings

Hyperparameter	Value	Notes
Optimiser	AdamW	Smart learning rate that adapts during training
Phase 1 Learning Rate	0.001	Used while training only the head
Phase 2 Learning Rate	0.0001	Used for full network fine-tuning (10 times smaller)
Weight Decay (L2)	0.0001	Penalizes large weights to keep the model simple
Batch Size	32	Number of images processed before updating weights
Loss Function	Cross-Entropy + Label Smoothing	Prevents overconfidence with $\epsilon = 0.1$
LR Scheduler	StepLR	Reduces learning rate further after 8 epochs
Gradient Clipping	Max norm = 1.0	Prevents unstable training due to large gradients
Early Stopping Patience	10 epochs	Stops if validation loss doesn't improve for 10 epochs
Input Resolution	224×224 pixels	Standard size for ResNet50
Total Epochs	15	Stopped when overfitting began

We chose the AdamW optimiser because it's smart about adjusting the learning rate for different parameters and handles weight decay better than older methods.

4.4.4 Model Export Formats

After training finished, we saved the model in multiple formats to ensure compatibility with different environments. The exported files, including their sizes and target use cases, are listed in Table 4.5.

Table 4.5: Model Export Formats

Format	Extension	Size	Use Case
PyTorch Native	.pth	~98 MB	The original format; used by FastAPI for serving
TorchScript	.pt	~97 MB	For production servers and systems using C++
ONNX	.onnx	~98 MB	For cross-platform use and systems not using PyTorch
Configuration	.json	~2 KB	Metadata about the model, class labels, and settings
DUS Reference DB	.csv	<1 MB	The 12 reference seed varieties for matching

4.5 Model Comparison and Selection

Before committing to ResNet50, we tested four popular neural network architectures on the same dataset, using the same training approach, to see which performed best. We compared ResNet50, MobileNetV2, DenseNet121, and GoogleNet, with the comparative results detailed in Table 4.6.

Table 4.6: Architecture Comparison on Test Set

Model	Accuracy	Precision	Recall	F1 Score	Time (min)
ResNet50	99.57%	99.57%	99.57%	99.57%	16.4
DenseNet121	99.57%	99.57%	99.57%	99.57%	15.0
GoogleNet	99.57%	99.57%	99.57%	99.57%	12.2
MobileNetV2	99.13%	99.15%	99.13%	99.13%	8.5

The results show that ResNet50, DenseNet121, and GoogleNet all achieved the same 99.57% accuracy, demonstrating that our dataset is rich enough that all three architectures can learn to classify seeds nearly perfectly. MobileNetV2 achieved slightly lower accuracy at 99.13% but trained much faster—in just 8.5 minutes compared to 16.4 minutes for ResNet50.

We chose ResNet50 for three practical reasons:

1. Its design with skip connections is well-understood in the research community, making it easier to diagnose training issues and fine-tune behavior.
2. It has excellent support across different deployment platforms—PyTorch, TorchScript, ONNX, and more—giving us flexibility in where and how we deploy it.
3. Most published agricultural image classification research uses ResNet50, so our results are directly comparable to other teams’ work, helping the scientific community benchmark progress.

The extra 4.4 minutes of training time is a small price to pay for these advantages, especially since training only happens once.

4.6 Final Model Performance Summary

After completing both training phases, we evaluated our final ResNet50 model on the test set—the images we’d held back and never shown the model during

training. The final evaluation metrics are compiled in Table 4.7, and a detailed summary of the training trajectory and learning stability is provided in Table 4.8.

Table 4.7: Final Test Set Performance—ResNet50

Metric	Value	Status
Test Accuracy	99.57%	Far exceeds industry standard (85–90%)
Precision	~0.99	Outstanding—almost no false positives
Recall	~0.99	Outstanding—catches almost all cases
F1 Score	~0.99	Excellent—perfectly balanced performance
Train-Validation Gap	<5%	Great generalization—model isn’t overfitting

Table 4.8: Training Behaviour Summary

Observation	Detail
Initial Convergence	The model started at about 86% accuracy and quickly improved to near its best performance within the first few epochs
Loss Trajectory	The loss (error) steadily decreased from about 1.0 to nearly 0.0 across all 15 epochs
Mid-training Stability	After the initial improvement phase, training was smooth and consistent with no wild fluctuations
Overfitting Control	The gap between training accuracy and validation accuracy never exceeded 5%, indicating the model learned general patterns rather than memorizing training data
Training Stopped at	Epoch 15—we detected the first signs of overfitting and stopped immediately, preserving the best weights we’d found

The final accuracy of 99.57% is significantly better than the typical industry standard of 85–90% for automated seed quality classification. The near-perfect precision (very few false alarms) and recall (catches nearly all problems) across all three seed quality categories mean our model is reliable and trustworthy for real agricultural use. This is exactly what we needed—a system that farmers and agricultural officers at ICAR Manipur can depend on to accurately identify problematic seeds before they’re planted.

Chapter 5

Model Serving API

Deploying a trained model as a live service introduces a different set of problems than training it. Model development focuses on architecture and accuracy; the serving layer determines response latency, concurrent throughput, and uptime under real load [20]. In AgriVerify, the PyTorch ResNet-50 classifier from Chapter 4 runs inside a dedicated inference microservice built specifically for this purpose.

This chapter describes the design of that microservice, implemented with FastAPI and Uvicorn [21]. It covers the asynchronous request lifecycle, the tensor normalization pipeline, the API endpoint contracts, and the containerised deployment approach using Docker multi-stage builds on GCP Cloud Run [22].

5.1 API Architecture and Request Orchestration

5.1.1 FastAPI and ASGI Serving Stack

The inference microservice runs on FastAPI atop the Uvicorn ASGI runtime. FastAPI was chosen for its native async I/O support, Pydantic-based request validation and automatic OpenAPI schema generation [21]. Uvicorn’s `uvloop`-backed event loop handles large numbers of concurrent connections with low memory overhead.

The stack separates network deserialization, Pydantic validation, OpenCV preprocessing and PyTorch execution into distinct layers. Computationally expensive inference operations run in dedicated thread pools and do not block the ASGI event loop.

5.1.2 End-to-End Model Request Flow

As illustrated in Figure 5.1, an inference request begins when the ASP.NET Core backend sends a multipart POST containing a seed image. FastAPI validates the MIME type and payload boundaries, then decodes the byte stream into an OpenCV matrix in memory.

From there, two operations run in parallel: a PyTorch inference engine — with ResNet-50 weights pre-loaded to avoid cold-start delays — computes classification logits [4], [23], while OpenCV contour routines extract 17 DUS (Distinctness, Uniformity, and Stability) morphological parameters [24]. The results are aggregated, scored against reference cultivars, and returned as JSON.

The sequential execution steps of this request flow are defined as follows:

1. **Request:** The FastAPI endpoint receives the image payload via a multipart POST request from the backend proxy.
2. **Preprocessing:** The image is decoded into an OpenCV matrix, resized to 224×224 pixels, and normalized using the standard ImageNet channel moments.
3. **Model Inference:** The pre-loaded ResNet50 model runs inference in a background thread pool to generate raw classification logits.
4. **Morphological Post-processing:** Simultaneously, OpenCV contour analysis extracts the 17 DUS physical measurements, and the system performs variety matching against the 12 reference profiles.
5. **Response:** The API aggregates the model predictions, confidence levels, variety distance scores, and DUS metrics into a structured JSON payload returned to the client.




Figure 5.1: Asynchronous model request processing flow.

5.2 API Specification and Endpoint Contracts

The microservice exposes a tightly constrained RESTful API. Tables 5.1 and 5.2 define the schema contracts for communication with the backend orchestrator.

Table 5.1: API Specification for the POST `/predict` Model Inference Endpoint

Attribute	Endpoint Constraint / Schema Definition
Route & Auth	POST <code>/predict</code> Requires Authorization: Bearer <code><token></code>
Payload	multipart/form-data (File: <code>image</code>). Types: JPEG/PNG. Max size: 10MB.
200 OK Response	<code>{"predicted_label": "Genuine", "confidence": 0.98, "inference_ms": 42.8, "dus_parameters": {"length_mm": 8.4, "aspect_ratio": 3.1}}</code>
Error Codes	400 : Malformed payload 413 : File too large 415 : Invalid MIME type

Table 5.2: API Specification for the GET /health Liveness Endpoint

Attribute	Endpoint Constraint / Schema Definition
Route	GET /health (Unauthenticated liveness/readiness probe)
Probing Logic	Confirms Uvicorn ASGI event loop and verifies ResNet-50 weights are loaded in memory.
200 OK	<code>{"status": "online", "model_loaded": true, "device": "cuda:0"}</code>
503 Error	Service unavailable (model weights corrupted or GPU allocation failure).

5.3 Containerisation and Deployment Architecture

The service uses a multi-stage Docker build to balance reproducibility against image size [22]. The **Builder Stage** compiles Python binaries (.whl) with `gcc`. The **Runtime Stage** copies those pre-compiled binaries into a minimal, security-hardened Debian base (`python:3.10-slim`), discarding the build toolchain entirely. The final image is approximately 850MB — down from over 3GB — with a correspondingly smaller attack surface. The automated containerization, versioning, and deployment pipeline is illustrated in Figure 5.2.




Figure 5.2: Automated CI/CD pipeline and Serverless GCP Deployment Architecture.

The image is deployed to GCP Cloud Run [25]. Cloud Run manages HTTPS termination and load balancing, and scales to zero when idle, so infrastructure costs track actual query volume rather than reserved capacity.

5.3.1 CI/CD Deployment Pipeline

The automated build and release lifecycle is managed via a continuous integration and continuous deployment (CI/CD) pipeline implemented with GitHub Actions. Whenever updates are merged into the repository’s production branch, the pipeline executes the following stages:

- **Build Validation:** Runs linter checks and automated tests to ensure code quality.
- **Docker Compilation:** Triggers a multi-stage Docker build, generating a security-hardened and size-optimized production container.
- **Registry Storage:** Authenticates with Google Cloud and pushes the compiled image to the secure GCP Artifact Registry with versioned tags.
- **Serverless Deployment:** Deploys the container to GCP Cloud Run, utilizing serverless auto-scaling to dynamically adjust resource allocation based on request volumes.

5.3.2 Live Production API Endpoint

To demonstrate the functionality of the serving layer under real-world conditions, the API is exposed as a live web service. Other client applications and services in the AgriVerify platform (such as the ASP.NET Core backend proxy) communicate with this service at:

`https://resnet-api-297419987783.asia-south1.run.app/`

Chapter 6

Backend System Architecture and Implementation

The backend of the AgriVerify system is responsible for managing business logic, AI model integration, authentication, database operations, and communication between frontend clients and cloud services. The system was developed using ASP.NET Core Web API following the Clean Architecture approach to achieve modularity, maintainability, and scalability.

The backend integrates multiple services including SQL Server, Azure AI services, JWT authentication, and cloud storage. The architecture separates core business logic from infrastructure and presentation concerns, enabling easier maintenance and future system expansion.

6.1 Architectural Design

The backend follows a layered architecture based on Clean Architecture principles, with the relationship and dependency flow between these layers illustrated in Figure 6.1. Each layer performs a dedicated responsibility while maintaining loose coupling between components.




Figure 6.1: Backend Clean Architecture Layers

The layered architecture improves:

- Separation of concerns
- Code maintainability
- Scalability
- Testability
- Reusability

6.2 Technology Stack

The backend uses modern technologies for API development, authentication, AI integration, and database management. The core technologies and libraries comprising this stack are summarized in Table 6.1.

Table 6.1: Backend Technology Stack

Technology	Purpose
ASP.NET Core Web API	Backend API framework
C# and .NET 8	Backend programming environment
Entity Framework Core	Database ORM and migrations
SQL Server	Relational database management
JWT Authentication	Secure user authentication
ASP.NET Core Identity	User and role management
Azure Custom Vision	Seed authenticity classification
Azure OCR	Text extraction from seed packets
Azure OpenAI	AI-generated explanations
Azure Blob Storage	Cloud image storage
FluentValidation	Input validation
AutoMapper	DTO mapping
Swagger	API documentation
Serilog	Request logging and monitoring

6.3 Project Structure

The backend source code is divided into multiple layers and modules to maintain clear responsibility boundaries, as outlined in Table 6.2.

Table 6.2: Backend Project Structure

Folder / Layer	Purpose
Domain	Business entities and interfaces
Application	Services, DTOs and validation
Infrastructure	Database and cloud integrations
Presentation	API controllers and middleware
Persistence	Repositories and migrations
Services	AI and Azure integrations
Configurations	Application configuration files
Middleware	Authentication and exception handling

6.4 Domain Layer

The Domain Layer contains the core business entities and contracts of the application. This layer is independent from external frameworks and infrastructure components.

6.4.1 Core Entities

The major entities in the system include:

- User
- Farmer
- Officer
- Seed
- SeedBatch
- DetectionResult
- Complaint
- AuditLog

These entities represent the primary agricultural and administrative components of the platform.

6.4.2 Interfaces

Interfaces define contracts between layers and improve abstraction.

- IRepository
- IVerificationService
- IAuthenticationService
- IComplaintService
- IAnalyticsService

This approach improves modularity and testing flexibility.

6.5 Application Layer

The Application Layer coordinates business workflows and request processing.

6.5.1 Application Services

The system includes several services responsible for handling business operations:

- Verification Service
- Authentication Service
- Complaint Service
- Analytics Service
- Admin Service

These services manage AI processing, validation, database interaction, and response generation.

6.5.2 DTOs

Data Transfer Objects (DTOs) are used for request and response communication between layers. The primary DTOs utilized in the system and their respective purposes are listed in Table 6.3.

Table 6.3: Important DTOs

DTO	Purpose
VerificationRequest	Seed verification request
VerificationResponse	AI verification result
LoginRequest	User login request
AuthResponse	JWT authentication response
CreateComplaintRequest	Complaint submission
OfficerDto	Officer information
AdminDashboardResponse	Dashboard statistics

6.5.3 Validation

FluentValidation is integrated into the application layer to validate:

- File formats
- Image sizes
- Email structure
- Password complexity
- Required fields

This reduces invalid requests and improves system reliability.

6.6 Infrastructure Layer

The Infrastructure Layer manages database operations, external integrations, and cloud services.

6.6.1 Database Management

Entity Framework Core is used for:

- Database migrations
- Entity mapping
- Relationship configuration
- Query execution
- Transaction management

6.6.2 Repository Pattern

The repository pattern abstracts database operations from business logic.

- Generic Repository
- Seed Repository
- Complaint Repository
- Detection Result Repository

6.6.3 AI and Cloud Integrations

The system integrates multiple Azure services for AI processing and storage, as detailed in Table 6.4.

Table 6.4: External Service Integrations

Service	Purpose
Azure Custom Vision	Seed authenticity classification
Azure OCR	Text extraction from seed packets
Azure OpenAI	AI-generated recommendations
Azure Blob Storage	Cloud image storage
ResNet Classifier	Paddy seed classification

6.7 Database Schema

The backend database maintains relationships between users, seed batches, verification results, complaints, and audit records. The database schema and its corresponding primary/foreign key mappings are represented in the Entity-Relationship Diagram in Figure 6.2.




Figure 6.2: Entity Relationship Diagram – FakeSeedDetection System

The database schema supports:

- User management
- Seed batch tracking
- AI verification history
- Complaint workflows

- Audit logging

Foreign key relationships maintain referential integrity across all entities.

6.8 Presentation Layer

The Presentation Layer exposes RESTful APIs for frontend communication.

6.8.1 API Controllers

The backend contains multiple API controllers responsible for different system modules. These controllers and their design objectives are summarized in Table 6.5.

Table 6.5: API Controllers

Controller	Purpose
AuthController	Authentication operations
VerificationController	Seed verification endpoints
ComplaintsController	Complaint management
AdminController	Administrative operations
AnalyticsController	Dashboard analytics
ReferenceDataController	Dropdown reference data

6.8.2 Middleware

Middleware components manage authentication, logging, and centralized exception handling.

- Authentication Middleware
- Request Logging Middleware
- Exception Handling Middleware

These middleware components improve system security and monitoring.

6.9 Authentication and Security

The backend uses JWT-based authentication with role-based authorization, following the flow depicted in Figure 6.3.




Figure 6.3: JWT Authentication Flow

Security features include:

- JWT Bearer authentication
- Role-based authorization
- Password hashing
- Refresh token support
- HTTPS communication
- Protected API endpoints

The system supports multiple user roles:

- Farmer
- Officer
- Administrator

6.10 Verification Workflow

The seed verification workflow integrates AI classification, OCR analysis, and database persistence, as shown in Figure 6.4.




Figure 6.4: Seed Verification Workflow

The workflow includes:

1. Image upload
2. Request validation
3. Image quality analysis
4. AI classification
5. OCR text extraction
6. Database persistence
7. Response generation

6.11 Configuration and Deployment

The backend uses multiple configuration files for different deployment environments.

- `appsettings.json`
- `appsettings.Development.json`
- `appsettings.Production.json`

These files manage:

- Database connections
- JWT settings
- Azure credentials
- Logging configuration
- API endpoints

Application startup is configured through `Program.cs`, which initializes:

- Database services
- Authentication services
- Dependency injection
- Middleware pipeline
- Swagger documentation
- Logging services

6.12 Summary

This chapter presented the backend architecture and implementation of the AgriVerify system. The backend was developed using ASP.NET Core Web API following the Clean Architecture pattern.

The system integrates database management, AI services, authentication, cloud storage, and REST APIs into a modular and scalable architecture. Entity Framework Core, JWT authentication, Azure AI services, and structured middleware improve system reliability, maintainability, and security.

The backend successfully supports seed verification, complaint management, user authentication, analytics, and administrative operations within the AgriVerify platform.

Chapter 7

Frontend System Architecture and Implementation

The AgriVerify frontend connects the inference engine and backend services to a browser interface that has to work for two quite different users: farmers on mid-range Android handsets with variable connectivity, and state administrators managing supply chain dashboards. The stack — Next.js 16, React 19, TypeScript 5, and Tailwind CSS — was chosen to handle that range without the interface becoming slow or hard to maintain [26], [27].

7.1 Frontend Architecture and Design Patterns

The architecture separates UI rendering from data fetching, API communication, and authorization. Each concern lives in its own layer.

7.1.1 Architectural Layers

A request passes through four layers before reaching backend services (Figure 7.1):

- **Presentation UI (React & Next.js):** Handles component rendering, form validation, and responsive layout [28].
- **Client-Side State (TanStack Query):** Manages async data fetching, caching, and optimistic mutations [29].
- **Service Client (Axios):** Wraps network calls, standardising payloads and error handling.

- **API Proxy (Next.js Routes):** Routes client requests to ASP.NET Core endpoints without exposing internal URLs to the browser.




Figure 7.1: Frontend Layered Architecture and Request Execution Flow

7.1.2 Secure Proxy Architecture

All outbound requests route through internal Next.js serverless endpoints (`/api/proxy/*`), which avoids CORS issues by keeping cross-origin communication server-side. Internal microservice addresses stay out of the browser, secrets remain on the server, and security headers are applied before responses reach the client [30].

7.2 Technology Stack Analysis

Table 7.1 summarises the stack and the reasoning behind each choice.

Table 7.1: Frontend Technology Stack

Technology	Architectural Reasoning
Next.js 14	Provides Server-Side Rendering and built-in API proxy routing [31].
React 18	Declarative UI rendering with efficient Virtual DOM reconciliation.
TypeScript 5	Compile-time type checking with backend DTO parity.
Tailwind CSS	Utility-first styling with minimal production bundle output [13].
Axios & TanStack	HTTP pipeline management, token injection, and cache synchronization [29].
Zustand 4.5	Lightweight atomic state slices for session persistence.

7.3 Modular Project Structure

The source tree follows domain-driven modularity. Key directories: `src/app` (Next.js routing), `src/components` (reusable UI primitives), `src/hooks` (data fetching logic), and `src/types` (TypeScript DTOs matching backend entities). Keeping these separate means business logic and presentation concerns do not bleed into each other.

7.4 User Role-Based Access Control (RBAC)

The frontend uses Role-Based Access Control (RBAC) to separate public and protected routes.

7.4.1 Role Partitioning

The application supports four user profiles: **Farmers** (verification and dispute submission), **Agricultural Officers** (inspection queues and evidence review), **Administrators** (system governance and telemetry), and **Public Users** (anonymous checks).

7.4.2 Route Protection Middleware

A `RoleGuard` component checks JWT claims against required permissions before mounting any protected React component, preventing privilege escalation at the route level.

7.5 Authentication and Secure Persistence

Sessions are secured using a dual-token JWT exchange designed to limit exposure to XSS attacks [30].

7.5.1 Token Storage Architecture

Short-lived access tokens are held in the Zustand in-memory store, keeping them out of reach of malicious scripts. Long-lived refresh tokens are persisted in encrypted storage or `HttpOnly` cookies, allowing background session renewal without exposing token values to client-side code.

7.5.2 Automated Token Rotation

Axios interceptors handle token expiry without user intervention. On a 401 `Unauthorized` response, the interceptor pauses pending requests, runs a refresh cycle, and replays the original request with the new token. From the user's side, the session continues uninterrupted.

7.6 API Integration and Type Safety

TypeScript DTOs enforce compile-time verification of HTTP payloads across all integration modules — Auth, Verification, Analytics, and Complaint Services. Each mirrors the corresponding ASP.NET Core schema, so data arriving at the frontend matches what the UI expects without runtime coercion or defensive casting.

7.7 State Management

State is split by volatility: server-driven data and client-local data are managed separately.

7.7.1 Asynchronous Server State (TanStack Query)

Network-driven data (dispute queues, telemetry) is managed by TanStack Query, which handles stale-while-revalidate caching, background polling, and optimistic UI updates. On slow connections, this keeps the interface usable even when the server is slow to respond [29].

7.7.2 Synchronous Client State (Zustand)

UI state that doesn't touch the network (active theme, session context) lives in Zustand. Its atomic slices limit unnecessary re-renders, which matters on lower-end devices where battery and memory are constrained.

7.8 Responsive UI Design and Field Usability

The UI is built mobile-first using Tailwind CSS. Most farmers will access the application on a mid-range Android handset, often outdoors, so the interface is designed around that constraint rather than treating it as an edge case [27].

7.8.1 Standardized UI Library

The interface is built from a set of reusable primitives (Table 7.2) to keep visual behaviour consistent across screens.

Table 7.2: Standardized UI Components

Component	Role
Data Table	Scrollable grids with sticky headers and dynamic pagination for inspection data.
Status Badge	Colour-coded indicators (e.g., Green/Genuine) for fast visual parsing.
Loading Skeleton	Animated placeholders that prevent layout shifts during network delays.
Input Modal	Focused overlays that trap keyboard navigation for accessible data entry.

7.8.2 Accessibility Enhancements

The UI meets WCAG AAA contrast ratios for readability in bright outdoor conditions, uses a minimum touch target size of 48×48 px to reduce input errors, and includes full ARIA semantic tagging for screen-reader compatibility.

Chapter 8

System Integration

This chapter explains how all the pieces of AgriVerify work together. We’ve designed the system to be flexible, modular, and secure—the frontend (what users see) is completely separate from the backend (where the hard work happens), and both are separated from the AI services. This separation makes the system easier to build, test, and scale.

8.1 System Integration Matrix

AgriVerify communicates across four main boundaries. To keep everything secure, all communication flows through a secure proxy layer. This proxy acts like a security guard: it sits between the browser and the backend, injecting authentication tokens and preventing users’ browsers from ever directly connecting to backend systems. Table 8.1 shows how everything talks to everything else, what format the data takes, and how we keep it all secure.

8.2 End-to-End Architectural Workflow

When a farmer submits a seed image, it travels through several stages: from the browser, through the proxy, into the backend for processing, out to AI services for analysis, and finally gets stored in the cloud. Meanwhile, our deployment system works behind the scenes automatically managing code updates. Figure 8.1 shows both journeys: how a verification request flows through the system, and how code automatically gets deployed to keep the system running.

Let’s walk through what happens at each stage:

1. **Browser to Proxy (HTTPS with file uploads):** A farmer opens AgriVer-

Table 8.1: How Different Parts of AgriVerify Communicate

Communication Path	Protocol	What Happens & Security
Browser to Proxy	HTTPS with file uploads	Axios automatically adds JWT security tokens and handles token refresh if they expire.
Proxy to Backend	REST over HTTPS	The proxy forwards all requests, checks data is valid, and hides real backend addresses from browsers.
Backend to AI	HTTPS or gRPC	Backend sends images, checks for blurriness/lighting issues, extracts text, and runs the ResNet50 model.
Cloud Services	Azure official libraries	Uses Azure Custom Vision for packaging, Azure OpenAI to write reports, and Azure Blob Storage to save images.
Deployment	Container systems	Multi-stage Docker builds prepare the code, store it securely, and deploy with automatic scaling.




Figure 8.1: How AgriVerify processes seed verification requests and automatically deploys updates. Top section shows what happens when a farmer submits images; bottom section shows how we keep the system updated.

ify on their phone or computer and takes a photo of their seed packet. They hit “upload” and the image travels over encrypted HTTPS. The Next.js proxy receives it and automatically adds their security token (so the backend knows who they are). If the security token has expired, the proxy handles getting a fresh one without bothering the farmer.

2. **Proxy to Backend (REST API over HTTPS):** The proxy forwards the farmer’s request to the ASP.NET backend, but hides the real backend address from the browser. The backend receives the image and validation data, checks that everything looks correct, and is ready to process it.
3. **Backend to AI Services (HTTPS / gRPC):** The backend doesn’t try to analyze the image itself—it’s too complex. Instead, it sends the image to specialized AI services. First, it checks if the image is blurry or too dark (if so, it warns the farmer). Then it runs three analyses in parallel: Azure’s text extraction pulls batch numbers from the packaging, Azure’s packaging classifier checks if it looks genuine, and our custom ResNet50 model extracts the 17 physical seed measurements. All three work simultaneously to save time.
4. **Cloud Storage (Azure Blob & SQL Database):** As results come back from the AI services, they’re immediately saved to the cloud: the image goes to Azure Blob Storage, and the results go into SQL Server. This ensures data is safely stored and the farmer’s verification history is preserved for future reference.
5. **Response Back to Browser:** The backend synthesizes all the AI results into a clear report using Azure OpenAI to write explanations in plain language. This report travels back through the proxy to the farmer’s browser, where they see the results within seconds.

Meanwhile, the **deployment pipeline** (shown in dashed lines) works independently in the background. Whenever our development team pushes new code to GitHub, Docker automatically builds a containerized version, stores it in Azure’s secure registry, and deploys it to Azure Container Apps—which automatically scales up or down based on how many farmers are using AgriVerify at any given time.

Chapter 9

Results and Evaluation

9.1 Model Performance

Upon completion of the training phase, the ResNet50 model was rigorously evaluated using an independent test dataset consisting of previously unseen images. Model performance was quantified using standard statistical metrics, namely Accuracy, Precision, Recall, and the F1 Score [32]. These metrics provide a multi-dimensional assessment of the model’s capacity to distinguish between genuine and defective (damaged or counterfeit) paddy seeds [4], [8].

The evaluation results demonstrate high classification performance. The model consistently classified seed samples into two primary categories: Pure (certified, legitimate seeds) and Impure (damaged, defective, or counterfeit seeds). The empirical results suggest that the AgriVerify system is suitable for deployment in real-world agricultural scenarios, particularly within resource-constrained environments.

9.1.1 Accuracy, Precision, Recall, and F1 Score

The overall performance metrics of the model on the test dataset are summarized in Table 9.1. Furthermore, the convergence behavior of the model, specifically the training/validation accuracy and loss profiles over 15 epochs, is illustrated in Figure 9.1.

An accuracy rate of 99.57% indicates that the model successfully generalizes classification features across varying seed textures and illumination conditions. The equivalence of precision and recall (both at 99.57%) is critical for agricultural applications; it signifies a minimized rate of both false positives (misclassifying

Table 9.1: Overall Model Evaluation Results

Performance Metric	Result	What This Means
Accuracy	99.57%	Out of every 100 seeds we tested, we correctly classified 99 or 100 of them
Precision	99.57%	When the system says a seed is problematic, it's right almost every time—virtually no false alarms
Recall	99.57%	The system catches almost all problematic seeds—very few slip through undetected
F1 Score	99.57%	Perfect balance between precision and recall—the system is equally good at both



Figure 9.1: How the model improved over 15 training cycles. The top line shows accuracy getting better; the bottom line shows the error getting smaller.

impure seeds as pure) and false negatives (misclassifying pure seeds as impure). The F1 Score confirms a robust balance between these two critical classification objectives [32].

Notably, the classification accuracy of 99.57% surpasses the typical baseline of 85–90% reported in existing literature for automated seed quality assessment systems [2], [6]. This performance gain validates the integration of transfer learning with tailored image augmentation and regularization techniques.

9.2 Seed Quality Results: Pure vs. Impure Seeds

To evaluate the class-specific performance of the classifier, metrics were computed independently for the pure and impure seed categories. The test set comprised 231 samples distributed across both classes, with the results summarized in Table 9.2.

Table 9.2: Performance on Each Seed Category

Seed Type	Precision	Recall	F1 Score	What This Means
Pure (Genuine)	100.0%	99.19%	99.60%	Perfect at recognizing genuine seeds—never wrongly approves a fake one
Impure (Damaged/-Fake)	99.07%	100.0%	99.53%	Perfect at detecting problems—catches every single defective seed

The model achieved 100.0% precision for the pure category, which can be attributed to the highly distinct visual markers of genuine paddy seeds, such as uniform golden-brown coloration, intact husks, and consistent surface textures. Conversely, the model demonstrated 100.0% recall for the impure category, successfully flagging every damaged, defective, or counterfeit seed. This asymmetric performance is highly desirable for agricultural quality control, ensuring that sub-standard seed batches are consistently intercepted while minimizing the false rejection of certified stock [7].

The high classification performance is primarily due to class-weighted loss functions and data augmentation strategies employed during training. These techniques mitigate class imbalance issues and prevent the network from exhibiting classification bias toward the majority class.

9.3 Confusion Matrix Analysis

To analyze the distribution of classification errors, a confusion matrix was constructed. The quantitative prediction distributions are detailed in Table 9.3, and the corresponding normalized confusion matrix is visualized in Figure 9.2.

Table 9.3: Detailed Look at Every Prediction

Actual Seed Type	What the Model Predicted	
	Impure	Pure
Impure	107	0
Pure	1	123




Figure 9.2: Visual representation of how often the model got predictions right or wrong.

Out of the 231 test samples evaluated, the model correctly classified 230 samples. The specific classification outcomes are analyzed below:

Impure Class (107 Samples): All 107 impure samples were correctly classified

as impure (True Positives), yielding zero false negatives. From an agricultural safety perspective, this prevents the inadvertent distribution and planting of defective or counterfeit seed batches.

Pure Class (124 Samples): The model correctly identified 123 pure samples (True Negatives), with only 1 sample misclassified as impure (False Positive). This conservative error is preferable in agricultural workflows; while a false positive merely necessitates a secondary manual inspection, a false negative (failing to identify impure seeds) could lead to complete crop failure and economic losses for the farmer.

9.4 Regularization Effectiveness

With a dataset size of 1,563 images, training deep neural networks introduces a significant risk of overfitting, where the model memorizes high-frequency noise in the training set rather than extracting generalizable morphological features. To mitigate this, several regularization techniques were evaluated, as compared in Table 9.4.

Table 9.4: How Our Anti-Overfitting Techniques Helped

Configuration	Train vs. Validation Gap	What Happened
No Protection	25–30%	Bad—the model memorized training data and performed much worse on new seeds
Dropout Only	8–12%	Better, but still some overfitting on tricky seed textures
All Techniques Combined	2.8%	Excellent—tight control with less than 5% difference

The validation-to-training accuracy gap was reduced to 2.8% when all regularization methods were active. This small generalization error indicates that the network successfully extracted generalizable features. This was achieved through the following architectural and optimization strategies:

- **Weight Decay (AdamW):** L2 regularization was applied to penalize large weight magnitudes, thereby constraining model complexity and improving generalization [33].

- **Label Smoothing:** Soft targets were introduced during cross-entropy loss computation to prevent the model from outputting overconfident logit predictions, which enhances calibration [34].
- **Dropout and Batch Normalization:** A dropout rate of 0.5 was applied to the fully connected layers to prevent co-adaptation of features, while batch normalization was utilized to stabilize internal covariate shift [4].
- **Learning Rate Scheduling:** A cosine annealing scheduler was implemented to dynamically decay the learning rate, allowing the optimizer to navigate complex loss landscapes and converge to flatter local minima.
- **Data Augmentation:** Diverse geometric and photometric transformations—including rotation, flipping, Gaussian blurring, and color jittering—were applied to simulate environmental variations encountered during field photography.

9.5 Model Comparison and Architecture Selection

Prior to selecting ResNet50 as the primary backbone, a comparative analysis was conducted using four prominent convolutional neural network architectures trained under identical hyperparameters and optimization settings. The quantitative evaluation metrics are compiled in Table 9.5, and the trade-off between classification accuracy and training computational overhead is illustrated in Figure 9.3.

Table 9.5: Comparing Four Different Neural Network Architectures

Model	Accuracy	Precision	Recall	F1 Score	Training Time
ResNet50	99.57%	99.57%	99.57%	99.57%	16.4 min
DenseNet121	99.57%	99.57%	99.57%	99.57%	15.0 min
GoogLeNet	99.57%	99.57%	99.57%	99.57%	12.2 min
MobileNetV2	99.13%	99.15%	99.13%	99.13%	8.5 min

ResNet50, DenseNet121, and GoogLeNet achieved identical classification accuracies of 99.57%. MobileNetV2 registered a slightly lower accuracy of 99.13% but required significantly less training time. ResNet50 was selected as the final architecture due to several specific advantages.

The residual connections (skip connections) in ResNet50 address the vanishing gradient problem, allowing efficient backpropagation through the 50-layer architecture




Figure 9.3: Visual comparison of how each model performed and how long training took.

and facilitating fine-tuning without information loss [4], [8]. Furthermore, ResNet50 exhibits superior compatibility across deployment frameworks, including PyTorch, TorchScript, ONNX, and edge-optimized runtime engines. This cross-platform compatibility supports future plans to deploy AgriVerify as a native mobile application. The marginal training time overhead (an additional 4.4 minutes compared to GoogLeNet) is offset by these deployment advantages, given that model training is an offline process [15].

9.6 System Performance and Production Latency

In addition to classification accuracy, computational efficiency and response latency are critical parameters for real-time applications deployed in agricultural fields.

9.6.1 Inference Latency Evaluation

Model inference latency refers to the duration required for the network to ingest a preprocessed image, perform a forward pass, and output the class probabilities. The latency profiles across different hardware configurations are summarized in Table 9.6.

On standard CPU architectures (without hardware acceleration), the model achieved an average inference latency of approximately 165 ms. When evaluated on GPU hardware (NVIDIA CUDA), the latency was reduced to 20 ms. Both execution speeds are well within acceptable thresholds for interactive, real-time application feedback.

Table 9.6: Speed on Different Hardware

Hardware	Speed	Good For What?
Regular CPU	~165 ms	Lightweight servers and edge devices; still plenty fast
GPU (NVIDIA CUDA)	~20 ms	Cloud servers handling many requests; 8 times faster than CPU

9.6.2 End-to-End System Performance Under Load

To evaluate the scalability of the backend system under high concurrent access conditions, stress testing was conducted. The system throughput and latency metrics under simulated concurrent loads are detailed in Table 9.7.

Table 9.7: Real-World System Performance Under Load

Metric	Observation	Significance
End-to-End Response	280–420 ms	Includes uploading the image, sending it through the system, running AI, and formatting the response
Concurrent Requests	40–60 per second	Very stable even when many farmers submit seeds simultaneously
Image Upload & Prep	80–120 ms	Network transfer and image pre-processing before AI runs
Database Queries	<50 ms	Fast retrieval of complaint and user records

The system demonstrated high stability and responsiveness under peak concurrent loads. The asynchronous proxy design pattern prevents thread blocking at the core application layer by managing intensive operations outside the main request-response cycle, thereby preserving database and API availability [9].

9.7 Functional Verification and Test Execution

System validation was conducted to verify functional correctness, security enforcement, and data integrity across all primary operational workflows.

9.7.1 Seed Classification Subsystem Validation

The image processing pipeline and classification subsystem were validated against diverse input scenarios. The specific test cases, expected behaviors, and execution results are detailed in Table 9.8.

Table 9.8: Seed Verification Test Cases

Test ID	What We Tested	Expected Behavior	Result
SV-01	Upload clear photos of genuine seed packaging	Correctly classified as Pure with physical measurements extracted	Pass
SV-02	Upload photos of cracked or discolored seeds	Classified as Impure (Damaged) with warnings shown	Pass
SV-03	Upload fake packaging with counterfeit branding	Classified as Impure (Fake) with high confidence	Pass
SV-04	Try uploading wrong file types (.exe, .zip, .txt)	System rejects and shows helpful error message	Pass
SV-05	Upload blurry or very dark seed photos	System warns user to retake photo or shows low confidence	Pass

9.7.2 Grievance Management Subsystem Validation

The complaint lodging and tracking functionalities were validated to verify database consistency and transactional security. The validation test cases are outlined in Table 9.9.

9.7.3 Role-Based Access Control and Security Verification

To ensure secure operations, the authorization and authentication mechanisms were tested against access violations. The results of the security test cases are compiled in Table 9.10.

9.8 System User Interfaces and Operational Workflows

AgriVerify provides simple, intuitive interfaces tailored to three different user types: farmers, agricultural officers, and system administrators. We built these using modern web technologies: Next.js 15, React 19, and Tailwind CSS [13], [31], [35]. Let’s walk through each interface and see how users interact with the system.

Table 9.9: Complaint System Test Cases

Test ID	What We Tested	Expected Behavior	Result
CM-01	Farmer files a complaint with description and evidence	Complaint saved to database and marked as Open	Pass
CM-02	Try filing complaint with missing required information	Form validation prevents submission and shows error	Pass
CM-03	Farmer views their complaint history	All past and current complaints displayed with pagination	Pass
CM-04	Officer changes complaint status from Open to Resolved	Update immediately visible in the system	Pass

Table 9.10: Security and Access Control Test Cases

Test ID	What We Tested	Expected Behavior	Result
RB-01	Valid farmer login	Receives security token and goes to farmer dashboard	Pass
RB-02	Farmer tries accessing admin page	System blocks access and redirects to farmer area	Pass
RB-03	Admin login with multi-factor authentication	Full admin privileges granted	Pass
RB-04	Security token expires during session	System automatically gets new token without user noticing	Pass
RB-05	Login with wrong password	System returns error and shows helpful message	Pass

9.8.1 Public Portal, Onboarding, and Multilingual Support

Upon accessing the AgriVerify platform, users are presented with the public landing page, which outlines the system’s value proposition and operational features (Figure 9.4). To accommodate rural demographics, the platform incorporates a localized language selection interface, enabling users to transition between English and regional dialects (Figure 9.5).




Figure 9.4: Platform landing page displaying value proposition and core features.




Figure 9.5: Multilingual selection interface supporting regional languages.

New users can register by selecting their role (Figure 9.6) and completing the onboarding registration form, which captures demographic, location, and agricul-

tural district details (Figure 9.7). This location data is critical for geo-tagging and regional analytics in the grievance management system.




Figure 9.6: Role selection interface for login.


For administrative personnel and agricultural officers, a secure portal interface enforces multi-factor authentication (MFA) via time-based security codes to prevent unauthorized administrative actions (Figure 9.8).

9.8.2 Farmer Portal and AI Seed Quality Assessment Workflow

Once authenticated, farmers are redirected to a personalized dashboard containing historical verification records, analytical statistics, and regional agricultural alerts (Figure 9.9). Profile parameters and verification badges can be managed within the user profile interface (Figure 9.10).


User details can be modified via the profile configuration interface (Figure 9.11). The seed assessment workflow is initiated by capturing and submitting high-resolution images of the front and back of the seed packaging using the mobile-optimized camera capture interface (Figure 9.12).

Following image processing, the model generates a comprehensive quality report. If the seed batch is certified as pure, the interface displays extracted physical characteristics and purity metrics (Figure 9.13). Conversely, if defects or counterfeit markers are detected, a warning interface highlights classification confidence and offers remedial instructions (Figure 9.14).



The image shows a 'Create Farmer Account' form. At the top is a green logo with a white leaf icon. Below it, the title 'Create Farmer Account' is centered, followed by the subtitle 'Join the AgriVerify community and protect your yield.' The form itself is a white card with a light gray border. It contains several input fields: 'Full Name' (with a person icon and the text 'John Doe'), 'Phone Number' (with a phone icon and the text '9876543210'), 'Email Address' (with an envelope icon and the text 'farmer@example.com'), 'District' (a dropdown menu with a location pin icon and the text 'Select District'), 'Password' (with a lock icon, masked text '*****', and an eye icon), and 'Confirm Password' (with a lock icon, masked text '*****', and an eye icon). At the bottom of the form is a green button with the text 'Create Account' and a right-pointing arrow. Below the button, there is a link that says 'Already have an account? Sign in here'.

Figure 9.7: Farmer registration form collecting demographic and location information.



The image shows a secure admin login interface. At the top center is a shield icon with an exclamation mark. Below it, the text 'ADMIN ACCESS' is in bold, followed by 'RESTRICTED' in green. Underneath that is the text 'SECURE MANAGEMENT CONSOLE'. The login form is a white card with a green border. It contains two input fields: 'ADMIN EMAIL' (with an envelope icon and the text 'admin@fakeseed.com') and 'MASTER PASSWORD' (with a lock icon, the text 'Admin@123456', and an eye icon). Below these fields is a gray button with a white circular arrow icon. At the bottom of the form is a link that says '< RETURN TO MAIN TERMINAL'.

Figure 9.8: Secure admin login requiring multi-factor authentication.




Figure 9.9: Farmer dashboard interface displaying verification statistics and history.




Figure 9.10: Farmer profile configuration interface with verification achievements.




Figure 9.11: User profile modification interface.




Figure 9.12: Mobile camera interface for uploading seed packaging images.




Figure 9.13: Verification result indicating pure seed batch with morphological measurements.




Figure 9.14: Verification alert for counterfeit or low-quality seed batch.

9.8.3 Grievance Redressal and Officer Investigation Interfaces

In instances of quality failure or suspected counterfeit distribution, farmers can lodge formal complaints detailing the batch source and severity (Figure 9.15). The status and resolution progress of these submissions can be monitored via the complaint queue interface (Figure 9.16).




Figure 9.15: Grievance registration interface for lodging seed quality complaints.




Figure 9.16: Farmer tracking interface for monitoring grievance status.

Agricultural officers access a dedicated monitoring dashboard summarizing regional

complaints and workload distribution metrics (Figure 9.17). Officers can query, filter, and prioritize the investigation queue based on severity, date, and location (Figure 9.18). Selecting a specific complaint displays the detailed investigation pane, allowing officers to examine user-submitted image evidence and update the resolution status (Figure 9.19).




Figure 9.17: Agricultural officer dashboard presenting regional complaint metrics.




Figure 9.18: Filtered and prioritized grievance investigation queue for officers.

COMPLAINT OPERATIONS						VIEW FULL QUEUE >
All statuses ▾		Search by batch, description, district...				
Batch	Issue	District	Severity	Status	Created	
1234567	cchnnb	imphal west	5	Closed	9/5/2026	Open Investigating Resolved Closed ✓
123456	poor quality	imphal west	5	Closed	9/5/2026	CLOSED ▾
12345	12345	bishnupur	5	Closed	8/5/2026	CLOSED ▾
12345	dwfsdghjkkjhgfdsa	N/A	3	Closed	4/5/2026	CLOSED ▾
12345	N/A	N/A	3	Resolved	4/5/2026	RESOLVED ▾

Figure 9.19: Officer grievance detail pane featuring user evidence and resolution controls.

9.8.4 Platform Administration, User Governance, and Telemetry

Platform administrators monitor system telemetry, API response times, active sessions, and database load via the real-time administrator dashboard (Figure 9.20). Administrative user governance, including account creation and role modification, is managed through the user governance portal (Figure 9.21).




Figure 9.20: Administrator real-time system performance telemetry dashboard.

System analytics are visualized through reporting dashboards that highlight verification trends and fake seed detection rates over time (Figure 9.22). Geographical visualization dashboards map the distribution of counterfeit seed incidents across different administrative districts, aiding policymakers in identifying counterfeit




Figure 9.21: User administration and access control interface.

hotspots (Figure 9.23).

9.9 Challenges Encountered and Mitigation Strategies

The development and deployment of AgriVerify encountered several technical challenges, which were addressed using specific mitigation strategies:

- Data Scarcity:** The training dataset comprised only 1,563 images, introducing a high risk of model overfitting. This constraint was mitigated by utilizing transfer learning with pre-trained ImageNet weights, applying data augmentation to simulate sample diversity, and employing weight decay to regularize model complexity [33], [36].
- Fine-Grained Feature Similarity:** Discriminating between mechanically damaged seeds and high-quality counterfeit packaging required high classification sensitivity. The system mitigated this by combining visual texture analysis with Optical Character Recognition (OCR) to extract text features from the packaging labels, providing multi-modal confirmation [32].
- Environmental Illumination Variations:** Field photographs taken by farmers under uncontrolled ambient lighting conditions could degrade classification accuracy. To resolve this, automatic preprocessing algorithms adjust image contrast and brightness before inference, and training images were augmented with random color jittering to enhance model robustness.




Figure 9.22: Analytical reporting dashboard showing quality trends and system statistics.




Figure 9.23: District-wise geospatial hot-spot mapping of counterfeit occurrences.

- **Resource Constraints and Model Scale:** Deploying a standard ResNet50 model (approximately 100 MB) on edge servers introduces latency and memory bottlenecks. This was mitigated by implementing half-precision floating-point quantization (FP16) and model optimization, reducing memory footprints without compromising classification accuracy [15].

9.10 Summary and System Limitations

This chapter presented a multi-dimensional evaluation of the AgriVerify platform, spanning model classification performance, system latency under concurrent loads, and functional validation. The ResNet50 classifier achieved an accuracy of 99.57% and an F1 Score of 99.57%, demonstrating performance in distinguishing certified seeds from damaged or counterfeit counterparts. Functional validation verified that the platform’s core workflows operate correctly and securely.

Despite these positive results, several system limitations remain to be addressed in future work:

- **Single-Crop Dependency:** The classifier is currently optimized exclusively for paddy seeds. Extending support to other major crops, such as wheat, maize, or soybeans, would require compiling specialized image datasets and retraining the network [8].
- **Network Connectivity Requirements:** The system relies on cloud-based API serving, restricting use in remote agricultural areas with poor internet connectivity. Quantizing models for offline, local execution on mobile devices represents a critical future development path.
- **Lack of Diagnostic Specificity:** While the system detects seed impurity, it does not identify specific seed-borne diseases or pathogens. Integrating phytopathological diagnostic capabilities would enhance agricultural value [32].

Addressing these limitations—through edge optimization, multi-crop expansion, and explainable AI (XAI) features to visualize classification attention maps—constitutes the primary roadmap for future research.

Chapter 10

Conclusion And Future Work

10.1 Summary of Work

AgriVerify was built to address a specific, practical problem: farmers in rural regions have no reliable way to verify paddy seed quality before planting. Existing methods manual inspection and laboratory testing are slow, inconsistent and largely out of reach for smallholder farmers. The system responds to that gap by combining deep learning classification, OCR-based packaging analysis, cloud deployment, and a farmer complaint workflow into a single platform.

The architecture reflects those requirements. A Next.js frontend delivers role-based dashboards for farmers, officers, administrators, and public users. An ASP.NET Core backend built on Clean Architecture principles handles API orchestration and business logic. A FastAPI microservice runs the machine learning inference. Azure AI Cognitive Services handles OCR and report synthesis.

The seed classification model uses ResNet50 transfer learning, trained on a curated dataset of healthy, damaged, and fake paddy seed images. Preprocessing and augmentation rotation, normalisation, brightness adjustment, cropping, flipping were applied to improve generalisation. The model extracts visual features including shape, texture, colour distribution, and grain structure to classify seed quality.

OCR through Azure Computer Vision extracts batch numbers, expiry dates, and manufacturer details from seed packaging. Combined with the classification output, Azure OpenAI synthesises these into plain-language summaries that non-technical users can act on without needing to interpret raw model scores.

Farmers can also submit complaints about suspicious or failed seed products directly through the platform. Agricultural officers investigate, record actions, update complaint status, and can trigger batch embargoes where warranted. This

turns what would otherwise be an isolated prediction system into a traceable quality monitoring workflow.

The ML API was containerised with Docker and deployed on scalable cloud infrastructure. The frontend and backend were designed for low-latency operation and remain usable under limited connectivity conditions common in rural agricultural regions.

10.2 Key Contributions and Findings

10.2.1 AI-Based Automated Paddy Seed Classification

The classification model distinguishes between healthy, damaged, and fake paddy seeds with high accuracy and low inference latency. Transfer learning from ResNet50 reduced the volume of labelled training data required while maintaining strong prediction performance, making real-time deployment viable.

10.2.2 Integration of Multiple AI Services

The project combined several AI capabilities into one platform:

- Deep learning for seed classification
- OCR for packaging text extraction
- Azure OpenAI for report synthesis
- Cloud-based REST APIs for orchestration

Each component addresses a distinct part of the verification problem. Together they produce outputs classification, provenance data, plain-language summary that a farmer can act on without technical knowledge.

10.2.3 DUS Parameter Evaluation

The system extracts and measures physical seed characteristics length, width, colour variance, circularity, surface texture and compares them against reference values for known cultivars. This moves the system beyond binary classification toward the kind of structured evaluation used in scientific seed testing, specifically the Distinctness, Uniformity, and Stability (DUS) framework.

10.2.4 Real-Time Farmer Assistance

Farmers upload a seed image and receive a classification result, confidence score, and advisory summary, typically within seconds. This replaces a process that previously required sending samples to a laboratory and waiting days for results, a wait that often extends past the planting window.

10.2.5 Secure and Scalable System Architecture

JWT authentication, role-based access control, repository patterns, and parallel AI task execution address both security and performance requirements. The cloud-ready architecture supports scaling to larger user bases without structural changes to the application.

10.2.6 Complaint Redressal and Governance Workflow

The complaint management system gives farmers a direct channel to report suspect products and gives officers the tools to investigate, document, and act. This creates an audit trail across the supply chain that manual reporting systems cannot provide.

10.2.7 Reduction of Manual Dependency

The model identifies defects, damaged grains, and fake seed varieties more consistently than manual inspection, which varies by observer experience and lighting conditions. Automating this step removes a significant source of variability from the verification process.

10.2.8 Cloud-Native AI Deployment

Docker containerisation, REST APIs, and scalable microservices make the system straightforward to update, monitor, and redeploy. Resource usage scales with actual query volume rather than reserved capacity.

10.3 Future Enhancements

10.3.1 Offline-First Progressive Web Application (PWA)

Intermittent connectivity is the norm in many of the regions this system targets. Implementing offline-first capabilities using PWA technologies and IndexedDB would allow farmers to scan seeds and queue submissions without an active connection, syncing when signal is available.

10.3.2 Enhanced Security Mechanisms

The current implementation stores JWT tokens in `localStorage`. Migrating to `HttpOnly` cookies with Content Security Policy enforcement would reduce exposure to XSS attacks and session hijacking, a worthwhile hardening step before any large-scale rollout.

10.3.3 Blockchain-Based Verification Ledger

Each verification report could be cryptographically hashed and written to a distributed ledger, producing a tamper-evident record of seed batch history across the supply chain. This would give regulators and buyers an independently verifiable audit trail.

10.3.4 IoT and Smart Agriculture Integration

Connecting the platform to soil sensors, humidity monitors, and fertiliser quality detectors would let the system combine seed verification data with environmental context, producing more specific agronomic recommendations than seed quality data alone can support.

10.3.5 Expansion to Multiple Crop Varieties

The current model and dataset cover paddy seeds. Extending training to wheat, maize, pulses, and vegetables would broaden the platform's utility considerably, though each crop would require its own labelled dataset and likely some model-specific tuning.

10.3.6 Mobile Application Development

A dedicated Android or iOS application would improve camera integration, offline caching, and overall usability for farmers who are more comfortable with native apps than browser interfaces.

10.3.7 Multilingual Farmer Assistance

English-only interfaces limit adoption in linguistically diverse agricultural regions. Regional language support and voice-based interaction would lower the literacy barrier and widen the platform’s reach.

10.3.8 Advanced Analytics Dashboard

Aggregated verification data could support trend analysis: counterfeit seed distribution patterns, regional fraud hotspots, seasonal quality variation. This would give agricultural officers and policymakers a clearer picture of where problems are concentrated.

10.3.9 Federated Learning for Privacy-Preserving AI

Training models collaboratively across decentralised devices, without pooling raw image data centrally, would address privacy concerns while allowing the model to improve from field data it would not otherwise have access to.

10.3.10 Explainable AI (XAI)

Grad-CAM visualisations or feature heatmaps would show users which parts of a seed image drove the classification decision. For agricultural officers in particular, understanding why a seed was flagged as fake — not just that it was — would make the system’s outputs easier to act on and defend.

10.4 Conclusion

AgriVerify demonstrates that deep learning classification, OCR-based packaging analysis, DUS parameter evaluation, and AI-generated summaries can together replace manual seed verification with something faster, more consistent, and more

accessible. Farmers get a result in seconds rather than days. Officers get a complaint management workflow with a traceable audit trail. The supply chain gets a monitoring layer it previously lacked.

The architecture is containerised and modular enough to extend — to new crops, new languages, new AI capabilities — without rebuilding from scratch. Offline support, explainability, and multilingual access remain the most significant gaps between the current prototype and something deployable at scale.

None of the individual components here — transfer learning, OCR, LLM summaries, complaint workflows — are new. The contribution is in how they are assembled: around the specific constraints of rural agricultural users, producing a coherent verification system that none of them could form on their own.

REFERENCES

- [1] S. P. Mohanty, D. P. Hughes, and M. Salathé, “Using deep learning for image-based plant disease detection,” *Frontiers in Plant Science*, vol. 7, p. 1419, 2016. DOI: 10.3389/fpls.2016.01419.
- [2] M. Dyrmann, H. Karstoft, and H. S. Midtiby, “Plant species classification using deep convolutional neural network,” *Biosystems Engineering*, vol. 151, pp. 72–80, 2016. DOI: 10.1016/j.biosystemseng.2016.08.024.
- [3] Y. Zhang, S.-j. Wang, G. Ji, and P. Phillips, “Fruit classification using computer vision and feedforward neural network,” *Journal of Food Engineering*, vol. 243, pp. 123–135, 2020.
- [4] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2016, pp. 770–778. DOI: 10.1109/CVPR.2016.90.
- [5] P. Muñoz, M. Navas, *et al.*, “ICT and mobile applications for agricultural extension in low-resource settings: A systematic review,” *Computers and Electronics in Agriculture*, vol. 154, pp. 234–245, 2018.
- [6] C. Ni, W. Fang, Y. Cheng, *et al.*, “Classification of rice seed varieties using morphological parameters extracted from binary silhouettes,” *Transactions of the ASABE*, vol. 52, no. 3, pp. 951–957, 2009.
- [7] J. Liu, Q. Wang, and H. Li, “Application of VGG16 transfer learning to maize seed defect detection,” *Computers and Electronics in Agriculture*, vol. 175, p. 105 566, 2020.
- [8] P. Xu, X. Chen, and L. Zhang, “ResNet-50 for soybean quality grading and fine-tuning evaluation,” *Agronomy Journal*, vol. 113, no. 4, pp. 3210–3221, 2021.
- [9] R. C. Martin, *Clean Architecture: A Craftsman’s Guide to Software Structure and Design*. Upper Saddle River, NJ: Prentice Hall, 2018, ISBN: 978-0134494166.

- [10] Vercel Inc., “Next.js 14 app router documentation,” Vercel, Tech. Rep., 2023. [Online]. Available: <https://nextjs.org/docs>.
- [11] TanStack, *TanStack Query (React Query) v5 documentation*, <https://tanstack.com/query/latest>, Accessed: 2024, 2023.
- [12] D. Kato, *Zustand: A small, fast, and scalable state management solution*, <https://github.com/pmndrs/zustand>, Accessed: 2024, 2023.
- [13] Tailwind Labs, *Tailwind CSS framework documentation*, 2024. [Online]. Available: <https://tailwindcss.com/docs>.
- [14] Microsoft Corporation, “ASP.NET Core 8 documentation,” Microsoft, Tech. Rep., 2023. [Online]. Available: <https://learn.microsoft.com/en-us/aspnet/core/>.
- [15] A. Paszke, S. Gross, F. Massa, *et al.*, “PyTorch: An imperative style, high-performance deep learning library,” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, pp. 8026–8037, 2019. [Online]. Available: <https://proceedings.neurips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html>.
- [16] G. Bradski, “The OpenCV library,” *Dr. Dobb’s Journal of Software Tools*, vol. 25, no. 11, pp. 120–125, 2000. [Online]. Available: <https://opencv.org>.
- [17] S. Ramírez, “FastAPI documentation,” Tiangolo, Tech. Rep., 2023. [Online]. Available: <https://fastapi.tiangolo.com>.
- [18] UPOV, “General introduction to the examination of distinctness, uniformity and stability and the development of harmonized descriptions of new varieties of plants,” International Union for the Protection of New Varieties of Plants (UPOV), Tech. Rep. TG/1/3, 2002. [Online]. Available: <https://www.upov.int/edocs/tgdocs/en/tg001.pdf>.
- [19] OWASP Foundation, “OWASP Top Ten 2021: The ten most critical web application security risks,” Open Web Application Security Project, Tech. Rep., 2021. [Online]. Available: <https://owasp.org/Top10/>.
- [20] R. C. Martin, *Clean Architecture: A Craftsman’s Guide to Software Structure and Design*. Boston, MA: Prentice Hall, 2017.
- [21] S. Ramírez, “FastAPI: High-performance, easy to learn, fast to code, ready for production,” *FastAPI Documentation*, 2023. [Online]. Available: <https://fastapi.tiangolo.com>.
- [22] D. Merkel, “Docker: Lightweight linux containers for consistent development and deployment,” *Linux Journal*, vol. 2014, no. 239, p. 2, 2014.

- [23] A. Paszke, S. Gross, F. Massa, *et al.*, “PyTorch: An imperative style, high-performance deep learning library,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019, pp. 8024–8035.
- [24] G. Bradski, “The OpenCV library,” *Dr. Dobb’s Journal of Software Tools*, vol. 25, no. 11, pp. 120–125, 2000.
- [25] S. Biswas *et al.*, “Serverless containerized microservices deployment using Google Cloud Run,” Google Cloud Infrastructure Technical Reports, Tech. Rep., 2022. [Online]. Available: <https://cloud.google.com/run/docs>.
- [26] M. Salim *et al.*, “Microservices-based system for automated data collection, sentiment analysis, and visualization in video game reviews across multiple platforms,” in *IEEE Conference Publication*, IEEE, 2024.
- [27] S. S. Singh, N. K. Mishra, and S. A. Singh, “Impact of online communication platforms on knowledge and adoption behaviour of farmers in eastern uttar pradesh: A study on major crops,” *Journal of Experimental Agriculture International*, 2026.
- [28] Meta Platforms, Inc., *React documentation: Concurrent rendering*, 2024. [Online]. Available: <https://react.dev/>.
- [29] TanStack, *TanStack Query: Asynchronous state management*, 2024. [Online]. Available: <https://tanstack.com/query/latest>.
- [30] W. G. Masue, D. Ngondya, and T. S. Kondo, “Quantifying vulnerabilities: A systematic review of the state-of-the-art web-based systems,” *Journal of ICT Systems*, vol. 2, no. 1, pp. 72–86, 2024. DOI: 10.56279/jicts.v2i1.51.
- [31] Vercel Inc., *Next.js documentation: App router & proxy middleware*, 2024. [Online]. Available: <https://nextjs.org/docs>.
- [32] S. P. Mohanty, D. P. Hughes, and M. Salathé, “Using deep learning for image-based plant disease detection,” *Frontiers in Plant Science*, vol. 7, p. 1419, 2016. DOI: 10.3389/fpls.2016.01419.
- [33] I. Loshchilov and F. Hutter, “Decoupled weight decay regularization,” *International Conference on Learning Representations (ICLR)*, 2019. [Online]. Available: <https://openreview.net/forum?id=Bkg6RiCqY7>.
- [34] R. Müller, S. Kornblith, and G. E. Hinton, “When does label smoothing help?” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019. [Online]. Available: <https://proceedings.neurips.cc/paper/2019/hash/f1748d6b0fd9d439f71450117eba2725-Abstract.html>.

- [35] J. Walke, “React: A JavaScript library for building user interfaces,” *Meta Open Source*, 2013. [Online]. Available: <https://reactjs.org>.
- [36] A. Krizhevsky, I. Sutskever, and G. E. Hinton, “ImageNet classification with deep convolutional neural networks,” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 25, pp. 1097–1105, 2012. [Online]. Available: <https://proceedings.neurips.cc/paper/2012/hash/c399862d3b9d6b76c8436e924a68c45b-Abstract.html>.