

AI-Based Paddy Seed Quality Assessment And Farmer Feedback System

*A thesis submitted in partial fulfilment of the requirements
for the award of the degree of*

Bachelor of Technology

By

Dayananda Laishram (2201CS0113)

Khangembam Yaiphaba Singh (2201CS0102)

Joymangol Chingangbam (2201CS0004)

Nilambar Elangbam (2201CS0005)

Under the Guidance of

Dr. Roshni Rajkumari

Assistant Professor



**Department of
Computer Science & Engineering**

Manipur Technical University

June, 2026

DEDICATION

This thesis is dedicated to all those who have inspired, supported and stood by us throughout the journey of this project.

To our families: Your unwavering love, patience, and encouragement have been the bedrock of our emotional strength during the most challenging phases. *Your belief in our potential has fortified our resilience and determination.*

To our mentors and instructors: Your expertise, thoughtful guidance, and commitment to our academic growth have been instrumental in shaping both this project and our understanding of the field. *Your encouragement and constructive feedback have been invaluable at each stage.*

To our friends and peers: Your camaraderie, moral support, and motivation have illuminated our path. *Your presence transformed moments of stress into opportunities for shared growth and learning.*

To the global community of researchers, engineers, and innovators in Artificial Intelligence and Computer Vision: Your pioneering work has laid the foundation for this thesis. *Your dedication to advancing knowledge continues to inspire and propel us to contribute meaningfully to this dynamic field.*

This thesis stands as a testament to the collective support, shared vision, and unwavering encouragement of all who have been part of our academic and personal journey.

AI-Based Paddy Seed Quality Assessment And Farmer Feedback System

*A thesis submitted in partial fulfilment of the requirements
for the award of the degree of*

Bachelor of Technology

By

Dayananda Laishram (2201CS0113)

Khangembam Yaiphaba Singh (2201CS0102)

Joymangol Chingangbam (2201CS0004)

Nilambar Elangbam (2201CS0005)

Under the Guidance of

Dr. Roshni Rajkumari

Assistant Professor



**Department of
Computer Science & Engineering**

Manipur Technical University

June, 2026



MANIPUR TECHNICAL UNIVERSITY

(A University established under the

Manipur Technical University Act, 2016)

Imphal, Manipur – 795004

BONAFIDE CERTIFICATE

This is to certify that the thesis entitled “**AI-Based Paddy Seed Quality Assessment and Farmer Feedback System**” submitted by:

Khangembam Yaiphaba Singh (2201CS0102)

Nilambar Elangbam (2201CS0005)

Dayananda Laishram (2201CS0113)

Joymangol Chingangbam (2201CS0004)

to the **Department of Computer Science & Engineering**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree in **Computer Science and Engineering**, is a genuine and original work carried out by them under my supervision during the academic year **2025–2026**.

The project embodies the results of their own research and development work and has not been submitted, either in part or in full, for the award of any other degree or diploma of this university or any other institution. The work has been carried out with sincerity, academic integrity, and adherence to research ethics.

We consider this work worthy of consideration for the partial fulfillment of the requirements for the said degree.

Supervisor

Dr. Roshni Rajkumari
Assistant Professor,
Computer Science & Engineering

Head of Department

Ma'am Tayenjam Aarena
Assistant Professor & Head,
Computer Science & Engineering

DECLARATION

I hereby declare that the thesis entitled “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**”, submitted to the **Department of Computer Science and Engineering, Manipur Technical University**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree, is a result of my own work and effort.

I affirm that this work has been carried out with academic honesty and integrity. All sources of information have been properly cited and acknowledged. This thesis has not been submitted earlier, either in part or in full, for the award of any degree or diploma at this or any other institution.

Date:

Place:

Nilambar Elangbam

Reg. No. - 2201CS0005

Computer Science & Engineering

DECLARATION

I hereby declare that the thesis entitled “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**”, submitted to the **Department of Computer Science and Engineering, Manipur Technical University**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree, is a result of my own work and effort.

I affirm that this work has been carried out with academic honesty and integrity. All sources of information have been properly cited and acknowledged. This thesis has not been submitted earlier, either in part or in full, for the award of any degree or diploma at this or any other institution.

Date:

Place:

Dayananda Laishram

Reg. No. - 2201CS0113

Computer Science & Engineering

DECLARATION

I hereby declare that the thesis entitled “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**”, submitted to the **Department of Computer Science and Engineering, Manipur Technical University**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree, is a result of my own work and effort.

I affirm that this work has been carried out with academic honesty and integrity. All sources of information have been properly cited and acknowledged. This thesis has not been submitted earlier, either in part or in full, for the award of any degree or diploma at this or any other institution.

Date:

Place:

Joymangol Chingangbam

Reg. No. - 2201CS0004

Computer Science & Engineering

DECLARATION

I hereby declare that the thesis entitled “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**”, submitted to the **Department of Computer Science and Engineering, Manipur Technical University**, in partial fulfilment of the requirements for the award of the **Bachelor of Technology** degree, is a result of my own work and effort.

I affirm that this work has been carried out with academic honesty and integrity. All sources of information have been properly cited and acknowledged. This thesis has not been submitted earlier, either in part or in full, for the award of any degree or diploma at this or any other institution.

Date:

Place:

Khangembam Yaiphaba Singh

Reg. No. - 2201CS0102

Computer Science & Engineering

ACKNOWLEDGEMENT

We sincerely acknowledge the collective support and guidance received throughout the successful completion of our thesis project “**AI-Based Paddy Seed Quality Assessment And Farmer Feedback System**” We are deeply grateful to the Department of Computer Science and Engineering, Manipur Technical University, for providing the academic environment, infrastructure and encouragement that fostered innovation and technical exploration. Our heartfelt thanks go to our project supervisor, **Dr. Roshni Rajkumari**, Assistant Professor, whose consistent guidance, insightful feedback and unwavering support played a pivotal role in shaping this project. We also extend our sincere appreciation to **Ma’am Tayenjam Aarena**, Head of Department, for his inspiring leadership and motivation.

We are equally thankful to all faculty members, technical staff and friends who provided their valuable input and moral support throughout the journey. A special note of gratitude goes to our families for their emotional strength, encouragement and faith in us during times of challenge. Above all, we are deeply thankful to the Almighty for granting us the perseverance and resilience needed to see this project through. This thesis stands as a reflection of the collective wisdom, collaboration and goodwill of everyone who contributed to its realization.

Date:

Place:

Sincerely,

Joymangol Chingangbam

Nilambar Elangbam

Dayananda Laishram

Khagembam Yaiphaba

ABSTRACT

AI-Based Paddy Seed Quality Assessment and Farmer Feedback System is a web platform that addresses counterfeit and substandard paddy seeds which is a recurring cause of crop failure and income loss for farmers across northeastern states of India. Farmers can upload or photograph seed samples through the interface; the system analyzes visual characteristics using CNN-based models running on cloud inference and returns a quality summary before sowing decisions are made, with no lab equipment required.

The platform also includes a Farmer Feedback and Complaint Module, where users can report poor germination, suspected counterfeits, or crop failures. Officers access an analytics dashboard that aggregates these reports to surface risk patterns at the batch and regional level.

Table of Contents

BONAFIDE CERTIFICATE	i
DECLARATION	ii
ACKNOWLEDGEMENT	vi
ABSTRACT	vii
List of Tables	xvii
List of Figures	xix
Symbols and Acronyms	xxi
1 Introduction	1
1.1 Background and Context	1
1.2 Motivation	1
1.3 Problem Statement	2
1.4 Objectives	2
1.5 Scope of the Project	3
2 Literature Review	4
2.1 Seed Quality Assessment: Traditional Methods	4
2.2 Deep Learning for Agricultural Image Classification	4
2.3 Transfer Learning and ResNet Architectures	5
2.4 Farmer Feedback Systems	5

2.5	Web-Based Agricultural Platforms	6
2.6	Review of Related Work	6
2.7	Summary and Research Gaps	7
3	System Design and Architecture	8
3.1	Overall System Architecture	8
3.1.1	Backend Architectural Layers	9
3.2	Functional Requirements	10
3.3	Non-Functional Requirements	12
3.4	Tools and Technologies	14
3.4.1	Deep Learning Stack (PyTorch and ResNet-50)	14
3.4.2	Backend Stack (ASP.NET Core 8, SQL Server, and Azure) . . .	15
3.4.3	Model Serving (FastAPI and Docker)	16
3.4.4	Frontend Stack (Next.js, TypeScript, and TanStack Query) . . .	17
3.5	System Architecture Diagrams	18
3.5.1	High-Level Architecture	18
3.5.2	Data Flow Diagram	19
3.5.3	Entity-Relationship Diagram	19
4	Deep Learning Model for Paddy Seed Quality Assessment	22
4.1	Dataset Description	22
4.1.1	Data Sources and Collection (Kaggle)	22
4.1.1.1	Public Source	22
4.1.1.2	Custom Source	23
4.1.2	Dataset Statistics and Class Distribution	23
4.1.3	Train/Validation/Test Split Strategy	24
4.1.4	Data Preprocessing and Augmentation	24
4.2	Model Architecture	25
4.2.1	ResNet50 Backbone (25.6M Parameters)	25
4.2.2	Custom Classification Head	25

4.2.3	Anti-Overfitting Techniques	26
4.3	DUS Parameter Integration	27
4.3.1	Feature Extraction—17 Physical Parameters	27
4.3.2	Pixel-to-Physical Unit Conversion	27
4.3.3	Variety Matching (12 Reference Varieties)	28
4.4	Training Pipeline	28
4.4.1	Phase 1—Classification Head Training (Epochs 1–8)	28
4.4.2	Phase 2—Full Fine-Tuning (Epochs 9–15)	28
4.4.3	Hyperparameter Configuration and Optimiser	29
4.4.4	Model Export Formats	29
4.5	Model Deployment	30
4.5.1	Serving Infrastructure	30
4.5.2	Containerisation and CI/CD Pipeline	30
4.6	Model Comparison and Selection	31
4.7	Final Model Performance Summary	32
5	Model Serving API	33
5.1	API Architecture and Request Orchestration	33
5.1.1	FastAPI and Uvicorn Serving Stack	33
5.1.2	End-to-End Model Request Flow	34
5.1.3	Concurrency and Performance Optimization	34
5.2	Inference Execution and Preprocessing Pipeline	35
5.2.1	Image Preprocessing and Tensor Normalization	35
5.2.2	Prediction, Confidence Scoring, and Latency Measurement . . .	36
5.3	API Specification and Endpoint Contracts	37
5.4	Containerisation and Deployment Architecture	39
5.4.1	Docker Multi-Stage Build Architecture	39
5.4.2	Serverless Deployment on GCP Cloud Run	40
5.5	Configuration, Observability, and Security Hardening	41
5.5.1	Environment Management and Secret Injection	41

5.5.2	Threat Model and Defensive Hardening	41
5.6	Summary	42
6	Frontend System Architecture and Implementation	43
6.1	Architecture and Design Principles	43
6.2	Technology Stack	44
6.3	Folder Structure and Repository Layout	47
6.4	Routing System and Access Control	48
6.4.1	App Router Namespace and Access Hierarchy	48
6.4.2	Role-Based Access Control Implementation	50
6.4.3	User Role Workflows and Portal Functionality	51
6.4.3.1	Farmer Portal — Verification and Complaint Workflows	51
6.4.3.2	Officer Portal — Regional Investigation Queue	51
6.4.3.3	Admin Portal — System Governance and Analytics .	52
6.4.3.4	Public Verification — Anonymous Field Scanner . . .	52
6.5	Authentication Flow and Session Management	52
6.5.1	Token Storage Strategy and Zustand Auth Store	52
6.5.2	Automated Refresh Token Recovery (Axios 401 Interceptor) . .	53
6.6	API Integration and Client Layer	54
6.6.1	Centralized Axios HTTP Client and Server Proxy Route	54
6.6.2	API Endpoint Integration Summary	55
6.7	State Management and UI Design System	60
6.7.1	Asynchronous Server State Management	60
6.7.2	UI Component Library and Design Tokens	60
6.7.3	Responsive Layouts and Accessibility	62
6.8	Summary	62
7	Backend System Architecture and Implementation	63
7.1	Architectural Overview	63
7.1.1	Architectural Layers	63

7.1.1.1	Domain Layer	63
7.1.1.2	Application Layer	64
7.1.1.3	Infrastructure Layer	64
7.1.1.4	Presentation Layer (API Layer)	64
7.1.2	Architectural Diagram	64
7.2	Technology Stack	65
7.3	Project Structure Analysis	66
7.3.1	Domain Layer Architecture	66
7.3.1.1	Entities Folder — Core Business Models	67
7.3.1.2	Enums Folder — System Classifications	67
7.3.1.3	Interfaces Folder — Service Contracts	68
7.3.1.4	Common Folder — Shared Base Classes	68
7.3.1.5	Exceptions Folder — Error Handling Strategy	69
7.3.2	Infrastructure Layer Architecture	69
7.3.2.1	Persistence Folder — Data Management	69
7.3.2.1.1	ApplicationDbContext.cs	69
7.3.2.1.2	UnitOfWork.cs	69
7.3.2.1.3	Repository Implementation	70
7.3.2.2	Configurations Folder — Database Mapping	70
7.3.2.3	Migrations Folder — Schema Version Control	70
7.3.2.4	Services Folder — External Integrations	71
7.3.2.5	Extensions Folder — Dependency Injection	71
7.3.3	Application Layer Architecture	72
7.3.3.1	Services Subfolder	72
7.3.3.1.1	Service Interfaces	72
7.3.3.1.2	Service Implementations	72
7.3.3.2	DTOs Subfolder — Data Transfer Objects	73
7.3.3.2.1	Request DTOs	73
7.3.3.2.2	Response DTOs	74

7.3.3.3	Validators Subfolder	74
7.3.3.4	Mappings Subfolder	75
7.3.3.5	Exceptions Subfolder	75
7.3.4	Presentation Layer (API) Architecture	75
7.3.4.1	Controllers Subfolder	75
7.3.4.1.1	BaseApiController.cs	75
7.3.4.1.2	AuthController.cs	76
7.3.4.1.3	VerificationController.cs	76
7.3.4.1.4	ComplaintsController.cs	76
7.3.4.1.5	AdminController.cs	76
7.3.4.1.6	AnalyticsController.cs	77
7.3.4.1.7	AnalyticsControllerV2.cs	77
7.3.4.1.8	ReferenceDataController.cs	77
7.3.4.2	Middleware Subfolder	78
7.3.4.2.1	ExceptionHandlerMiddleware.cs	78
7.3.4.2.2	RequestLoggingMiddleware.cs	78
7.3.4.2.3	AuthenticationMiddleware.cs	78
7.3.4.3	Configuration Files	79
7.3.4.3.1	appsettings.json	79
7.3.4.3.2	appsettings.Development.json	79
7.3.4.3.3	appsettings.Production.json	79
7.3.4.4	Program.cs — Application Entry Point	80
7.4	Data Flow and Integration Points	80
7.4.1	Verification Workflow	80
7.4.2	Security Architecture	81
8	System Integration	82
8.1	End-to-End Request Flow (Browser to ML Model)	82
8.1.1	Overview of the Request Lifecycle	83
8.1.2	Browser-Level Request Generation	83

8.1.3	Backend Request Reception	84
8.1.4	AI Inference Execution Pipeline	85
8.1.5	Request Flow Diagram	86
8.2	Frontend - Backend Communication (Proxy Architecture)	86
8.2.1	Centralized Communication Architecture	87
8.2.2	Axios Client Integration	87
8.2.3	Proxy Route Workflow	88
8.2.4	Authentication and Session Recovery	88
8.2.5	Proxy Architecture Diagram	89
8.3	Backend - Model Serving API Integration	89
8.3.1	AI Service Integration Architecture	89
8.3.2	Image Validation Pipeline	89
8.3.3	OCR and Metadata Extraction	90
8.3.4	Seed Classification Workflow	90
8.3.5	Model Serving Workflow Diagram	92
8.4	Azure Service Integration Workflow	92
8.4.1	Azure Services Used in the Platform	93
8.4.2	Azure Blob Storage Workflow	93
8.4.3	Azure AI Service Coordination	94
8.4.4	Cloud-Based AI Processing Pipeline	94
8.4.5	Azure Integration Diagram	95
8.5	Deployment Architecture (Docker, Azure Container Apps)	95
8.5.1	Containerization Strategy	95
8.5.2	Docker-Based Build Workflow	96
8.5.3	Azure Container Registry Integration	96
8.5.4	Azure Container Apps Deployment	97
8.5.5	Deployment Workflow Diagram	98
9	Results and Evaluation	99
9.1	Model Performance	99

9.1.1	Accuracy, Precision, Recall, and F1 Score	99
9.2	Per-Class Performance	100
9.3	Confusion Matrix Analysis	100
9.4	Regularization Effectiveness	101
9.5	Model Comparison and Analysis	102
9.6	System Performance	103
9.6.1	API Inference Latency (165 ms CPU / 20 ms GPU)	103
9.6.2	Backend Response Times and Throughput	104
9.7	Functional Test Cases	104
9.7.1	Seed Verification Test Cases	105
9.7.2	Complaint Management Test Cases	105
9.7.3	Role-Based Access and Admin Test Cases	106
9.7.4	Output Samples and Screenshots	107
9.7.5	Challenges Encountered	107
9.7.6	Known Limitations of the Current System	108
10	Conclusion And Future Work	109
10.1	Summary of Work	109
10.2	Key Contributions and Findings	110
10.2.1	AI-Based Automated Paddy Seed Classification	110
10.2.2	Integration of Multiple AI Services	111
10.2.3	DUS Parameter Evaluation	111
10.2.4	Real-Time Farmer Assistance	111
10.2.5	Secure and Scalable System Architecture	111
10.2.6	Complaint Redressal and Governance Workflow	112
10.2.7	Reduction of Manual Dependency	112
10.2.8	Cloud-Native AI Deployment	112
10.3	Future Enhancements	112
10.3.1	Offline-First Progressive Web Application (PWA)	112
10.3.2	Enhanced Security Mechanisms	113

10.3.3 Blockchain-Based Verification Ledger	113
10.3.4 IoT and Smart Agriculture Integration	113
10.3.5 Expansion to Multiple Crop Varieties	113
10.3.6 Mobile Application Development	113
10.3.7 Multilingual Farmer Assistance	113
10.3.8 Advanced Analytics Dashboard	114
10.3.9 Federated Learning for Privacy-Preserving AI	114
10.3.10 Explainable AI (XAI)	114
10.4 Conclusion	114

References	116
-------------------	------------

List of Tables

3.1	Backend Architectural Layers and Responsibilities	9
3.2	Comprehensive Functional Requirements Specification	10
3.3	Comprehensive Non-Functional Requirements Specification	12
3.4	Backend Technology Stack	16
4.1	Dataset Source Summary	23
4.2	Class-wise Distribution of the Combined Dataset	23
4.3	Dataset Split Statistics	24
4.4	Hyperparameter Configuration and Training Settings	29
4.5	Model Export Formats	29
4.6	Architecture Comparison on Test Set	31
4.7	Final Test Set Performance—ResNet50	32
4.8	Training Behaviour Summary	32
5.1	Contractual Specification for the <code>POST /predict</code> Classification End-point	38
5.2	Contractual Specification for the <code>GET /health</code> Operational Monitoring Endpoint	39
6.1	Core Technology Stack of the AgriVerify Frontend Application	46
6.2	Directory Structure and Component Responsibilities within <code>src/</code>	47
6.3	Frontend Application Route Namespace and Access Hierarchy	49
6.4	Authentication Service API Endpoint Integration Contracts	56
6.5	Seed Verification Service API Endpoint Integration Contracts	57
6.6	Complaint Management Service API Endpoint Integration Contracts	58

6.7	Administration and Analytics Service API Endpoint Integration Contracts	59
6.8	Shared Core UI Primitive Components in <code>src/components/ui/</code> . . .	61
7.1	Backend Technology Stack Components	66
9.1	Overall Model Performance Metrics	99
9.2	Per-Class Performance Analysis	100
9.3	Confusion Matrix Analysis	101
9.4	Regularization Effectiveness	101
9.5	Comparative Analysis of CNN Architectures	102
9.6	Inference Latency Analysis	104
9.7	Backend Performance Evaluation	104
9.8	Seed Verification Test Cases	105
9.9	Complaint Management Test Cases	106
9.10	Role-Based Access Control Test Cases	106

List of Figures

3.1	High-level system architecture of the AgriVerify platform. Solid arrows indicate request direction; dashed arrows indicate response direction. . .	18
3.2	Data flow diagram for the seed verification pipeline. The three inference calls (Custom Vision, OCR, ResNet-50) execute in parallel; results are aggregated before OpenAI synthesis.	20
3.3	Simplified entity-relationship diagram for the AgriVerify domain model. Cardinality notation follows UML conventions.	21
5.1	End-to-end model request processing flow. Raw image bytes undergo ASGI validation and tensor transformation before parallel deep learning classification and morphological DUS extraction.	35
5.2	Containerisation and automated deployment workflow. Source code changes trigger multi-stage container builds that are pushed to Artifact Registry before serverless GCP Cloud Run deployment.	40
6.1	Layered frontend architecture illustrating data flow from presentation UI through custom query hooks and domain services to the server-side Next.js API proxy route.	45
6.2	Logical folder organization of the AgriVerify frontend source tree, highlighting the separation between App Router navigation and core domain modules.	48
6.3	Authentication lifecycle diagram detailing JWT token persistence in Zustand and automated refresh token recovery via Axios HTTP 401 interception.	54
7.1	Clean Architecture Layer Hierarchy	65
8.1	End-to-End Request Flow from Browser to Machine Learning Pipeline .	86

8.2	Frontend and Backend Proxy Communication Workflow	89
8.3	Backend Model Serving and AI Inference Workflow	92
8.4	Azure Service Integration Workflow	95
8.5	Docker and Azure Container Apps Deployment Architecture	98

Symbols and Acronyms

$\approx$	Approximately equal to, used in performance evaluation
$\leq$	Less than or equal to, used in joint angle thresholding
π	Pi constant used in degree–radian conversion
θ	Angle between three landmarks
x, y	Coordinates of a detected keypoint in the frame
2D	2-Dimensional
3D	3-Dimensional
%	Percentage
AI	Artificial Intelligence
BGR	Blue-Green-Red color format
CNN	Convolutional Neural Network
COVID	CoronaVirus Disease of 2019
CV	Computer Vision
FPS	Frames Per Second
GB	Gigabyte
GHUM & GHUML	Generative 3D Human Shape and Articulated Pose Models
GUI	Graphical User Interface
iOS	iPhone Operating System
ML	Machine Learning
OpenCV	Open Source Computer Vision Library

RAM Random Access Memory

RGB Red-Green-Blue color format

Chapter 1

Introduction

1.1 Background and Context

Agriculture remains the foundational backbone of many global economies, heavily influencing food security, employment, and national income. At the core of all agricultural productivity is the quality of the seed planted; it is the most critical and non-negotiable input that determines crop yield, resilience to disease, and overall harvest quality. However, the agricultural sector is currently facing a severe crisis due to the proliferation of counterfeit or "fake" seeds in the market.

Counterfeit seeds are often packaged to resemble premium brands or are visually indistinguishable from high-yield varieties, yet they exhibit drastically lower germination rates and poor genetic traits. Traditional methods of verifying seed authenticity rely on destructive laboratory testing or subjective physical inspection, both of which are inaccessible, expensive, and far too slow for the average smallholder farmer. Consequently, farmers often discover they have been deceived only after planting, leading to catastrophic losses in time, resources, and income.

1.2 Motivation

The primary motivation for this project stems from the devastating socio-economic impact that fake seeds have on farming communities. When farmers unknowingly purchase and plant counterfeit seeds, they face severe financial distress, which cascades into broader issues of local food insecurity and loss of trust in agricultural supply chains. There is an urgent, unmet need for accessible technological tools that empower grassroots agricultural workers to protect themselves against fraud.

Simultaneously, the rapid advancement of deep learning, computer vision, and cloud computing presents an unprecedented opportunity to address this crisis. By leveraging Convolutional Neural Networks (CNNs), it is now possible to achieve laboratory-grade visual analysis automatically. The motivation is to harness these cutting-edge AI technologies and package them into a simple, web-based tool, putting the power of automated, real-time seed verification directly into the hands of the farmers.

1.3 Problem Statement

Despite the devastating effects of agricultural fraud, farmers currently lack a reliable, accessible, and rapid method to accurately distinguish between genuine and counterfeit seeds before the planting season begins. Visual similarities make manual detection nearly impossible, and existing institutional supply chains lack a transparent, real-time verification system.

Therefore, the core problem is the absence of an integrated, technology-driven solution that can automatically analyze seed imagery for subtle fraudulent anomalies, while also capturing localized, crowdsourced feedback from farmers to map and identify the spread of counterfeit seeds in real-time.

1.4 Objectives

The overarching goal of this project is to create an intelligent ecosystem for seed verification. To achieve this, the project outlines the following specific objectives:

1. To design and train a robust deep learning model (utilizing architectures such as ResNet) capable of accurately classifying seed images and identifying visual anomalies associated with counterfeit seeds.
2. To develop and deploy a highly accessible, user-friendly web-based platform that allows farmers to seamlessly upload images of their seeds for real-time AI inference and verification.
3. To integrate a comprehensive "Farmer Feedback System" within the platform, enabling users to report their post-purchase and post-planting experiences (such as

germination failure rates).

4. To utilize the crowdsourced feedback and newly uploaded images to continuously refine the database, creating an early warning system against emerging counterfeit supply chains.

1.5 Scope of the Project

The scope of this project encompasses the end-to-end development of an AI-driven web application tailored for agricultural quality control. This includes the collection, curation, and preprocessing of a targeted seed image dataset, followed by the training, validation, and optimization of a Convolutional Neural Network (CNN).

Additionally, the scope covers the full-stack development of the web platform, including an intuitive frontend interface for image uploads and a robust backend to handle model inference and database management. The project is strictly limited to the non-destructive, visual assessment of seeds via digital imagery; it does not extend to biochemical, genetic, or DNA-based testing methods. The final deliverable will be a functional prototype that demonstrates the feasibility of combining deep learning with crowdsourced feedback to combat agricultural fraud.

Chapter 2

Literature Review

2.1 Seed Quality Assessment: Traditional Methods

Conventional seed quality assessment relies on two broad approaches: manual visual inspection and laboratory-based biochemical testing. In field practice, inspectors evaluate seeds by examining size, shape, colour, and surface texture. These judgements are inherently subjective, vary between observers, and are difficult to standardise across districts or seasons. Laboratory tests such as the Tetrazolium (TZ) test and standard germination assays provide more defensible viability estimates, but both are destructive, the seed cannot be planted after testing and typically require four to seven days to complete under controlled conditions. This turnaround makes them impractical for screening large consignments at intake points, or for use by farmers who need a decision before the planting window closes.

Near-infrared (NIR) spectroscopy and X-ray transmission imaging have been investigated as non-destructive alternatives. NIR detects internal moisture gradients and biochemical composition without damaging the sample; X-ray imaging can reveal embryo defects invisible on the seed surface. Both methods perform well in laboratory conditions, but the equipment costs typically **USD 15,000–80,000** per unit and the need for trained operators confine their use to national seed-testing stations. Smallholder farmers and district-level inspectors in regions such as northeast India have no practical access to either technology, which is precisely where the need for affordable, rapid screening is most acute.

2.2 Deep Learning for Agricultural Image Classification

Convolutional Neural Networks (CNNs) have been applied to a range of agricultural

classification problems over the past decade, including leaf disease identification [1], weed recognition [2], and post-harvest quality grading [3]. Their main advantage over classical machine-learning pipelines is that spatial features are learned directly from raw pixel data during training, removing the need to manually engineer feature extractors for each new crop or defect type.

Applying CNNs to seed classification follows naturally from this body of work. Paddy seeds, for example, exhibit subtle differences in surface texture and colour between certified and counterfeit lots differences that are difficult to quantify with handcrafted descriptors but emerge reliably from sufficiently large training sets. The computational cost of running a trained CNN at inference time is also modest enough to support server-side web deployment, which matters when the target users have intermittent electricity but reasonable mobile-data coverage.

2.3 Transfer Learning and ResNet Architectures

Training a CNN from random initialization requires on the order of tens of thousands of labeled examples per class to achieve stable convergence. Annotated seed image datasets of that scale do not yet exist for most crop varieties, and certainly not for the specific task of detecting counterfeits. Transfer learning is the standard response: a model pre-trained on ImageNet (1.28 million labeled images across 1,000 categories) is fine-tuned on the smaller domain-specific dataset. The early convolutional layers, which encode low-level features such as edges and textures, are retained; only the later task-specific layers are updated. In practice, this approach consistently reaches accuracy close to what a from-scratch model would achieve on a dataset roughly ten times larger. ResNet architectures are well-suited to this fine-tuning scenario. The residual skip connections introduced by He et al. [4] allow gradients to propagate through very deep networks without vanishing, making it feasible to fine-tune models with 50 or more layers on datasets that would cause shallower architectures to overfit. ResNet-50 has become a common baseline in agricultural image classification because it balances parameter count (≈ 25 million) against accuracy on standard benchmarks, and pre-trained weights are publicly available. For this project, it provides a well-validated starting point that reduces both training time and the risk of overfitting on the available paddy seed dataset.

2.4 Farmer Feedback Systems

A detection model trained on a fixed dataset will degrade over time if counterfeiting practices evolve and no new training data are collected. Field-level feedback from farmers is one practical mechanism for capturing these shifts. When a farmer submits a

seed batch for image-based verification and later reports germination failure or abnormal yield from that batch, the report constitutes a labeled example a model prediction paired with a ground-truth outcome that can be incorporated into periodic retraining cycles.

Feedback also serves a more immediate function. Aggregating reports across users can flag a suspect lot or supplier before the model itself is updated: three independent failure reports from the same district within a fortnight is a meaningful signal, regardless of what the image classifier predicts. This kind of rule-based alert complements the model rather than replacing it. The practical difficulty is sustaining farmer participation; most deployments in low-resource contexts find that single-tap rating interfaces produce far higher response rates than free-text fields or structured forms [5].

2.5 Web-Based Agricultural Platforms

Deploying a trained model as a web service separates the computational workload from the end user’s device. A farmer uploads a seed image through a browser; the server runs inference and returns a result. This architecture has two concrete advantages over distributing a standalone mobile application: the model can be updated centrally without requiring client-side installations, and every prediction is logged in a structured format that supports monitoring and retraining.

For farmers in Manipur and similar regions, browser-based access via a mid-range Android handset is more realistic than assuming a dedicated app installation. Keeping the frontend lightweight and avoiding heavy JavaScript bundles and large media assets reduces page-load times on 4G connections with variable signal quality. Containerised (Dockerized) inference endpoints can handle the computationally intensive tasks without requiring a permanently running server, which reduces operating costs during off-season periods when query volume is low. Together, these design choices bring ML-assisted quality screening within reach of users who would otherwise have no access to any testing infrastructure.

2.6 Review of Related Work

Early automated seed classification used Support Vector Machines (SVMs) and Random Forests operating on handcrafted features derived from shape measurements, colour histograms, and texture descriptors. Ni et al. [6] classified rice seed varieties using morphological parameters extracted from binary silhouettes, achieving 94% accuracy across five varieties under controlled lighting. These methods perform adequately when intra-class variation is low and imaging conditions are uniform, but generalise poorly to photographs taken in field settings.

Deep learning approaches have largely replaced handcrafted methods in more recent work. Liu et al. [7] applied VGG16 transfer learning to maize seed defect detection and reported 97.3% accuracy on a 12-class dataset. Xu et al. [8] used ResNet-50 for soybean quality grading and found that fine-tuning from ImageNet weights outperformed training from scratch by approximately 8 percentage points when training data were limited to fewer than 2,000 images per class. However, both studies address naturally occurring defects — cracking, mould, insect damage rather than intentional substitution. A counterfeit seed is designed to resemble the genuine article; distinguishing the two is a harder classification problem than separating a damaged kernel from an intact one, because the visual gap is smaller and the adversarial nature of counterfeiting means it may shift over time. No published study directly evaluates CNN performance on verified counterfeit-versus-genuine seed pairs, and none integrates a structured feedback collection mechanism into the evaluation pipeline.

2.7 Summary and Research Gaps

The reviewed literature establishes that transfer-learned CNNs are effective for agricultural image classification, including seed quality assessment. Three specific gaps are directly relevant to this work. First, no published study has targeted intentional counterfeiting as distinct from natural quality defects; the two problems differ in that counterfeits may pass casual visual inspection and require the model to detect very subtle cues. Second, existing models are evaluated on laboratory-quality images collected under controlled conditions; performance on photographs taken by farmers using mobile devices in variable lighting has not been reported. Third, no end-to-end system has been published that connects model inference to a structured feedback collection mechanism and uses that feedback to retrain the model iteratively.

This project addresses all three gaps. The classifier is trained and evaluated specifically on a fake-versus-genuine paddy seed dataset; the web platform is tested with mobile-captured images; and the farmer feedback module is built directly into the application, with logged responses stored in a format suitable for future retraining rounds.

Chapter 3

System Design and Architecture

3.1 Overall System Architecture

The AgriVerify platform is engineered as a distributed, multi-tier web system that integrates deep learning inference, cloud-based AI orchestration, and a role-aware web application into a single coherent service. The architecture adheres to the **Clean Architecture** principle (Onion Architecture variant), which ensures that the core business logic remains entirely independent of infrastructure concerns such as databases, cloud vendors, and UI frameworks. Dependencies flow strictly inward: the outer layers (Infrastructure and API) depend on abstractions exposed by the inner layers (Application and Domain), never the reverse [9].

At the highest level, the platform comprises three major subsystems:

1. **Deep Learning Microservice** — A PyTorch/ResNet-50 model trained on a 1,563-image paddy seed corpus and served as a containerised FastAPI endpoint on Google Cloud Platform (GCP) Cloud Run [4].
2. **Backend API** — An ASP.NET Core 8 Web API acting as the primary orchestration layer. It authenticates users, mediates access to the ML microservice and Azure AI services (Custom Vision, AI Vision OCR, and Azure OpenAI), persists results to Microsoft SQL Server via Entity Framework Core, and exposes a versioned REST API consumed by the frontend [10].
3. **Frontend Web Application** — A Next.js 14 App Router application providing role-based dashboards for Farmers, Officers, and Administrators. All browser-to-backend traffic is forwarded through a server-side proxy route to centralise authentication token injection and prevent credential exposure in the browser [11].

3.1.1 Backend Architectural Layers

The backend follows the four-layer Onion model, each layer encapsulating a distinct concern as outlined in Table 3.1.

Table 3.1: Backend Architectural Layers and Responsibilities

Architectural Layer	Core Responsibilities and Encapsulated Components
Domain Layer	Contains core entities (User, VerificationHistory, ProductComplaint), value objects, enumerations (UserRole, VerificationStatus, ComplaintStatus), and domain exceptions. This layer carries no dependency on any external framework or service.
Application Layer	Orchestrates business workflows through application services (VerificationService, UserService, ComplaintService), defines Data Transfer Objects (DTOs), AutoMapper profiles, and FluentValidation rules. It depends only on the Domain layer.
Infrastructure Layer	Provides concrete implementations for interfaces declared in the Application layer: Entity Framework Core repositories, the ApplicationDbContext, Azure service wrappers (CustomVisionService, ComputerVisionService, OpenAIService, BlobStorageService), and migration management. This layer bridges the application to external systems without surfacing implementation details to upper layers.
API Layer	Hosts ASP.NET Core controllers (AuthController, VerificationController, ComplaintController, AdminController), the JWT middleware pipeline, Swagger/OpenAPI documentation, and the global exception handling middleware.

3.2 Functional Requirements

The functional requirements presented below were derived from stakeholder analysis involving farmers, agricultural officers, and domain experts. They form the contractual basis for all implementation decisions across the platform.

Table 3.2: Comprehensive Functional Requirements Specification

Req ID	Module / Role	Functional Specification
FR-01.1	Farmers (Auth)	Must register using their name, email address, state, and district, and authenticate via an email–password credential pair.
FR-01.2	Officers (Auth)	Must authenticate using an officer ID and a time-based one-time password (TOTP) as a second authentication factor.
FR-01.3	Administrators	Accounts are provisioned through database seeding; credentials are managed directly by the platform.
FR-01.4	Session Control	The system must issue JWT access tokens and refresh tokens upon successful login, and must invalidate both upon logout.
FR-02.1	Image Submission	Authenticated farmers must be able to upload front and back images of seed packaging for AI-powered authenticity analysis.
FR-02.2	AI Inference	The system must perform parallel inference using Azure Custom Vision for image classification, Azure AI Vision for OCR text extraction, and Azure OpenAI for contextual report synthesis.
FR-02.3	Result Output	A confidence score, a verification status (Genuine / Counterfeit / Suspicious), extracted OCR text, and a natural-language summary must be returned to the caller and persisted to the database.

Continued on next page

Table 3.2 continued from previous page

Req ID	Module / Role	Functional Specification
FR-02.4	Public Access	An anonymous public verification endpoint must be available to unauthenticated users.
FR-03.1	Farmers (Complaints)	Must be able to file complaints against seed batches, specifying issue type, affected district, severity level, and optional image evidence.
FR-03.2	Officers (Complaints)	Must be able to view assigned complaints, update their status (Open / Under Investigation / Resolved / Closed), and record investigation actions with timestamps.
FR-03.3	Admin (Complaints)	Must be able to view all complaints, filter by district and status, and assign unassigned complaints to available officers.
FR-04.1	Farmer board	Dash- Must have access to a personal dashboard displaying their verification history, complaint status, and summary metrics.
FR-04.2	Officer board	Dash- Must be able to view a complaint queue, perform batch risk assessments, and access officer-level analytics.
FR-04.3	Admin board	Dash- Must have access to platform-wide analytics, user management controls, and reference data configuration.
FR-05.1	Data Aggregation	The system must maintain a <code>BatchRiskRegistry</code> that aggregates complaint statistics per batch number.
FR-05.2	Risk Analytics	Must track affected districts, compute average severity scores, and assign categorical risk levels.

Continued on next page

Table 3.2 continued from previous page

Req ID	Module / Role	Functional Specification
FR-06.1	Feature Inference	The ResNet-50 microservice must accept seed images, apply the 224×224 preprocessing pipeline, perform forward inference, and return class probabilities alongside 17 extracted DUS (Distinctness, Uniformity, and Stability) physical parameters such as length, width, colour, texture, and shape metrics.
FR-06.2	Variety Matching	The microservice must perform variety matching against a reference database of 12 registered paddy varieties and return the top-3 matches with their similarity percentages.

3.3 Non-Functional Requirements

Non-functional requirements govern the quality attributes of the system and are treated as first-class design constraints throughout the implementation.

Table 3.3: Comprehensive Non-Functional Requirements Specification

Req ID	Quality Attribute	Architectural Constraint & Target Specification
NFR-01	Performance	The end-to-end verification pipeline (image upload $\rightarrow$ AI orchestration $\rightarrow$ database write $\rightarrow$ response) must complete within a latency budget suitable for interactive web use. Parallel execution of OCR and Custom Vision calls via <code>Task.WhenAll</code> is mandated to reduce blocking latency. The FastAPI microservice must serve single-image inference in sub-second wall-clock time on GCP Cloud Run hardware.

Continued on next page

Table 3.3 continued from previous page

Req ID	Quality Attribute	Architectural Constraint & Target Specification
NFR-02	Scalability	The backend must be stateless and deployable as a Dockerised container, enabling horizontal scaling without code changes [12]. The Next.js frontend must support server-side rendering and edge caching to handle high concurrent client traffic efficiently.
NFR-03	Security	All HTTP traffic must be served strictly over TLS (HTTPS). JWT access tokens must carry minimal claims and expire within a configurable time-to-live (TTL). Passwords must be stored as cryptographic hashes using robust key derivation functions such as PBKDF2 or bcrypt. Sensitive credentials (Azure keys, database connection strings, and API keys) must reside in environment variables and must never be hard-coded. Role-based endpoint protection must be strictly enforced through ASP.NET Core [Authorize] policies [13].
NFR-04	Reliability & Data Integrity	All grouped repository operations must execute within a Unit of Work transaction with atomic commit-or-rollback semantics. Soft-delete functionality must be implemented for complaint and verification records to support future legal auditing. Health check endpoints (/health and /api/health) must expose the connectivity status of all critical dependencies.

Continued on next page

Table 3.3 continued from previous page

Req ID	Quality Attribute	Architectural Constraint & Target Specification
NFR-05	Maintainability	The codebase must adhere to Clean Architecture layering with clear separation of concerns. Infrastructure services must be registered through an <code>InfrastructureServiceCollectionExtensions</code> extension class rather than inline in <code>Program.cs</code> . All public API endpoints must be fully documented via the Swagger/OpenAPI 3.0 specification.
NFR-06	Usability & Accessibility	The frontend must deliver a responsive layout usable on both desktop browsers and mobile devices, supporting farmers who access the platform via smartphones. The UI must provide immediate loading feedback through skeleton placeholders and toast notifications for asynchronous operations.
NFR-07	Portability	The deep learning microservice must be exported in multiple formats (<code>.pth</code> , TorchScript <code>.pt</code> , and ONNX <code>.onnx</code>) to support deployment across heterogeneous runtime environments. The entire application stack must be fully containerisable via Docker.

3.4 Tools and Technologies

3.4.1 Deep Learning Stack (PyTorch and ResNet-50)

The classification model is implemented entirely in **Python 3.10** using the following libraries [14]:

- **PyTorch** (`torch`, `torchvision`) — The primary deep learning framework used for model definition, training, and inference. PyTorch’s dynamic computation graph simplifies debugging and permits flexible architecture modifications during

transfer-learning fine-tuning phases.

- **ResNet-50** — A 50-layer deep residual network pre-trained on ImageNet (1.2M images, 1,000 classes) sourced from `torchvision.models` [4]. The backbone contributes approximately 23.5M parameters; a custom three-layer classification head ($2,048 \rightarrow 512 \rightarrow 128 \rightarrow 3$ units) adds a further 1.1M trainable parameters.
- **OpenCV** (`cv2`) — Used in the DUS parameter extraction pipeline for contour detection, pixel-to-millimetre calibration, and morphological analysis of seed images to compute the 17 physical parameters (colour mean/standard deviation, circularity, solidity, grain length/width, and related metrics) [15].
- **NumPy / Pandas** — Employed for numerical array manipulation and dataset statistics management [16].
- **scikit-learn** — Provides stratified train/validation/test split utilities and evaluation metric helpers [17].
- **Matplotlib / Seaborn** — Used for training curve visualisation and confusion matrix rendering.

Training was conducted on a GPU-enabled environment using the **AdamW** optimiser with weight decay $\lambda = 10^{-4}$, cross-entropy loss with label smoothing $\varepsilon = 0.1$, and a two-phase transfer learning schedule. Phase 1 (Epochs 1–8, backbone frozen, learning rate $= 10^{-3}$) allows the classification head to converge quickly, while Phase 2 (Epochs 9–15, full fine-tuning, learning rate $= 10^{-4}$) enables end-to-end weight adaptation. The final ResNet-50 model achieved a test accuracy of **99.57%** on the combined 1,563-image dataset, with a train-to-validation accuracy gap below 5%.

3.4.2 Backend Stack (ASP.NET Core 8, SQL Server, and Azure)

The backend REST API is built on the **Microsoft .NET 8** ecosystem. Table 3.4 summarises the key technology components and their roles within the system.

Azure AI services are consumed through dedicated service wrappers (`CustomVisionService`, `ComputerVisionService`, and `OpenAIService`) registered in the dependency injection container as singletons. The parallel orchestration pattern (`await Task.WhenAll(ocrTask, classificationTask)`) ensures that OCR extraction and image classification execute concurrently, minimising per-request latency [10].

Table 3.4: Backend Technology Stack

Component	Technology / Library
Web Framework	ASP.NET Core 8.0 / .NET 8 (Kestrel)
Database	Microsoft SQL Server
ORM	Entity Framework Core (Code-First, Migrations)
Authentication	ASP.NET Core Identity + JWT Bearer
Object Mapping	AutoMapper
Validation	FluentValidation + DataAnnotations
Logging	Serilog (Console and File sinks)
API Documentation	Swagger / OpenAPI 3.0
AI Services	Azure Custom Vision, Azure AI Vision (OCR), Azure OpenAI
Cloud Storage	Azure Blob Storage
Deployment	Docker / Render Cloud

3.4.3 Model Serving (FastAPI and Docker)

The trained PyTorch model is wrapped in a **FastAPI** microservice and deployed as a Docker container on **GCP Cloud Run** [18]. The inference pipeline within the microservice follows five sequential stages:

1. **Request** — The FastAPI endpoint receives a multipart image payload from the ASP.NET backend.
2. **Preprocessing** — The image is resized to 224×224 pixels and normalised using the ImageNet channel-wise mean $\mu = [0.485, 0.456, 0.406]$ and standard deviation $\sigma = [0.229, 0.224, 0.225]$.
3. **Model Inference** — The ResNet-50 model processes the normalised tensor and produces class logits.
4. **Post-processing** — Softmax converts logits to class probabilities; OpenCV routines extract the 17 DUS parameters; variety matching is performed against the 12-variety reference database.
5. **Response** — A structured JSON payload containing the prediction label, confidence score, physical parameters, DUS classification, and top-3 variety matches is returned to the backend.

The model is exported in three formats to maximise deployment flexibility: PyTorch native format (`.pth`, ≈ 98 MB), TorchScript (`.pt`, ≈ 97 MB, for C++ runtimes), and

ONNX (.onnx, ≈ 98 MB, for cross-platform inference). Continuous delivery is automated via a **GitHub Actions** CI/CD pipeline that triggers a Docker build-and-push to GCP Artifact Registry on every commit to the `main` branch.

3.4.4 Frontend Stack (Next.js, TypeScript, and TanStack Query)

The user-facing application is a **Next.js 14 App Router** single-page application written in **TypeScript**. The key technology choices are described below.

- **Next.js 14 App Router** — Provides file-system-based routing with route groups that cleanly separate the three role namespaces (`(farmer)`, `(officer)`, and `(admin)`). Server components handle initial data fetching; client components manage interactive UI behaviour.
- **TypeScript** — Enforces strict compile-time type safety across all layers, including shared DTO interfaces (`src/types/`), domain service modules (`src/services/`), and custom React hooks (`src/hooks/`).
- **TanStack Query (React Query v5)** — Manages all asynchronous server state, including cache invalidation, background refetch, pagination, and optimistic mutation rollback. Query hooks are co-located with domain modules; the global `QueryClient` disables `refetch-on-window-focus` to reduce incidental network requests [19].
- **Zustand** — A lightweight client-side state store for authentication state (`user`, `token`, `role`, `isAuthenticated`). The store is initialised from `localStorage` on the client side and exposes `setAuth` and `clearAuth` actions [20].
- **Axios** — Centralised HTTP client configured with a base URL of `/api/proxy`. A request interceptor automatically injects the Bearer token from storage; a response interceptor handles 401 errors by transparently attempting a token refresh before replaying the failed request.
- **Next.js API Route Proxy** — A server-side proxy route at `/api/proxy/[...path]` forwards all browser requests to the ASP.NET backend using the URL stored in the `BACKEND_API_URL` environment variable, preventing credential exposure in the browser.
- **Tailwind CSS** — Utility-first styling framework with extended design tokens (colours, border radii, and shadows) defined in `tailwind.config.ts`.

- **Shaden/UI + Custom Primitives** — A compact local component library providing Button, Input, Select, Dialog, Sheet, Table, Badge, Tabs, Skeleton, and Progress components built on native HTML elements with minimal abstraction overhead.

3.5 System Architecture Diagrams

3.5.1 High-Level Architecture

Figure 3.1 illustrates the end-to-end system topology. The user's browser communicates exclusively with the Next.js frontend server; no backend URLs are ever exposed to the client. All backend calls are proxied server-side to the ASP.NET Core API, which delegates to the ML microservice on GCP and to the suite of Azure AI services. Verification records are persisted in Microsoft SQL Server via Entity Framework Core, and seed images are stored in Azure Blob Storage.




Figure 3.1: High-level system architecture of the AgriVerify platform. Solid arrows indicate request direction; dashed arrows indicate response direction.

3.5.2 Data Flow Diagram

Figure 3.2 traces the data flow for the primary seed verification use case. When a farmer submits seed images, the request travels from the browser through the Next.js proxy, enters the ASP.NET verification pipeline, and is fanned out in parallel to the ML microservice and two Azure AI services. The combined results are synthesised by Azure OpenAI, persisted to the database, and returned as a structured JSON response.

3.5.3 Entity-Relationship Diagram

Figure 3.3 presents the simplified entity-relationship (ER) diagram for the core domain model persisted in Microsoft SQL Server. The model captures the principal entities and the cardinality of their associations.

Several design decisions embedded in the ER model merit explicit discussion:

- **Soft delete** — `IsDeleted` and `DeletedAt` fields on `VerificationHistory` and `ProductComplaint` preserve audit trails and support future legal review without physically removing records.
- **Third Normal Form (3NF) normalisation** — Seed package information, verification records, and complaint data are stored in separate tables, eliminating data redundancy and update anomalies.
- **Indexed columns** — `User.Email`, `VerificationHistory.UserId`, `ProductComplaint.UserId`, and `ProductComplaint.District` carry non-clustered indices to support efficient dashboard query execution.
- **BatchRiskRegistry** — Updated by a background aggregation job triggered on every complaint state change, this denormalised table provides fast, pre-computed risk summaries for officer-facing analytics without requiring expensive real-time aggregation queries.




Figure 3.2: Data flow diagram for the seed verification pipeline. The three inference calls (Custom Vision, OCR, ResNet-50) execute in parallel; results are aggregated before OpenAI synthesis.




Figure 3.3: Simplified entity-relationship diagram for the AgriVerify domain model. Cardinality notation follows UML conventions.

Chapter 4

Deep Learning Model for Paddy Seed Quality Assessment

4.1 Dataset Description

4.1.1 Data Sources and Collection (Kaggle)

The dataset used in this project was assembled from two distinct sources on Kaggle: one publicly available and one custom-collected by our team in collaboration with the Indian Council of Agricultural Research (ICAR), Manipur. Combining these two sources improved both the diversity and the real-world representativeness of the training data, ensuring that the model is exposed to paddy seed images captured under a variety of conditions and geographic contexts.

4.1.1.1 Public Source

The first source is the Paddy Seeds Quality Classification dataset published on Kaggle by Ayyan Arkadalkani:

```
https://www.kaggle.com/datasets/ayyanarkadalkani/  
paddy-seeds-quality-classification-dataset
```

This dataset comprises 1,213 labeled images distributed across three quality classes—Healthy, Damaged, and Fake/Low Quality—with a total storage footprint of approximately 183 MB. It provided a strong baseline of labeled examples captured under varied lighting and background conditions, making it suitable as a general-purpose foundation for transfer learning.

4.1.1.2 Custom Source

The second source is a custom dataset collected directly from ICAR, Manipur. Images of paddy seeds were captured under controlled laboratory conditions at the ICAR facility to reflect the specific seed varieties cultivated in the northeastern region of India. This collection totals 350 images and occupies approximately 24.7 MB of storage. The dataset has been published on Kaggle as an open-access resource to support future research in regional paddy seed analysis:

[https://www.kaggle.com/datasets/nilambarelangbam/
paddy-seed-images-collected-from-icar-manipur](https://www.kaggle.com/datasets/nilambarelangbam/paddy-seed-images-collected-from-icar-manipur)

By merging both datasets, we obtained a combined corpus of 1,563 images that is both larger and more geographically diverse than either source alone, better reflecting the range of seed conditions encountered in real agricultural settings across India.

4.1.2 Dataset Statistics and Class Distribution

After merging both datasets, the combined corpus contains 1,563 paddy seed images distributed across three quality classes. Table 4.1 summarises the source-wise composition, and Table 4.2 presents the class-wise distribution.

Table 4.1: Dataset Source Summary

Source	Images	Size	Kaggle Identifier
Paddy Seeds Quality Classification (Ayyan Arkadalkani)	1,213	~183 MB	ayyanarkadalkani/ paddy-seeds-quality-...
ICAR Manipur Custom Dataset (Our Collection)	350	~23.7 MB	nilambarelangbam/ paddy-seed-images-...
Combined Total	1,563	~208 MB	—

Table 4.2: Class-wise Distribution of the Combined Dataset

Class	Number of Images	Percentage (%)
Healthy Paddy Seeds	~750	~48.0
Damaged Paddy Seeds	~510	~32.6
Fake/Low Quality Seeds	~303	~19.4
Total	1,563	100.0

The dataset exhibits a moderate class imbalance, with Healthy seeds being the most represented class and Fake/Low Quality seeds the least. To prevent the model from

developing a bias toward the majority class, weighted random sampling was applied during training, ensuring that each class contributes proportionally to every training batch regardless of its raw frequency in the dataset.

4.1.3 Train/Validation/Test Split Strategy

We partitioned the 1,563 images into three non-overlapping subsets following a stratified split strategy. Stratified splitting ensures that the class proportions observed in the full dataset are preserved in each subset, preventing any split from being inadvertently dominated by a single quality class. Table 4.3 summarises the split.

Table 4.3: Dataset Split Statistics

Split	Proportion	Approx. Images	Purpose
Training Set	80%	~1,250	Model weight optimisation via back-propagation
Validation Set	10%	~157	Hyperparameter tuning and early stopping
Test Set	10%	~156	Final unbiased evaluation on unseen data

The training set is used exclusively for updating model weights. The validation set is monitored at the end of every epoch to guide hyperparameter decisions and to detect the onset of overfitting. The test set is held out entirely until the final evaluation stage, ensuring that all reported performance metrics reflect genuine generalisation to previously unseen examples.

4.1.4 Data Preprocessing and Augmentation

Before feeding images into the network, a standard preprocessing pipeline was applied to ensure compatibility with the ResNet50 architecture. All images were rescaled to 224×224 pixels—the standard input resolution for ImageNet-pretrained ResNet models. Pixel values were then normalised using the ImageNet channel-wise mean and standard deviation:

$$\mu = [0.485, 0.456, 0.406], \quad \sigma = [0.229, 0.224, 0.225] \quad (4.1)$$

to align the input distribution with that seen during backbone pre-training.

To expand the effective diversity of the training set and improve the model’s robustness to real-world image variability, eight data augmentation techniques were applied exclusively during training:

1. **Random Resized Crop**—Randomly crops a region of the image and resizes it to 224×224 pixels, simulating varying camera distances and seed positions within

the capture frame.

2. **Horizontal Flip**—Randomly mirrors the image along the vertical axis with probability 0.5, exploiting the bilateral symmetry common to seed grains.
3. **Vertical Flip**—Randomly mirrors the image along the horizontal axis with probability 0.5, further increasing orientation diversity in the training set.
4. **Random Rotation**—Rotates the image by a randomly sampled angle within $\pm 30^\circ$, accounting for seeds that are not perfectly aligned during imaging.
5. **Translation/Random Shift**—Randomly shifts the image content along the x and y axes, preventing the model from relying on the absolute position of the seed within the frame.
6. **Colour Jitter**—Randomly perturbs brightness, contrast, saturation, and hue within specified bounds, enhancing robustness to lighting variations in agricultural settings.

4.2 Model Architecture

4.2.1 ResNet50 Backbone (25.6M Parameters)

ResNet50 was selected as the backbone architecture for this classification task. This deep residual network contains 50 layers and 25.6 million trainable parameters. The architecture is built upon residual blocks, which employ skip connections to enable the training of very deep networks by mitigating the vanishing gradient problem.

The pre-trained weights from the ImageNet dataset were used as the initial checkpoint, providing a rich set of learned feature representations for natural images. Transfer learning was then applied: the backbone layers were initially frozen during Phase 1 training, with only the custom classification head being optimised. This approach leverages the generalisable features learned on ImageNet while adapting the model to the specific task of seed quality classification.

4.2.2 Custom Classification Head

Upon the frozen ResNet50 backbone, a custom classification head was appended to perform the final three-way classification (Healthy, Damaged, Fake/Low Quality). The head comprises:

1. Global Average Pooling layer (reducing the spatial dimensions from 7×7 to a single vector)
2. Fully connected layer with 512 hidden units
3. Batch Normalisation layer
4. ReLU activation function
5. Dropout layer (dropout rate: 0.5)
6. Output layer with 3 units (one per class) and softmax activation

This architecture ensures that the model learns class-specific patterns while maintaining numerical stability during training.

4.2.3 Anti-Overfitting Techniques

To prevent overfitting and improve generalisation to unseen data, multiple regularisation techniques were employed:

- **Dropout**—Applied in the classification head (rate: 0.5) to randomly deactivate neurons during training, reducing co-adaptation.
- **Batch Normalisation**—Applied after linear transformations to normalise layer inputs, stabilising training and acting as a regulariser.
- **Label Smoothing**—Cross-entropy loss was computed with label smoothing ($\epsilon = 0.1$), preventing the model from assigning overconfident predictions.
- **Weight Decay (L2 Regularisation)**—A weight decay coefficient of 0.0001 was applied to all trainable parameters, penalising large weights.
- **Data Augmentation**—As described in Section 4.1.4, eight augmentation techniques were applied to increase training set diversity.
- **Early Stopping**—Training was monitored on the validation set, and the best model weights were saved. Training was terminated upon detection of overfitting (validation loss increase).

4.3 DUS Parameter Integration

4.3.1 Feature Extraction—17 Physical Parameters

The Distinctiveness, Uniformity, and Stability (DUS) parameters are standard morphological descriptors used in agricultural research to characterise seed varieties. Our model was extended to extract 17 physical parameters from each seed image using image processing techniques. These parameters include:

- Seed length and width
- Aspect ratio (length/width)
- Perimeter and area
- Circularity and solidity
- Eccentricity
- Moments of inertia
- Hu moments (7-dimensional shape descriptor invariant to rotation, scale, and translation)

These features provide a direct mapping from visual characteristics to standardised agricultural descriptors, enabling variety matching and quality assessment beyond class labels.

4.3.2 Pixel-to-Physical Unit Conversion

To convert pixel-based measurements to real-world physical units (millimetres), a calibration step was performed. A reference object of known size (calibration ruler) was included in the imaging setup. The pixel-to-millimetre conversion factor was computed as:

$$\text{scale} = \frac{\text{reference_length_mm}}{\text{reference_length_pixels}} \quad (4.2)$$

All extracted measurements were then scaled using this factor, ensuring that reported physical parameters are in standard units independent of camera resolution or imaging distance.

4.3.3 Variety Matching (12 Reference Varieties)

A reference database of 12 authoritative paddy seed varieties was compiled from ICAR documentation. Each reference variety is represented by median values of the 17 physical parameters computed from multiple certified images. When processing a new image, the extracted parameters are compared against all 12 reference varieties using Euclidean distance:

$$d_i = \sqrt{\sum_{j=1}^{17} (p_j - r_{i,j})^2} \quad (4.3)$$

where p_j is the j -th extracted parameter, $r_{i,j}$ is the j -th median reference value for variety i , and d_i is the distance to variety i . The variety with the smallest distance is reported as the predicted variety.

4.4 Training Pipeline

4.4.1 Phase 1—Classification Head Training (Epochs 1–8)

During Phase 1, the backbone weights were frozen and only the custom classification head was optimised. This phase lasted 8 epochs with a learning rate of 0.001. The objective was to rapidly converge the head weights to a reasonable initialisation before fine-tuning the entire network.

The training batch size was set to 32 samples, and the loss function was Cross-Entropy with label smoothing ($\epsilon = 0.1$). Validation performance was monitored at the end of each epoch to ensure stable convergence and to detect early signs of overfitting.

4.4.2 Phase 2—Full Fine-Tuning (Epochs 9–15)

During Phase 2 (epochs 9–15), all backbone layers were unfrozen and the entire network was jointly optimised. The learning rate was reduced by a factor of 10 to 0.0001, a common strategy to avoid catastrophic forgetting of pre-trained weights. This phase refined the backbone features to be better suited to the seed classification task.

Training was halted at epoch 15 when early stopping criteria were triggered (validation loss increased, indicating onset of overfitting). The best model weights from the entire training run were restored and used for final evaluation.

4.4.3 Hyperparameter Configuration and Optimiser

Table 4.4 summarises the complete hyperparameter configuration used during training.

Table 4.4: Hyperparameter Configuration and Training Settings

Hyperparameter	Value	Notes
Optimiser	AdamW	Decoupled weight decay
Phase 1 Learning Rate	0.001	Head-only training
Phase 2 Learning Rate	0.0001	Full fine-tuning (10× reduction)
Weight Decay (L2)	0.0001	Applied to all trainable weights
Batch Size	32	GPU-optimised mini-batch
Loss Function	Cross-Entropy + Label Smoothing	$\epsilon = 0.1$, prevents overconfidence
LR Scheduler	StepLR	step=8, $\gamma = 0.1$
Gradient Clipping	Max norm = 1.0	Prevents exploding gradients
Early Stopping Patience	10 epochs	Monitors validation loss
Input Resolution	224×224 pixels	ResNet50 standard input
Total Epochs	15	Halted at overfitting onset

The AdamW optimiser was selected for its adaptive learning rate properties and decoupled weight decay, which provides better generalisation compared to standard L2 regularisation in the original Adam implementation.

4.4.4 Model Export Formats

Upon completion of training, the final model was exported in multiple formats to support deployment across different runtime environments. Table 4.5 lists the export formats and their intended use cases.

Table 4.5: Model Export Formats

Format	Extension	Size	Use Case
PyTorch Native	.pth	~98 MB	Primary format; FastAPI serving
TorchScript	.pt	~97 MB	Production deployment; C++ runtimes
ONNX	.onnx	~98 MB	Cross-platform; non-PyTorch frameworks
Configuration	.json	~2 KB	Model metadata, labels, hyperparameters
DUS Reference DB	.csv	<1 MB	12 reference varieties for matching

4.5 Model Deployment

Following training, the model was integrated into a production-ready web service and deployed to Google Cloud Platform (GCP). The deployment architecture is designed for scalability, reproducibility, and continuous delivery, ensuring that the model is accessible as a live API endpoint for downstream applications and testing.

4.5.1 Serving Infrastructure

The trained model (`.pth`) is loaded and served via FastAPI, a high-performance Python web framework. A single API endpoint accepts an image as input, runs it through the ResNet50 inference pipeline, and returns the quality prediction, confidence scores, and extracted physical parameters in a structured JSON response. The inference flow is as follows:

1. **Request**—FastAPI endpoint receives the image payload.
2. **Preprocessing**—Image is resized to 224×224 pixels and normalised using ImageNet statistics.
3. **Model Inference**—ResNet50 processes the image and produces class logits.
4. **Post-processing**—Softmax converts logits to probabilities; DUS parameters are extracted and variety matching is performed.
5. **Response**—Structured JSON containing prediction, confidence, physical parameters, DUS classification, and variety matches.

4.5.2 Containerisation and CI/CD Pipeline

The entire application—model weights, FastAPI server, and all dependencies—is packaged inside a Docker container, ensuring a consistent and reproducible runtime environment across development and production. Deployment is fully automated through a GitHub Actions CI/CD pipeline that executes the following steps on every push to the main branch:

1. **Build**—Docker image is built from the project Dockerfile.
2. **Push**—The built image is pushed to GCP Artifact Registry for versioned storage.
3. **Deploy**—The new image is deployed to GCP Cloud Run, which automatically provisions a serverless, auto-scaling container instance.

GCP Cloud Run was selected for its serverless auto-scaling model, which eliminates idle resource costs while handling concurrent inference requests elastically. The live API endpoint is publicly accessible at:

`https://resnet-api-297419987783.asia-south1.run.app/`

4.6 Model Comparison and Selection

To justify the selection of ResNet50 as the final architecture, a comparative evaluation of four widely used convolutional neural network architectures was conducted on the same dataset, training split, and hyperparameter configuration. The models evaluated were ResNet50, MobileNetV2, DenseNet121, and GoogleNet. Table 4.6 presents the results.

Table 4.6: Architecture Comparison on Test Set

Model	Accuracy	Precision	Recall	F1 Score	Time (min)
ResNet50	99.57%	99.57%	99.57%	99.57%	16.4
DenseNet121	99.57%	99.57%	99.57%	99.57%	15.0
GoogleNet	99.57%	99.57%	99.57%	99.57%	12.2
MobileNetV2	99.13%	99.15%	99.13%	99.13%	8.5

ResNet50, DenseNet121, and GoogleNet achieved identical peak accuracy of 99.57% on the test set, demonstrating that the dataset is well within the representational capacity of all three architectures. MobileNetV2 achieved a slightly lower accuracy of 99.13% but completed training in only 8.5 minutes—approximately half the time of ResNet50—making it a viable alternative in resource-constrained deployment scenarios.

ResNet50 was selected as the primary model for three reasons:

1. Its residual skip connections provide well-understood gradient flow properties that simplify fine-tuning behaviour analysis.
2. It has the broadest ecosystem support across deployment frameworks (PyTorch, TorchScript, ONNX, TensorRT).
3. Its architecture aligns with the majority of published benchmarks in agricultural image classification, making our results directly comparable to prior work.

The 4.4-minute training overhead compared to GoogleNet is a negligible cost given that training is performed only once.

4.7 Final Model Performance Summary

The final ResNet50 model was evaluated on the held-out test set following the completion of Phase 2 training. Table 4.7 presents the overall performance metrics, and Table 4.8 summarises key training behaviour observed across both phases.

Table 4.7: Final Test Set Performance—ResNet50

Metric	Value	Status
Test Accuracy	99.57%	Exceeds industry benchmark (85–90%)
Precision	~0.99	Excellent—minimal false positives
Recall	~0.99	Excellent—minimal false negatives
F1 Score	~0.99	Excellent—strong harmonic balance
Train-Validation Gap	<5%	Excellent generalisation; overfitting controlled

Table 4.8: Training Behaviour Summary

Observation	Detail
Initial Convergence	Rapid accuracy improvement from ~86% in Epoch 1 to near-peak within ea
Loss Trajectory	Decreased from ~1.0 to near 0.0 across 15 epochs
Mid-training Stability	Stable, consistent learning after mid-training phase with no erratic oscillation
Overfitting Control	Train-validation gap remained below the critical 5% threshold throughout tra
Training Stopped at	Epoch 15—onset of overfitting detected; best weights preserved

The final accuracy of 99.57% significantly exceeds the typical industry benchmark of 85–90% for automated seed quality classification systems. The near-perfect precision, recall, and F1 score across all three quality classes confirm that the model can reliably distinguish healthy seeds from damaged and fake/low quality seeds—a critical requirement for practical agricultural deployment at Manipur Technical University and partner ICAR facilities.

Chapter 5

Model Serving API

Modern agricultural machine learning deployments require robust, low-latency, and highly scalable inference infrastructure to translate offline trained models into operational real-time services. While model development focuses on architectural optimization and empirical accuracy, the deployment tier dictates the real-world responsiveness, concurrent throughput, and operational resilience of the verification platform [9]. Within the AgriVerify architecture, the PyTorch ResNet-50 classification model (detailed in Chapter 4) is encapsulated within a dedicated, high-performance microservice. This microservice exposes automated endpoints consumed by the ASP.NET Core orchestration backend, isolating deep learning computation from general business logic.

This chapter details the architectural design and software engineering implementation of the Model Serving microservice constructed using FastAPI and Uvicorn [18]. The exposition examines the asynchronous request processing lifecycle, the rigorous tensor normalization and parameter extraction pipeline, and comprehensive API specifications formulated as exact contractual tables. Furthermore, containerised deployment strategies using Docker multi-stage builds on Google Cloud Platform (GCP) Cloud Run are presented alongside configuration management, observability, and defensive security hardening techniques [12], [21].

5.1 API Architecture and Request Orchestration

5.1.1 FastAPI and Uvicorn Serving Stack

The inference microservice is built upon the FastAPI framework running on top of the Uvicorn Asynchronous Server Gateway Interface (ASGI) runtime [22]. FastAPI was selected over traditional WSGI frameworks due to its exceptional asynchronous I/O concurrency, built-in validation powered by Pydantic type annotations, and automatic

OpenAPI 3.0 schema generation [23], [24]. Uvicorn provides a lightning-fast event loop implementation based on `uvloop`, enabling the microservice to handle thousands of concurrent network connections with minimal memory overhead.

The serving stack is strictly modularised into discrete operational layers: network request deserialization, Pydantic schema validation, OpenCV image preprocessing, PyTorch tensor execution, and JSON payload formatting. This decoupled design ensures that computationally intensive PyTorch tensor operations execute without blocking the underlying ASGI web server event loop [25].

5.1.2 End-to-End Model Request Flow

The lifecycle of an inference request begins when the ASP.NET Core backend issues an authenticated multipart POST request to the microservice. The request payload contains the raw binary stream of a seed packaging image captured by a farmer or agricultural officer. Upon arrival, FastAPI's routing layer intercepts the payload and validates the multipart boundaries and MIME type constraints. Once validated, the byte stream is decoded in-memory into a continuous OpenCV matrix representation.

The image matrix is subjected to standardization transformations before being converted into a normalized PyTorch tensor. The PyTorch inference engine, pre-loaded with frozen ResNet-50 weights during worker initialization, performs a forward propagation pass across the tensor to compute raw classification logits [4], [14]. In parallel with the deep network forward pass, OpenCV contour detection algorithms isolate the physical seed morphology within the image frame to extract 17 distinct Distinctness, Uniformity, and Stability (DUS) parameters [15]. The raw logits are passed through a softmax activation function to derive normalized class probabilities, identifying whether the sample is Genuine, Counterfeit, or Suspicious. Finally, the extracted physical parameters are cross-referenced against a reference database of 12 registered paddy varieties to establish top-3 genetic matches. The aggregated metrics are serialized into a strict JSON schema and returned to the caller.

5.1.3 Concurrency and Performance Optimization

Serving complex convolutional neural networks under real-time constraints requires strict resource management to prevent thread starvation and memory saturation. The microservice implements several explicit performance optimizations:

- **Model Warm-Loading:** The ResNet-50 PyTorch weights (`.pth` format) are loaded into system memory exactly once during Uvicorn worker startup using




Figure 5.1: End-to-end model request processing flow. Raw image bytes undergo ASGI validation and tensor transformation before parallel deep learning classification and morphological DUS extraction.

FastAPI’s lifespan event handlers. This architectural pattern completely eliminates cold-start latency during subsequent inference requests.

- **Asynchronous Non-Blocking Execution:** Disk I/O, network communication, and database logging are executed asynchronously via Python’s `asyncio` event loop. Computationally heavy PyTorch inference calls are offloaded to dedicated thread pools using `run_in_executor()` to preserve server responsiveness.
- **Multi-Worker Process Scaling:** In multi-core production environments, Uvicorn is launched with multiple worker processes (`-workers 4`) managed by Gunicorn. Each worker maintains an independent Python runtime and model instance, maximizing CPU utilization across virtual machine hardware.

5.2 Inference Execution and Preprocessing Pipeline

5.2.1 Image Preprocessing and Tensor Normalization

Deep neural network classifiers are highly sensitive to variations in input distribution. To guarantee perfect alignment between the training domain and operational inference, incoming image frames are subjected to a deterministic preprocessing sequence using OpenCV and PyTorch Vision transforms [15], [16].

The pipeline accepts standard compressed raster formats (`image/jpeg`, `image/png`) and executes the following sequential transformations:

1. Raw incoming HTTP byte buffers are decoded directly into an uncompressed 8-bit

OpenCV BGR array without intermediate file-system disk I/O.

2. The color space is converted from BGR to RGB to conform to standard PyTorch training pipelines.
3. The image canvas is resized to exactly 224×224 pixels using bicubic interpolation to preserve fine-grained structural details of the seed packaging texture.
4. Pixel intensity values are scaled from the integer domain $[0, 255]$ to the continuous floating-point interval $[0.0, 1.0]$.
5. Channel-wise standardization is applied using the ImageNet statistical moments:

$$\mathbf{X}_{\text{norm}} = \frac{\mathbf{X} - \boldsymbol{\mu}}{\boldsymbol{\sigma}} \quad (5.1)$$

where $\boldsymbol{\mu} = [0.485, 0.456, 0.406]$ and $\boldsymbol{\sigma} = [0.229, 0.224, 0.225]$.

6. The tensor is expanded along the zero-th axis (`unsqueeze(0)`) to create a batch dimension of size $B = 1$, conforming to the expected input shape $(1, 3, 224, 224)$ of the ResNet-50 architecture.

5.2.2 Prediction, Confidence Scoring, and Latency Measurement

Following forward propagation through the ResNet-50 network, the unnormalized output logits $\mathbf{z} = [z_1, z_2, z_3]$ corresponding to the three classification categories (Genuine, Counterfeit, Suspicious) are passed through a softmax activation function:

$$P(y = i \mid \mathbf{z}) = \frac{e^{z_i/T}}{\sum_{j=1}^3 e^{z_j/T}} \quad (5.2)$$

where $T = 1.0$ represents the temperature parameter. The class index i^* yielding the maximum probability is selected as the top-1 prediction.

To prevent overconfident classifications when encountering out-of-distribution or heavily corrupted images, the microservice enforces an uncertainty rejection threshold. If the top-1 confidence score $P(y = i^*) < 0.75$, the verification status is automatically flagged as `Suspicious` requiring manual officer inspection.

Operational observability is maintained by recording precise wall-clock execution times. The microservice utilizes Python's high-resolution monotonic clock (`time.perf_counter()`) to compute latency across the preprocessing, model execution, and post-processing phases. The total elapsed time is embedded directly into the response payload as `inference_ms`, providing downstream orchestrators with continuous performance telemetry.

5.3 API Specification and Endpoint Contracts

The microservice exposes a tightly constrained RESTful API governed by strict Pydantic validation contracts [24]. Tables 5.1 and 5.2 detail the precise structural definitions, request headers, payload schemas, and HTTP response codes for the primary classification and health monitoring endpoints.

Table 5.1: Contractual Specification for the `POST /predict` Classification Endpoint

Specification Item	Endpoint Details and Contractual Constraints
HTTP Method & URL	<code>POST /predict</code> (Accessible strictly via HTTPS on port 443 / 8000)
Headers	<code>Authorization: Bearer <token></code> (Secret inter-service API key) <code>Content-Type: multipart/form-data</code>
Request Parameters	image (File): The binary raster file stream. Permitted MIME types are <code>image/jpeg</code> and <code>image/png</code> . Maximum file size constraint is 10 MB.
Success Response (Schema Definition)	200 OK — Returned as a JSON payload: <pre>{ "predicted_label": "Genuine", "confidence": 0.9842, "inference_ms": 42.8, "dus_parameters": { "grain_length_mm": 8.42, "grain_width_mm": 2.65, "aspect_ratio": 3.17, "solidity": 0.981 }, "variety_matches": [{"variety": "IR64", "similarity": 96.4}, {"variety": "Swarna", "similarity": 82.1}] }</pre>
Client Error Codes	400 Bad Request: Malformed multipart payload or missing image field. 413 Payload Too Large: Uploaded file exceeds the 10 MB size limit. 415 Unsupported Media Type: Uploaded file is not a valid JPEG or PNG. 422 Unprocessable Entity: Failed Pydantic schema validation.

Table 5.2: Contractual Specification for the `GET /health` Operational Monitoring Endpoint

Specification Item	Endpoint Details and Contractual Constraints
HTTP Method & URL	<code>GET /health</code> (Unauthenticated endpoint for automated probes)
Headers	<code>Accept: application/json</code>
Probing Semantics	Liveness Probe: Verifies the Uvicorn ASGI process is executing. Readiness Probe: Confirms ResNet-50 weights are successfully loaded into memory and capable of receiving tensor operations.
Success Response	200 OK — Lightweight JSON operational status payload: <pre>{ "status": "online", "model_loaded": true, "device": "cuda:0", "version": "1.2.0", "uptime_seconds": 34821 }</pre>
Failure Status	503 Service Unavailable: Process running but model weights corrupted or failed to initialize during startup lifecycle.

5.4 Containerisation and Deployment Architecture

5.4.1 Docker Multi-Stage Build Architecture

To ensure perfect reproducibility across diverse compute environments while maintaining minimal production footprint, the Model Serving microservice is packaged using a multi-stage Docker build pipeline [12]. Packaging heavy machine learning frameworks such as PyTorch, OpenCV, and NumPy traditionally results in monolithic container images exceeding several gigabytes. The multi-stage architecture systematically isolates compilation toolchains from runtime execution.

1. **Builder Stage:** Built upon a fully featured Python 3.10 Debian base image, this stage installs compilation tools (`build-essential`, `gcc`) and compiles all dependencies declared in `requirements.txt` into binary Python wheels (`.whl`).
2. **Runtime Stage:** Constructed from an ultra-slim, security-hardened base image (`python:3.10-slim-bookworm`), this stage copies only the pre-compiled wheel binaries from the builder stage. It installs minimal runtime system libraries required by OpenCV (`libgl1`, `libglu1-mesa`), establishes a non-root system user (`appuser`), and copies the application code alongside pre-trained ResNet-50 weights.

This decoupling reduces the final production container image from over 3.2 GB to approximately 850 MB, significantly reducing network transfer latency and vulnerability attack surface during automated deployments.




Figure 5.2: Containerisation and automated deployment workflow. Source code changes trigger multi-stage container builds that are pushed to Artifact Registry before serverless GCP Cloud Run deployment.

5.4.2 Serverless Deployment on GCP Cloud Run

The resulting production image is deployed to Google Cloud Platform (GCP) Cloud Run, a fully managed serverless container runtime [21]. Cloud Run automatically abstracts underlying infrastructure management, providing immediate horizontal elasticity. Under zero-traffic conditions, the container instances scale to zero to minimize operational costs. During traffic spikes, Cloud Run dynamically provisions concurrent container instances, routing traffic through an automated HTTPS load balancer with TLS termination.

Each Cloud Run instance is provisioned with 4 vCPUs and 8 GB of RAM, ensuring sufficient memory capacity for in-memory PyTorch tensor execution while preventing out-of-memory kernel terminations.

5.5 Configuration, Observability, and Security Hardening

5.5.1 Environment Management and Secret Injection

In strict adherence to the 12-Factor application methodology, all runtime configuration parameters and sensitive credentials are fully decoupled from the application source code [13]. Runtime behavior is governed by environment variables managed through GCP Secret Manager and injected securely into the container at runtime.

Key environment variables include:

- `MODEL_WEIGHTS_PATH`: File path locating the pre-trained `resnet50_paddy.pth` model.
- `API_KEY_SECRET`: Cryptographic secret required in the Bearer token authorization header to authenticate requests from the ASP.NET Core backend.
- `ALLOWED_CORS_ORIGINS`: Strictly defined whitelist of permitted origins to prevent cross-site request forgery.
- `MAX_IMAGE_SIZE_MB`: Configurable constraint limiting upload payloads to 10 MB.

5.5.2 Threat Model and Defensive Hardening

Exposing machine learning inference endpoints to external traffic introduces unique attack vectors, including denial-of-service via massive payload bombs, memory corruption through malformed image matrices, and reverse engineering of model parameters. The microservice enforces multiple layers of defensive security:

- **Input Sanitization and Bounds Checking:** Incoming byte streams are intercepted at the ASGI middleware layer. If the `Content-Length` header exceeds 10 MB, the connection is instantly terminated before allocating memory buffer space. Pydantic validation guarantees that parameters adhere strictly to expected data types [24].

- **Transport Security:** Plaintext HTTP communication is strictly blocked. Cloud Run enforces TLS 1.3 encryption across all endpoints, protecting data-in-transit against man-in-the-middle eavesdropping.
- **Graceful Exception Handling:** Unhandled server errors are intercepted by global exception handlers. Standardized JSON error responses (as defined in Table 5.1) are returned to the client, suppressing internal Python stack traces or file-system paths that could facilitate vulnerability exploitation.

5.6 Summary

This chapter presented the comprehensive software architecture and engineering implementation of the Model Serving microservice for the AgriVerify platform. By leveraging FastAPI and Uvicorn, the system achieves exceptional asynchronous request concurrency while isolating PyTorch ResNet-50 tensor computation and OpenCV morphological DUS feature extraction. Contractual API specifications were established through explicit tables for prediction and health endpoints, ensuring robust inter-service communication. Finally, the integration of Docker multi-stage builds and serverless deployment on GCP Cloud Run, reinforced by 12-factor configuration management and defensive security hardening, delivers a publication-ready, enterprise-grade AI inference service capable of real-world agricultural verification.

Chapter 6

Frontend System Architecture and Implementation

Modern agricultural platforms operating across rural and semi-urban environments demand responsive, intuitive, and highly reliable user interfaces. While the deep learning microservice (Chapter 5) executes complex classification algorithms and the ASP.NET Core backend (Chapter 7) manages business orchestration and relational persistence, the frontend application serves as the tangible contact point for farmers, agricultural officers, and system administrators [26]. The primary software engineering objective for the AgriVerify frontend is to deliver a role-partitioned, highly responsive web application capable of capturing high-resolution seed imagery in the field, displaying complex multi-model AI inference reports, and managing multi-stage administrative workflows over varying network bandwidths.

This chapter details the architectural design and technical implementation of the AgriVerify frontend built using Next.js 15, React 19, TypeScript, and Tailwind CSS [27]–[30]. The discussion explores the decoupled component-service architecture, the file-system-based routing hierarchy leveraging Next.js route groups, and rigorous client-side role enforcement. Furthermore, robust authentication workflows including Zustand in-memory state management and automated Axios refresh-token recovery are presented alongside comprehensive API integration contracts and UI component design systems [31]–[33].

6.1 Architecture and Design Principles

The AgriVerify frontend is structured around a decoupled, layered component-service architecture [26]. In contrast to tightly coupled web applications where data fetching, state mutation, and business logic are embedded directly within presentation components, AgriVerify enforces strict separation of concerns across four discrete functional layers:

presentation UI, query orchestration, domain service modules, and transport proxying.

1. **Presentation UI Layer:** Composed of thin React Server Components and interactive Client Components. This layer declares layout structures and visual styling but delegates all asynchronous state transitions to dedicated custom hooks.
2. **Query Orchestration Layer:** Encapsulated within custom React hooks powered by TanStack Query (React Query v5) [32]. This layer manages server-state caching, background refetching, cache invalidation, and optimistic UI updates.
3. **Domain Service Layer:** A collection of modular TypeScript service files (`authService.ts`, `verificationService.ts`, `complaintService.ts`, `adminService.ts`) that encapsulate backend REST API endpoints and data transfer object (DTO) validation.
4. **Transport and Proxy Layer:** Configured via a centralized Axios HTTP client that automatically handles authentication token injection and error interception, forwarding requests to a secure Next.js server-side proxy route.

This architectural separation significantly enhances maintainability and testability. UI components can be developed and tested in isolation using mock query hooks; backend endpoint modifications require updates only within the corresponding domain service file without disrupting the presentation layer.

6.2 Technology Stack

The technology stack underpinning the AgriVerify frontend was selected to maximize compile-time type safety, client-side rendering performance, and developer productivity. Table 6.1 details the primary frameworks and libraries integrated into the frontend ecosystem.




Figure 6.1: Layered frontend architecture illustrating data flow from presentation UI through custom query hooks and domain services to the server-side Next.js API proxy route.

Table 6.1: Core Technology Stack of the AgriVerify Frontend Application

Component / Library	Version	Architectural Role and System Function
Next.js	15.2.4	Core application framework providing file-system routing, server components, route groups, and the API proxy pipeline [27].
React	19.2.5	UI rendering engine supporting concurrent mode and reactive client state [28].
TypeScript	5.7.3	Enforces rigorous compile-time type safety across DTO contracts, UI props, and state models [29].
Tailwind CSS	3.4.17	Utility-first CSS framework managing responsive layouts and design tokens [30].
Axios	1.7.9	Centralized HTTP client managing request headers and token interception [31].
TanStack Query	5.59.0	Manages asynchronous server state, cache invalidation, and optimistic mutations [32].
Zustand	5.0.1	Lightweight in-memory global state store managing authentication profiles [33].
Framer Motion	12.38.0	Orchestrates fluid page transitions and micro-animations across dashboard views.
Lucide React	1.14.0	Comprehensive iconography library providing standardized visual glyphs.
react-webcam	7.2.0	HTML5 media device interface enabling mobile and desktop camera capture for seed scanning.

Application typography is managed via Next.js native Google Font optimization (`next/font/google`). The root layout preloads `DM Sans` for body copy and `Sora` for display headings, injecting both font families as CSS custom properties into the DOM. This technique eliminates render-blocking external web font requests, ensuring instant text rendering and preventing cumulative layout shifts (CLS).

6.3 Folder Structure and Repository Layout

The source code repository is organized inside a dedicated `src/` root directory to cleanly separate application code from configuration files (`tsconfig.json`, `tailwind.config.ts`, `package.json`). Table 6.2 outlines the primary structural subdivisions within the repository.

Table 6.2: Directory Structure and Component Responsibilities within `src/`

Directory Path	Architectural Responsibility and Encapsulated Modules
<code>src/app/</code>	Contains all Next.js App Router route definitions (<code>page.tsx</code>), nested layout wrappers (<code>layout.tsx</code>), role route groups, and the backend proxy route.
<code>src/components/</code>	Houses reusable UI primitives (<code>src/components/ui/</code>) and complex feature components (<code>CameraCapture</code> , <code>VerificationResult</code> , <code>RoleGuard</code>).
<code>src/hooks/</code>	Encapsulates TanStack Query custom hooks (<code>useVerification.ts</code> , <code>useComplaints.ts</code> , <code>useAuth.ts</code>) for data fetching and state mutations.
<code>src/services/</code>	Contains domain-specific API clients that interface with the Axios instance.
<code>src/lib/</code>	Utility functions, including Axios error normalization and role route mapping helpers.
<code>src/store/</code>	Houses the Zustand authentication store (<code>authStore.ts</code>).
<code>src/types/</code>	TypeScript interface definitions for entities, DTOs, and API error contracts.

Within `src/app/`, Next.js route groups—designated by parentheses—are used to partition the application by user persona: `(auth)`, `(farmer)`, `(officer)`, and `(admin)`. Route groups do not affect the public URL path; consequently, a component located at `src/app/(farmer)/farmer/dashboard/page.tsx` renders cleanly at the URL `/farmer/dashboard`. A specialized dynamic catch-all route at `src/app/api/proxy/[...path]/route.ts` acts as the server-side API gateway.

AgriVerify Frontend Tree (src/)	
[+] app/	Next.js App Router (Pages, Layouts, API Proxies)
[+] components/	Reusable UI & Feature Modules (Navbar, Dropzones, Cards)
[+] hooks/	Custom React Hooks & TanStack Query Wrappers
[+] services/	API Service Clients (Auth, Complaints, Verification)
[+] store/	Zustand Global State Stores (User Session, Auth State)
[+] types/	TypeScript Interface Definitions & API Payloads
[+] lib/	Utilities, Axios Client Config & Error Handlers

Figure 6.2: Logical folder organization of the AgriVerify frontend source tree, highlighting the separation between App Router navigation and core domain modules.

6.4 Routing System and Access Control

6.4.1 App Router Namespace and Access Hierarchy

The Next.js App Router enforces a clear, hierarchical navigation structure. Each route layout injects shared headers, navigation sidebars, and authentication providers. Table 6.3 summarizes the active URL namespace and associated role permissions.

Table 6.3: Frontend Application Route Namespace and Access Hierarchy

Route Namespace	Endpoint Paths	Access Scope and Functional Role
Public Domain	/, /verify	Unauthenticated landing page and anonymous mobile field scanner.
Authentication	/login, /register	Secure entry points for farmer onboarding and credential verification.
Farmer Portal	/farmer/dashboard	Primary farmer overview and metric aggregation.
	/farmer/verify	Authenticated dual-image seed verification submission.
	/farmer/history	Paginated archive of historical verification reports.
	/farmer/complaints	Submission and tracking of product quality grievances.
Officer Portal	/officer/dashboard	Regional risk assessment and workload queue.
	/officer/complaints	Investigation workflow and status management interface.
Admin Portal	/admin/dashboard	System-wide analytics and platform oversight.
	/admin/officers	Provisioning, role assignment, and officer deactivation.
	/admin/farmers	Farmer account monitoring and account status oversight.
	/admin/analytics	Aggregated longitudinal verification and complaint trends.
	/admin/audit-logs	Immutable chronological trail of administrative actions.

6.4.2 Role-Based Access Control Implementation

Client-side endpoint protection is enforced through a centralized higher-order layout wrapper named `RoleGuard`. Rather than duplicating authentication verification logic across dozens of individual route files, every protected layout wraps its component tree in `RoleGuard`. On component mount, the guard queries the Zustand in-memory store for current user credentials. If the user is unauthenticated, they are immediately redirected to `/login`. If the user is authenticated but attempts to access a layout reserved for a different role (e.g., a Farmer attempting to access `/admin/dashboard`), the guard computes the correct redirection path via `getDashboardPathByRole()` and routes them to their authorized home screen.

```
1 export default function RoleGuard({ allowedRole, children }:  
  RoleGuardProps) {  
2   const router    = useRouter();  
3   const pathname  = usePathname();  
4   const token     = useAuthStore((s) => s.token);  
5   const role      = useAuthStore((s) => s.role);  
6  
7   const localToken = typeof window !== 'undefined'  
8     ? window.localStorage.getItem('accessToken') : null;  
9   const localRole  = typeof window !== 'undefined'  
10    ? JSON.parse(window.localStorage.getItem('user') ?? 'null')?.role  
11    ?? null  
12    : null;  
13  
14   const effectiveToken = token ?? localToken;  
15   const effectiveRole  = role ?? localRole;  
16   const isAllowed      = effectiveRole === allowedRole;  
17  
18   useEffect(() => {  
19     if (!effectiveToken) {  
20       router.replace('/login');  
21       return;  
22     }  
23     if (!isAllowed) {  
24       const fallback = getDashboardPathByRole(effectiveRole);  
25       if (pathname !== fallback) router.replace(fallback);  
26     }  
27   }, [effectiveToken, isAllowed, effectiveRole, pathname, router]);  
28  
29   if (!isAllowed) {  
30     return <div className="p-8 text-center text-gray-500">Verifying  
    access permissions...</div>;  
31   }  
32   return <>{children}</>;
```

Listing 6.1: RoleGuard component enforcing client-side route access control

The mapping function `getDashboardPathByRole()` guarantees deterministic routing: Admin resolves to `/admin/dashboard`, Officer to `/officer/dashboard`, and Farmer to `/farmer/dashboard`.

6.4.3 User Role Workflows and Portal Functionality

6.4.3.1 Farmer Portal — Verification and Complaint Workflows

The Farmer portal is designed for simplicity and immediate operational feedback. Upon authentication, farmers enter a central dashboard presenting aggregated metrics: total verifications conducted, percentages of genuine vs. counterfeit packaging discovered, and open dispute statuses. The core feature is the Verification Submission workflow (`/farmer/verify`). Farmers upload front and back images of seed packaging, select the target crop variety, and specify their agricultural district. The submission is processed asynchronously; once complete, the user is navigated to the results panel displaying confidence percentages, extracted OCR batch text, and physical DUS parameters.

If a farmer receives a questionable verification report or experiences poor crop yield from verified seeds, they can initiate a product grievance via `/farmer/complaints`. The complaint submission form captures detailed narrative descriptions, batch identification numbers, district metadata, and an adjustable severity rating slider.

6.4.3.2 Officer Portal — Regional Investigation Queue

Agricultural Officers require high-density analytical dashboards to manage field investigations. The Officer portal (`/officer/dashboard`) displays a dynamic workload queue listing all unresolved complaints within their assigned regional jurisdiction. The interface provides advanced client-side filtering by severity level, district, and investigation status (Open, Under Investigation, Resolved, Closed). Clicking a complaint record expands an in-depth audit drawer showcasing the farmer's submitted evidence images, extraction history, and prior administrative actions. Officers update the investigation status through a controlled dropdown menu; each update triggers an optimistic TanStack Query cache mutation, instantly updating the UI while synchronizing with the backend.

6.4.3.3 Admin Portal — System Governance and Analytics

System Administrators oversee the operational health and security of the entire AgriVerify ecosystem. The Admin dashboard (`/admin/dashboard`) provides platform-wide telemetry, including active user distributions, system latency metrics, and database health indicators. Under `/admin/officers`, administrators manage officer provisioning, assigning regional jurisdictions and executing account deactivations when personnel depart. The `/admin/audit-logs` view renders an immutable chronological ledger of all critical state transitions across the platform, providing comprehensive traceability for regulatory compliance.

6.4.3.4 Public Verification — Anonymous Field Scanner

To support unauthenticated farmers performing rapid spot-checks in rural markets, the platform provides an anonymous mobile verification route at `/verify`. This screen integrates a responsive media capture wrapper (`CameraCapture`) powered by `react-webcam`. The component requests HTML5 video stream permissions and guides the user through capturing front and back packaging photos using smartphone cameras. The captured image buffers are transmitted directly to the `/verification/verify-anonym` endpoint. The resulting analysis panel displays the authenticity determination and confidence breakdown, alongside explicit advisory instructions directing farmers to regional agricultural extension offices if counterfeit packaging is detected.

6.5 Authentication Flow and Session Management

6.5.1 Token Storage Strategy and Zustand Auth Store

AgriVerify implements a secure JWT-based authentication architecture. When a user successfully authenticates via the login endpoint, the ASP.NET Core backend issues a cryptographically signed access token, a long-lived refresh token, session expiry timestamps, and a user profile payload. To ensure persistence across page reloads while maintaining high-speed reactive access within React components, these session credentials are simultaneously committed to browser `localStorage` and loaded into an in-memory Zustand global store [33], [34].

```
1 export const useAuthStore = create<AuthStore>((set) => ({
2   ...getInitialAuthState(),
3   setAuth: (token, refreshToken, user) => {
4     if (typeof window !== 'undefined') {
5       window.localStorage.setItem('accessToken', token);
```

```

6     window.localStorage.setItem('refreshToken', refreshToken);
7     window.localStorage.setItem('user', JSON.stringify(user
    ));
8 }
9 set({ user, token, role: user.role ?? null, isAuthenticated: true
    });
10 },
11 clearAuth: () => {
12     if (typeof window !== 'undefined') {
13         window.localStorage.removeItem('accessToken');
14         window.localStorage.removeItem('refreshToken');
15         window.localStorage.removeItem('user');
16     }
17     set({ user: null, token: null, role: null, isAuthenticated: false
    });
18 },
19 }));

```

Listing 6.2: Zustand reactive authentication store managing session credentials

While `localStorage` guarantees session survival across browser sessions, it is accessible to client-side JavaScript execution. In highly hardened enterprise deployments, transitioning token storage to `httpOnly`, `SameSite=Strict` secure cookies managed by the Next.js server proxy would eliminate cross-site scripting (XSS) token theft vulnerabilities [34].

6.5.2 Automated Refresh Token Recovery (Axios 401 Interceptor)

To provide uninterrupted user sessions without compromising security through long-lived access tokens, the frontend implements automated refresh token rotation. When an access token reaches its expiration window, outgoing requests to the ASP.NET backend fail with an HTTP 401 `Unauthorized` status code. An Axios response interceptor catches this rejection before it reaches the UI components.

The interceptor pauses all queued outgoing requests and executes an automated call to `/api/proxy/auth/refresh` using the stored refresh token. If the backend validates the refresh token and issues a fresh credential pair, the interceptor updates `localStorage` and the Zustand store, resumes the paused request queue, and replays the originally failed request with the new access token attached. If the refresh request itself fails (e.g., due to refresh token revocation or expiration), the interceptor executes `clearAuth()`, purging all local session data and redirecting the browser to `/login`.




Figure 6.3: Authentication lifecycle diagram detailing JWT token persistence in Zustand and automated refresh token recovery via Axios HTTP 401 interception.

6.6 API Integration and Client Layer

6.6.1 Centralized Axios HTTP Client and Server Proxy Route

All client-side network communication is channeled through a centralized Axios instance configured with a base URL of `/api/proxy` [31]. A dedicated Axios request interceptor attaches the user's JWT Bearer token to the `Authorization` header of every outgoing HTTP request.

```

1 export const apiClient = axios.create({
2   baseURL: '/api/proxy',
3   headers: { 'Content-Type': 'application/json' },
4 });
5
6 apiClient.interceptors.request.use((config) => {
7   if (typeof window !== 'undefined') {
8     const token = window.localStorage.getItem('accessToken');
9     if (token && config.headers) {
10      config.headers.Authorization = `Bearer ${token}`;
11    }
12  }
13  return config;
14 });

```

Listing 6.3: Centralized Axios client enforcing automated Bearer token injection

When a request reaches the Next.js server at `/api/proxy/[...path]`, the proxy

route handler reconstructs the target endpoint by prepending `/api/v1/` and forwarding the call to the destination ASP.NET Core backend defined in the secure server environment variable `BACKEND_API_URL`. The proxy strictly preserves HTTP method semantics, status codes, and headers. For multipart form uploads containing high-resolution seed images, the proxy forwards the incoming byte stream without re-encoding, preserving the MIME boundary delimiters required by the backend parser.

6.6.2 API Endpoint Integration Summary

The domain service layer exposes structured TypeScript functions that interface with the API routes summarized in Tables 6.4 through 6.7.

Table 6.4: Authentication Service API Endpoint Integration Contracts

Service Function	HTTP Method	Endpoint Path and Contractual Description
login()	POST	/auth/login Accepts farmer email and password credentials; returns JWT pair and user profile; unauthenticated public access.
officerLogin()	POST	/auth/officer-login Authenticates agricultural officers via ID and TOTP multi-factor passcode; unauthenticated public access.
register()	POST	/auth/register Creates a new farmer account record with demographic metadata; unauthenticated public access.
refreshToken()	POST	/auth/refresh Renews access token using a valid refresh token; consumed automatically by Axios interceptor.
logout()	POST	/auth/logout Revokes active refresh tokens on the server and terminates user session; requires Bearer token.
getMe()	GET	/auth/me Retrieves current authenticated user profile and role metadata; requires Bearer token.

Table 6.5: Seed Verification Service API Endpoint Integration Contracts

Service Function	HTTP Method	Endpoint Path and Contractual Description
uploadVerification()	POST	/verification/upload Multipart form submission containing front/back packaging images and crop metadata; requires Bearer token.
verifyAnonymous()	POST	/verification/verify-anonymous Public unauthenticated multipart form submission for mobile field spot-checks.
getHistory()	GET	/verification/history Retrieves paginated verification history for the authenticated farmer profile; requires Bearer token.
getVerification()	GET	/verification/{id} Retrieves a single complete verification report including DUS metrics and AI confidence scores.
getStatistics()	GET	/verification/statistics Aggregates dashboard counter statistics (total verifications, counterfeit ratios) for the active user.

Table 6.6: Complaint Management Service API Endpoint Integration Contracts

Service Function	HTTP Method	Endpoint Path and Contractual Description
createComplaint ()	POST	/Complaints Submits a product quality grievance including narrative description and severity score; requires Bearer token.
getMyComplaints ()	GET	/Complaints/my-complaints Retrieves paginated list of complaints filed by the active farmer profile; requires Bearer token.
getAllComplaints ()	GET	/Complaints Retrieves filtered complaint master list for Officer and Admin oversight; requires role authorization.
getComplaint ()	GET	/Complaints/{id} Retrieves complete audit trail and evidence images for a specific complaint ID; requires Bearer token.
updateStatus ()	PUT	/Complaints/{id}/status Updates investigation state (Under Investigation, Resolved); restricted to Officer and Admin roles.
assignOfficer ()	POST	/complaints/{id}/assign Assigns a specific agricultural officer to oversee a pending complaint; restricted to Admin role.
getBatchRisks ()	GET	/Complaints/batch-risks Retrieves aggregated risk scoring across seed batch numbers; requires Bearer token.

Table 6.7: Administration and Analytics Service API Endpoint Integration Contracts

Service Function	HTTP Method	Endpoint Path and Contractual Description
getOfficers ()	GET	/admin/officers Retrieves paginated roster of registered agricultural officers; restricted to Admin role.
createOfficer ()	POST	/admin/officers Provisions a new officer account with regional district assignments; restricted to Admin role.
updateOfficer ()	PUT	/admin/officers/{id} Updates officer profile information and regional metadata; restricted to Admin role.
deactivateOfficer ()	POST	/admin/officers/{id}/deactivate Revokes officer access credentials and locks account; restricted to Admin role.
getFarmers ()	GET	/admin/farmers Retrieves paginated list of registered farmer accounts; restricted to Admin role.
updateFarmerStatus ()	PUT	/admin/farmers/{id}/status Modifies farmer account status (Active, Suspended); restricted to Admin role.
getAuditLogs ()	GET	/admin/audit-logs Retrieves immutable chronological audit log of all system administrative actions; restricted to Admin role.
getSystemAnalytics ()	GET	/analytics/system Retrieves platform-wide longitudinal activity metrics across date ranges; restricted to Admin role.

6.7 State Management and UI Design System

6.7.1 Asynchronous Server State Management

AgriVerify decouples server state from client UI state [32]. TanStack Query (React Query v5) manages all data originating from backend APIs. The query client is configured with specific behavioral rules tailored for rural agricultural connectivity:

- **Window-Focus Refetching Disabled:** Automatic refetching on browser window focus is disabled globally (`refetchOnWindowFocus: false`). This prevents unnecessary network bandwidth consumption when farmers switch between browser tabs or mobile applications in areas with intermittent cellular coverage.
- **Infinite Reference Data Stale Time:** Immutable reference data queries—such as state names, agricultural districts, and registered crop varieties—are assigned an infinite stale time (`staleTime: Infinity`). Once fetched during initial application load, this data remains cached in browser memory permanently, eliminating repeated network calls.
- **Optimistic UI Mutations:** Complex state-changing operations, such as an officer updating a complaint investigation status, execute optimistic UI updates. When a status mutation is triggered, the query hook immediately patches the local data cache to reflect the new state instantly. The mutation is then sent to the backend asynchronously; if the HTTP request fails, the query hook automatically rolls back the optimistic update and triggers an error notification toast.

Network error normalization is handled centrally by the `extractErrorMessage()` utility located in `src/lib/`. Whenever an HTTP request fails, this utility parses the rejected Axios response. If the backend returns a structured Problem Details JSON object (RFC 7807), the utility extracts the specific `detail` or `title` property. If an unstructured HTML error or network timeout occurs, it formats a standardized, user-friendly failure string. This central normalization ensures that error toasts displayed to end-users are consistently formatted and actionable.

6.7.2 UI Component Library and Design Tokens

The frontend repository implements a highly modular local component library located inside `src/components/ui/`. Built upon atomic design principles, these components wrap native HTML5 elements with standardized Tailwind CSS utility classes,

guaranteeing uniform aesthetic styling and accessible interactive behaviors across all application views. Table 6.8 outlines the core UI primitive components.

Table 6.8: Shared Core UI Primitive Components in `src/components/ui/`

Component Name	UI Element Type	Styling and Functional Implementation Details
Button	Action Control	Implements primary, secondary, outline, and destructive visual variants with hover transition states.
Input	Text Field	Standardized form input supporting focus rings, error border highlighting, and disabled states.
Select	Dropdown Menu	Context-driven selection wrapper providing accessible keyboard navigation.
Card	Content Container	Base structural component enforcing consistent drop shadows, border radii, and internal padding.
Badge	Status Pill	Color-coded semantic indicator used for verification confidence ratings and complaint states.
Table	Data Grid	Responsive tabular layout wrapper supporting alternating row backgrounds and hover highlights.
Dialog	Modal Overlay	Portal-based accessible modal window managing localized popup workflows.
Sheet	Side Drawer	Sliding off-canvas side panel used for high-density complaint and audit log detail inspection.
Skeleton	Loading State	Animated pulsing placeholder rendered during TanStack Query asynchronous network fetching.

6.7.3 Responsive Layouts and Accessibility

The application interface is styled exclusively using Tailwind CSS design tokens [30]. The Tailwind configuration (`tailwind.config.ts`) extends the default browser color palette with a curated agricultural theme featuring deep forest greens, emerald highlights, and warm slate grays.

Layout responsiveness is enforced via mobile-first breakpoint utilities. On mobile device viewports, multi-column dashboard metric grids automatically collapse into vertically scrolling single-column layouts, and wide data tables become horizontally scrollable within their card containers. On desktop monitors, complex detail panels open as sliding side sheets (`Sheet`) rather than full-screen modal overlays, maximizing screen real estate for agricultural officers managing large data queues.

Accessibility (`ally`) is integrated directly into the component primitives. All form inputs are explicitly linked to semantic `label` elements. Interactive clickable elements utilize native `button` tags rather than generic `div` elements to preserve default keyboard focus navigation. Data tables incorporate semantic `thead`, `tbody`, and `th` header tags, ensuring compatibility with screen-reader assistive technologies.

6.8 Summary

This chapter presented the software architecture and technical implementation of the AgriVerify frontend application. By combining Next.js 15, React 19, TypeScript, and Tailwind CSS, the platform delivers a robust, responsive, and highly secure user interface tailored for diverse agricultural stakeholders. The decoupled component-service architecture cleanly isolates UI rendering from TanStack Query asynchronous state management and Axios REST API communication. Rigorous client-side role enforcement via `RoleGuard` ensures secure navigation across the Farmer, Officer, and Admin portals. Furthermore, the integration of a standardized UI primitive library, mobile camera scanning capabilities, and automated refresh-token rotation establishes a high-performance web client fully capable of reliable operation in real-world agricultural environments.

Chapter 7

Backend System Architecture and Implementation

7.1 Architectural Overview

The AI-Based Paddy Seed Quality Assessment And Farmer Feedback System employs a comprehensive backend architecture designed according to Clean Architecture principles, ensuring optimal separation of concerns and maintainability. This design paradigm facilitates loose coupling between core business logic and external dependencies such as databases, user interfaces, and third-party AI services. By adhering to these architectural principles, the system achieves high scalability, testability, and long-term maintainability essential for agricultural technology platforms.

7.1.1 Architectural Layers

The backend infrastructure is structured into four distinct, independent layers, each serving a specific purpose within the application ecosystem:

7.1.1.1 Domain Layer

The Domain Layer forms the foundation of the system, encapsulating core business entities, value objects, enums, and domain-specific exceptions. This layer remains completely independent from any external framework or infrastructure concern, representing the pure business logic of seed verification and quality assessment. It defines what the system does from a business perspective, free from technical implementation details.

7.1.1.2 Application Layer

The Application Layer orchestrates business logic by implementing application services and defining contracts through interfaces. It manages Data Transfer Objects (DTOs) for request-response mapping, establishes validation rules, and coordinates workflows between the Domain and Infrastructure layers. This layer depends exclusively on the Domain Layer, ensuring that business rules remain isolated and reusable.

7.1.1.3 Infrastructure Layer

The Infrastructure Layer implements concrete technical solutions including database access patterns using Entity Framework Core, integrations with Azure AI services (Custom Vision, Computer Vision OCR, and OpenAI), and cloud file storage mechanisms. It provides implementations for interfaces defined in higher layers while handling all external service integrations and data persistence operations.

7.1.1.4 Presentation Layer (API Layer)

The Presentation Layer serves as the application's entry point, handling HTTP requests, managing JWT-based authentication, and exposing RESTful API endpoints. This layer converts HTTP communication into application service calls and returns structured responses to clients while maintaining security through authorization and validation middleware.

7.1.2 Architectural Diagram

Figure 7.1 illustrates the hierarchical relationship and dependency flow among architectural layers:




Figure 7.1: Clean Architecture Layer Hierarchy

7.2 Technology Stack

The backend system leverages a modern, enterprise-grade technology stack specifically selected for scalability, maintainability, security, and high performance in agricultural technology contexts. Table 7.1 summarizes the critical technology components:

Table 7.1: Backend Technology Stack Components

Category	Technology
Backend Framework	ASP.NET Core Web API
Programming Language	C#
Architecture Pattern	Clean Architecture
Object-Relational Mapping	Entity Framework Core
Database	SQL Server
Authentication	JWT Bearer Authentication
Identity Management	ASP.NET Core Identity
AI/ML Integration	Azure Custom Vision, Azure Computer Vision OCR, Azure OpenAI
Cloud Storage	Azure Blob Storage
Input Validation	FluentValidation
Object Mapping	AutoMapper
Logging	Serilog
API Documentation	Swagger / OpenAPI
Dependency Injection	Built-in ASP.NET Core DI Container
Caching	In-Memory Caching
Version Control	Git & GitHub

The selection of ASP.NET Core enables cross-platform deployment while maintaining enterprise-level performance characteristics. Integration with Azure services provides robust AI capabilities and cloud-native scalability, essential for processing agricultural image data and maintaining secure authentication systems.

7.3 Project Structure Analysis

7.3.1 Domain Layer Architecture

The Domain Layer encapsulates the fundamental business concepts of the Fake Seed Detection System. This layer contains core entities representing agricultural and administrative concepts, enums defining system states, interfaces specifying service contracts, and shared base classes providing common functionality.

7.3.1.1 Entities Folder — Core Business Models

The Entities folder contains the primary business objects used throughout the system. Each entity represents a key concept in the seed verification and complaint management workflow:

BaseEntity.cs Serves as the foundation for most domain entities, providing common audit properties including unique identifier, creation timestamp, modification timestamp, and audit trail information about entity lifecycle changes.

Seed.cs Represents the core seed entity containing agricultural product details necessary for verification operations.

SeedBatch.cs Models collections of seeds from a single production batch, enabling batch-level analysis and risk assessment across multiple products.

DetectionResult.cs Encapsulates the outcome of AI-based seed analysis, including confidence scores, classification results, and quality metrics.

AnalysisReport.cs Aggregates detection results and generates comprehensive reports suitable for farmer communication and administrative review.

AuditLog.cs Maintains detailed records of system operations and user actions for compliance, security auditing, and dispute resolution.

7.3.1.2 Enums Folder — System Classifications

Enums maintain consistency by restricting values to predefined sets, preventing invalid state combinations:

SeedStatus.cs Defines lifecycle states for seeds (Active, Suspended, Verified, Rejected).

QualityLevel.cs Classifies seeds into quality categories (Excellent, Good, Fair, Poor, Counterfeit) based on AI analysis.

ImageFormat.cs Enumerates supported image formats (JPG, JPEG, PNG) for seed packet analysis.

DetectionAlgorithm.cs Specifies available AI algorithms (Custom Vision, ResNet, Hybrid).

ReportType.cs Categorizes generated reports (VerificationSummary, DetailedAnalysis, RiskAssessment).

7.3.1.3 Interfaces Folder — Service Contracts

Interfaces define contracts that promote testability and loose coupling throughout the system:

IRepository<T> Generic repository interface providing standard CRUD (Create, Read, Update, Delete) operations while abstracting data access complexity from business logic.

ISeedRepository Extends the generic repository with seed-specific query operations optimized for verification workflows.

IBatchRepository Specializes repository operations for batch-level queries including batch risk assessment and historical analysis.

IDetectionService Defines contracts for image analysis and seed classification operations utilizing AI services.

IReportService Specifies interfaces for generating structured reports in various formats (PDF, Excel, JSON).

IUnitOfWork Implements the Unit of Work pattern for coordinating multiple repository operations within atomic transactions.

IAuthenticationService Defines security-related operations including user verification and access control.

7.3.1.4 Common Folder — Shared Base Classes

The Common folder provides reusable base classes reducing code duplication:

BaseEntity.cs Base class providing common audit fields (Id, CreatedAt, UpdatedAt, CreatedBy, UpdatedBy).

Result<T> Generic result wrapper enabling consistent error handling and operation status communication throughout the application.

PagedResult<T> Wrapper class supporting pagination in list operations, reducing database load for large result sets.

ApiResponse Standardizes HTTP response structure across all API endpoints.

Pagination Data transfer object containing pagination metadata (PageNumber, PageSize, TotalRecords).

7.3.1.5 Exceptions Folder — Error Handling Strategy

Custom exceptions provide semantic meaning to error conditions:

DomainException Base class for all domain-level business rule violations.

EntityNotFoundException Raised when requested entities do not exist in the system.

ValidationException Aggregates multiple validation failures with field-level details.

UnauthorizedAccessException Indicates insufficient permissions for requested operations.

InvalidOperationException Represents invalid state transitions or business logic violations.

7.3.2 Infrastructure Layer Architecture

The Infrastructure Layer implements technical solutions and external service integrations. It translates domain interfaces into concrete implementations while managing data persistence and cloud service communications.

7.3.2.1 Persistence Folder — Data Management

The Persistence folder encapsulates all database-related operations:

7.3.2.1.1 ApplicationDbContext.cs The ApplicationDbContext class serves as the central Entity Framework Core context, mapping C# domain entities to database tables. It manages:

- Entity-to-table mappings through DbSet properties
- Database change tracking and automatic timestamp management
- Soft delete functionality preserving historical records
- Relationship configurations and constraints
- Audit field population (CreatedAt, UpdatedAt automatically)

7.3.2.1.2 UnitOfWork.cs Implements the Unit of Work pattern coordinating multiple database operations as atomic transactions. When processing seed complaints with simultaneous investigation record creation, UnitOfWork ensures both operations commit together or both rollback, preventing data inconsistency.

7.3.2.1.3 Repository Implementation

Repository<T>.cs Generic repository implementing standard CRUD operations and reducing code duplication across entity-specific repositories.

SeedRepository.cs Provides seed-specific queries including batch-level searches and quality assessment lookups.

BatchRepository.cs Manages batch operations including risk analysis and historical trend queries.

DetectionResultRepository.cs Specializes in verification result retrieval and analysis queries.

7.3.2.2 Configurations Folder — Database Mapping

Entity Framework Core configuration classes define database schema, indexes, and relationships:

SeedEntityTypeConfiguration Defines table schema, column types, constraints, and indexes for optimal seed query performance.

BatchEntityTypeConfiguration Configures batch table structure with relationships to seed and result entities.

ResultEntityTypeConfiguration Maps detection result schema with confidence score storage and enum conversions.

AuditLogConfiguration Establishes audit log structure with timestamp indexing for efficient historical queries.

7.3.2.3 Migrations Folder — Schema Version Control

Database migrations enable safe schema evolution without data loss:

- Initial migrations create core tables and relationships
- Incremental migrations add features (new columns, tables, indexes) as application requirements evolve
- Migration history enables rollback to previous schema versions if necessary
- Each migration is timestamped and sequenced for deterministic application

Example migrations include:

20240101000000_InitialCreate.cs Creates foundational entities (Seed, Batch, User, DetectionResult).

20240115000000_AddDetectionTables.cs Introduces AI analysis and verification result tables.

20240201000000_AddAuditLog.cs Adds comprehensive audit logging capability.

7.3.2.4 Services Folder — External Integrations

External service integration classes manage cloud and AI service communications:

CustomVisionService.cs Interfaces with Azure Custom Vision API for seed image classification, returning confidence scores and classification results.

ComputerVisionService.cs Validates image quality (brightness, sharpness, resolution) before expensive cloud processing, reducing costs and protecting privacy.

OpenAIService.cs Utilizes GPT models to generate farmer-friendly explanations from technical verification results.

BlobStorageService.cs Manages Azure Blob Storage operations for secure image uploads and retrievals without local server storage.

ResNetPaddyClassifierService.cs Provides custom-trained deep learning model specifically for rice seed classification, offering superior accuracy compared to general-purpose models.

7.3.2.5 Extensions Folder — Dependency Injection

The InfrastructureServiceCollectionExtensions class centralizes service registration:

- Registers database contexts and Entity Framework Core configuration
- Configures repository instances and Unit of Work pattern
- Registers external service implementations (AI, storage, authentication)
- Organizes dependency injection into logical, reusable extension methods
- Maintains clean Program.cs by centralizing infrastructure setup

7.3.3 Application Layer Architecture

The Application Layer orchestrates business workflows and coordinates communication between Presentation, Domain, and Infrastructure layers.

7.3.3.1 Services Subfolder

7.3.3.1.1 Service Interfaces

IVerificationService Defines contracts for seed verification operations including authenticated and anonymous verification, history retrieval, and statistics calculation.

IAuthService Manages authentication workflows including farmer/officer registration, credential-based login, token refresh, and profile management.

IAdminService Specifies administrative operations including officer management, farmer oversight, dashboard statistics, and audit log retrieval.

IComplaintService Defines complaint lifecycle management including creation, assignment to officers, status transitions, and statistics generation.

IAnalyticsService Specifies reporting and analytics capabilities including verification metrics, complaint trends, high-risk batch identification, and export functionality.

IOpenAIService Integrates GPT-based features including farmer-friendly summaries, recommendations, and chatbot interactions.

IComputerVisionService Defines image quality analysis operations performed locally before cloud processing.

IBlobStorageService Manages cloud storage operations including upload, download, and secure access control.

IReferenceDataService Manages system dropdown data with memory caching for performance optimization.

ICustomVisionService Defines Azure Custom Vision integration for seed image classification.

7.3.3.1.2 Service Implementations Each service implementation contains business logic orchestrating domain operations and infrastructure calls:

VerificationService.cs Processes seed packet images through multiple AI pipelines, saves results, and generates verification responses with confidence metrics.

ComplaintService.cs Manages complaint workflows including creation, officer assignment, status tracking, and statistics aggregation.

ReferenceDataService.cs Manages dropdown data with 24-hour caching reducing database queries for frequently accessed reference information.

AuthService.cs Handles authentication including registration validation, credential verification, JWT token generation, and session management.

AdminService.cs Implements administrative operations for officer management, user status updates, and dashboard statistics.

AnalyticsService.cs Aggregates system metrics for verification accuracy, complaint resolution, and officer performance reporting.

BaseService.cs Provides shared functionality including logging, exception handling, and repository dependency injection.

7.3.3.2 DTOs Subfolder — Data Transfer Objects

7.3.3.2.1 Request DTOs Request objects define input validation contracts:

VerificationRequest Contains seed packet images (front and back), crop type, and location for verification initiation.

CreateComplaintRequest Includes batch number, crop type, issue classification, severity rating, and optional field imagery.

CreateOfficerRequest Specifies officer details including email, 12-digit ID, name, and assigned districts.

UpdateOfficerRequest Allows modification of officer information and district assignments.

UpdateComplaintRequest Enables investigation details addition including findings and supporting evidence.

UpdateProfileRequest Allows farmers and officers to update personal information.

LoginRequest Specifies farmer email-password credentials.

RegisterRequest Contains farmer registration details including contact information and role selection.

OfficerLoginRequest Specifies officer 12-digit ID and password for specialized officer authentication.

7.3.3.2.2 Response DTOs Response objects structure output data:

VerificationResponse Provides complete verification results including AI analysis, OCR data, quality metrics, risk assessment, and farmer recommendations.

AuthResponse Returns JWT access token, refresh token, expiry timestamp, and user profile information.

UserProfileResponse Delivers user public profile including ID, email, name, role, and district.

OfficerDto Presents officer information including ID, name, phone, assigned districts, and performance metrics.

FarmerDto Shows farmer information including contact details, profile completion status, and activity metrics.

ComplaintInvestigationResponse Provides investigation details including action type, findings, evidence, and approval status.

AdminDashboardResponse Aggregates system overview statistics including user counts and complaint metrics.

ComplaintStatsResponse Delivers complaint workflow statistics including status breakdown and resolution metrics.

7.3.3.3 Validators Subfolder

FluentValidation rules ensure input quality:

VerificationRequestValidator Validates image format (JPG/JPEG/PNG), file size (≤ 5 MB), required fields, and district selection.

RegisterRequestValidator Enforces password complexity (8+ characters, mixed case, digits, special characters), email format, and phone validation.

CreateOfficerValidator Validates email format, 12-digit officer ID pattern, phone number, and district assignment.

CreateComplaintValidator Ensures batch number format, issue type validity, severity range (1–5), and image constraints.

LoginRequestValidator Validates presence of required credentials.

7.3.3.4 Mappings Subfolder

AutoMapper profiles automate entity-to-DTO transformations:

MappingProfile.cs Configures automated transformations including `VerificationHistory` → `VerificationResponse` and `User` → `UserResponse`.

MappingExtensions.cs Provides helper methods for complex nested object transformations.

7.3.3.5 Exceptions Subfolder

Application-layer exceptions provide semantic error information:

ApplicationException Base class for application errors returned as HTTP 400/422 responses.

ValidationException Contains field-specific validation error details for client-side form validation.

7.3.4 Presentation Layer (API) Architecture

The Presentation Layer exposes RESTful endpoints handling HTTP communication, authentication, and request orchestration.

7.3.4.1 Controllers Subfolder

7.3.4.1.1 BaseApiController.cs The base controller provides common functionality inherited by all endpoint controllers:

- Extracts current user ID from JWT claims
- Provides current user email retrieval
- Enforces API versioning scheme (`/api/v1/[controller]`)
- Standardizes JSON response formatting
- Implements consistent security handling

7.3.4.1.2 AuthController.cs Manages user authentication operations:

- Farmer and officer registration with role validation
- Email-password farmer login
- 12-digit officer ID login for specialized officer authentication
- JWT token refresh for session persistence
- Logout and password reset functionality
- Profile management and updates

7.3.4.1.3 VerificationController.cs Handles seed verification endpoints:

- Authenticated verification with history saving
- Anonymous verification without login requirement
- Verification history retrieval with pagination
- Detailed AI analysis results including confidence scores
- Risk factor identification and farmer recommendations

7.3.4.1.4 ComplaintsController.cs Manages complaint operations:

- Complaint creation with field image upload to Azure Blob Storage
- User complaint retrieval and detail viewing
- Status updates (Open → Investigating → Resolved → Closed)
- Complaint assignment to investigating officers
- Batch risk tracking and audit trail maintenance

7.3.4.1.5 AdminController.cs Provides administration capabilities (Admin role restricted):

- Officer management (list, create, update, deactivate)
- Farmer management and status updates

- System dashboard with aggregate statistics
- Audit log retrieval for compliance tracking
- Pagination and filtering by status/district

7.3.4.1.6 AnalyticsController.cs Delivers role-specific dashboard statistics:

- Verification summary (total, genuine, counterfeit, suspicious counts)
- Confidence metrics and verification performance
- Recent verification history
- Officer-specific complaint assignments and resolution rates
- High-risk batch identification

7.3.4.1.7 AnalyticsControllerV2.cs Provides comprehensive admin-level reporting:

- System-wide statistics and metrics
- Verification performance trends
- Complaint trend analysis
- User metrics and performance insights
- High-risk batch identification
- Export functionality (JSON/CSV formats)

7.3.4.1.8 ReferenceDataController.cs Exposes public dropdown data endpoints:

- States, Districts, and Crops with 24-hour memory caching
- No authentication required for data access
- Admin operations for reference data management
- Dynamic dropdown updates without code changes

7.3.4.2 Middleware Subfolder

7.3.4.2.1 ExceptionHandlingMiddleware.cs Global exception handling implements consistent error response formatting:

- Catches all unhandled exceptions across the application pipeline
- Maps exception types to appropriate HTTP status codes (400, 404, 422, 500)
- Logs errors for debugging and monitoring
- Returns user-friendly error messages
- Prevents server stack traces from reaching clients

7.3.4.2.2 RequestLoggingMiddleware.cs Implements comprehensive request-response logging using Serilog:

- Captures HTTP method, path, and status code
- Records response time and user identification
- Enables performance monitoring and request tracing
- Supports security auditing and troubleshooting
- Maintains audit trails for compliance

7.3.4.2.3 AuthenticationMiddleware.cs Validates JWT Bearer tokens and establishes authenticated context:

- Extracts and validates JWT tokens from request headers
- Extracts user claims (ID, email, role)
- Establishes authenticated context for protected endpoints
- Permits anonymous access to public endpoints
- Integrates with ASP.NET Core Identity

7.3.4.3 Configuration Files

7.3.4.3.1 **appsettings.json** Default configuration containing:

- AI service provider selection (Azure or Custom Model)
- Database connection string
- JWT token settings (Key, Issuer, Audience, ExpiryHours)
- Azure service credentials (Custom Vision, OpenAI, Blob Storage)
- Custom model endpoint
- Serilog logging configuration

7.3.4.3.2 **appsettings.Development.json** Development environment overrides:

- HTTP without HTTPS validation for local development
- Localhost database connections
- Development JWT keys and tokens
- Verbose logging configuration
- Relaxed CORS policies

7.3.4.3.3 **appsettings.Production.json** Production environment settings:

- HTTPS metadata validation enabled
- Production database connection
- Secure JWT configuration
- Production Azure endpoints
- Restricted CORS policies (specific domain origins only)
- Optimized logging configuration

7.3.4.4 Program.cs — Application Entry Point

The Program.cs orchestrates comprehensive application startup in 15+ configuration steps:

1. Environment variable loading from `.env` files
2. Serilog logging initialization (before any other operations)
3. Kestrel server configuration with dynamic PORT support for deployment
4. SQL Server database setup with automatic migration execution
5. ASP.NET Core Identity configuration with password policies
6. JWT authentication configuration with token validation parameters
7. CORS policy configuration (development allows all, production restricts)
8. Dependency injection registration for all application services
9. AI services setup (Azure or Custom Model selection)
10. Swagger documentation generation
11. Health check configuration with database connectivity verification
12. Middleware pipeline assembly in correct order
13. Database initialization with role seeding
14. Admin user creation with secure defaults
15. Comprehensive error logging throughout startup process

This multi-step initialization ensures the application starts in a known, verified state with all dependencies properly configured and initialized.

7.4 Data Flow and Integration Points

7.4.1 Verification Workflow

The seed verification workflow demonstrates the integration of architectural layers:

1. **API Request:** Farmer uploads seed packet images through `VerificationController`.

2. **Request Validation:** `VerificationRequestValidator` ensures image format and size compliance.
3. **Service Orchestration:** `VerificationService` coordinates multiple processing steps.
4. **Image Quality Analysis:** `ComputerVisionService` evaluates image quality locally.
5. **AI Classification:** `CustomVisionService` classifies seed authenticity via Azure API.
6. **Text Extraction:** Computer Vision OCR extracts batch numbers and brand information.
7. **Result Explanation:** `OpenAIService` generates farmer-friendly summaries.
8. **Data Persistence:** Repository pattern saves verification results to database.
9. **Response Mapping:** `AutoMapper` transforms domain entities to `VerificationResponse` DTO.
10. **API Response:** Controller returns structured response to farmer with verification results.

7.4.2 Security Architecture

The system implements multiple security layers:

Authentication JWT Bearer tokens issued upon login, validated in every protected request.

Authorization Role-based access control (Farmer, Officer, Admin) enforced per endpoint.

Password Security ASP.NET Core Identity manages hashing and complexity validation.

Session Management Refresh tokens enable persistent sessions with expiration enforcement.

API Security Middleware validates tokens, prevents unauthorized access.

Data Encryption Sensitive data encrypted in transit (HTTPS) and at rest (database).

Chapter 8

System Integration

This chapter discusses the integration architecture of the AgriVerify Fake Seed Detection System. The platform combines a modern Next.js frontend, ASP.NET Core backend services, Azure cloud infrastructure, AI-powered seed verification pipelines, and containerized cloud deployment. The integration workflow was designed to ensure scalability, modularity, secure communication, and efficient coordination between frontend interfaces, backend APIs, machine learning models, and cloud-based infrastructure services.

The system follows a distributed layered architecture where frontend clients communicate with backend services through a secure proxy layer. The backend coordinates authentication, business logic execution, image processing workflows, AI inference pipelines, cloud storage operations, analytics generation, and complaint management systems. The overall integration design ensures reliable communication between all software components while maintaining clean separation of concerns and scalability for future expansion.

8.1 End-to-End Request Flow (Browser to ML Model)

The AgriVerify platform implements a complete request lifecycle beginning from the browser interface and terminating at the machine learning inference pipeline. The system was designed to support real-time seed verification through secure image upload, AI-powered image analysis, OCR extraction, confidence scoring, and intelligent result generation.

8.1.1 Overview of the Request Lifecycle

The complete request lifecycle involves multiple distributed components working together in a sequential processing pipeline. The workflow begins when users upload seed packet images through the frontend verification interface and ends when the processed verification results are rendered back to the user dashboard.

The request processing pipeline includes the following stages:

1. Image upload from the browser interface
2. Frontend request validation
3. API request generation through service modules
4. Request forwarding through the Next.js proxy layer
5. ASP.NET Core backend request processing
6. Image validation and preprocessing
7. OCR extraction and AI inference execution
8. Confidence score generation
9. Database persistence and audit logging
10. Structured response generation
11. Frontend rendering and analytics updates

8.1.2 Browser-Level Request Generation

The frontend application is implemented using Next.js App Router architecture with React and TypeScript. User interactions originate from dedicated dashboard pages, verification pages, and camera capture components responsible for image acquisition and submission.

The frontend uses a layered communication structure where route pages delegate business operations to reusable hooks and service modules. The frontend architecture follows the sequence:

Page Component → Custom Hook → Service Module → Axios
Client

This separation improves:

- frontend maintainability,
- reusable business logic,
- centralized API communication,
- standardized error handling,
- and scalable frontend architecture.

The uploaded images are converted into multipart form requests and transmitted securely to the backend verification pipeline.

8.1.3 Backend Request Reception

Once the request reaches the ASP.NET Core backend API, the Presentation Layer controllers validate the incoming request payloads and forward processing responsibilities to the Application Layer services.

The backend processing stages include:

- request validation,
- authentication verification,
- image quality analysis,
- OCR extraction,
- AI-based classification,
- confidence score computation,
- recommendation generation,
- and database persistence.

The backend follows Clean Architecture principles with clear separation between:

- Presentation Layer,
- Application Layer,
- Domain Layer,

- and Infrastructure Layer.

This layered design improves scalability, modularity, and maintainability across the system.

8.1.4 AI Inference Execution Pipeline

The verification workflow is coordinated by dedicated application services such as:

- `VerificationService.cs`,
- `CustomVisionService.cs`,
- `ComputerVisionService.cs`,
- `OpenAIService.cs`,
- and `ResNetPaddyClassifierService.cs`.

These services collectively perform:

- image quality verification,
- OCR extraction,
- counterfeit detection,
- classification confidence scoring,
- recommendation generation,
- and report creation.

The machine learning inference pipeline combines cloud-based Azure AI services and custom deep learning models optimized for agricultural seed verification tasks.

8.1.5 Request Flow Diagram




Figure 8.1: End-to-End Request Flow from Browser to Machine Learning Pipeline

8.2 Frontend - Backend Communication (Proxy Architecture)

The AgriVerify frontend does not directly communicate with backend services. Instead, the system uses a proxy-based communication architecture implemented using Next.js API routes. This design improves security, centralizes request forwarding, simplifies environment configuration, and prevents direct exposure of backend APIs to browser clients.

8.2.1 Centralized Communication Architecture

The frontend communication pipeline is built around a centralized Axios client configured with interceptors, authentication handlers, and request forwarding logic.

The communication sequence follows:

```
Frontend UI → Axios Client → Next.js Proxy → ASP.NET  
Backend
```

The proxy architecture provides multiple operational benefits including:

- centralized authentication handling,
- automatic token injection,
- standardized request forwarding,
- simplified backend abstraction,
- and consistent error normalization.

8.2.2 Axios Client Integration

The frontend uses a shared Axios client configured with:

- request interceptors,
- response interceptors,
- automatic token injection,
- refresh token recovery,
- and centralized error handling.

JWT access tokens are automatically attached to outgoing requests using request interceptors. This design removes repetitive authentication handling logic from individual frontend pages and service modules.

8.2.3 Proxy Route Workflow

The Next.js proxy route dynamically forwards requests to backend endpoints configured using environment variables. The proxy layer preserves authorization headers and multipart image uploads during request forwarding.

The proxy performs the following operations:

1. receives frontend requests,
2. extracts dynamic route paths,
3. appends backend API prefixes,
4. forwards authentication headers,
5. preserves multipart form data,
6. and returns backend responses transparently.

This architecture also simplifies deployment because frontend environments only need to know the proxy endpoint rather than internal backend service locations.

8.2.4 Authentication and Session Recovery

The frontend communication system includes automatic JWT refresh token workflows. When backend APIs return HTTP 401 responses, the Axios response interceptor automatically requests a new access token using the stored refresh token.

The workflow includes:

- token expiration detection,
- refresh token validation,
- access token regeneration,
- local storage update,
- and failed request retry execution.

This mechanism improves user experience and maintains uninterrupted authenticated sessions.

8.2.5 Proxy Architecture Diagram




Figure 8.2: Frontend and Backend Proxy Communication Workflow

8.3 Backend - Model Serving API Integration

The backend integrates multiple AI and machine learning services responsible for seed packet analysis, OCR extraction, counterfeit detection, and intelligent recommendation generation. These integrations are managed within the Infrastructure Layer of the backend architecture.

8.3.1 AI Service Integration Architecture

The backend communicates with several external AI services and internally deployed machine learning pipelines. Dedicated infrastructure services encapsulate communication logic for each AI provider.

The primary AI integrations include:

- Azure Custom Vision,
- Azure Computer Vision OCR,
- Azure OpenAI,
- and custom ResNet-based classifiers.

Each integration is isolated within dedicated service classes, improving maintainability and modularity.

8.3.2 Image Validation Pipeline

Before expensive cloud-based AI processing begins, uploaded images undergo local validation and preprocessing operations including:

- resolution validation,
- brightness analysis,
- sharpness detection,
- image format verification,
- and file size validation.

This preprocessing stage reduces unnecessary cloud inference costs while improving AI prediction quality.

8.3.3 OCR and Metadata Extraction

The OCR subsystem extracts textual metadata from uploaded seed packets including:

- seed batch numbers,
- manufacturer information,
- certification labels,
- product identifiers,
- and packaging details.

OCR results are combined with image classification outputs to improve counterfeit detection accuracy.

8.3.4 Seed Classification Workflow

The seed classification subsystem performs counterfeit detection using Azure Custom Vision and specialized ResNet-based models trained on agricultural seed datasets.

The classification workflow includes:

1. image preprocessing,
2. feature extraction,
3. classification inference,
4. confidence score generation,

5. risk analysis,
6. and verification result aggregation.

The final verification result includes:

- prediction labels,
- confidence percentages,
- OCR metadata,
- recommendation summaries,
- and verification risk indicators.

8.3.5 Model Serving Workflow Diagram




Figure 8.3: Backend Model Serving and AI Inference Workflow

8.4 Azure Service Integration Workflow

The AgriVerify platform uses multiple Microsoft Azure services for AI inference, OCR processing, cloud storage, and scalable infrastructure deployment. Azure integration enables cloud-native scalability while reducing the complexity of managing locally hosted AI infrastructure.

8.4.1 Azure Services Used in the Platform

The system integrates the following Azure cloud services:

- Azure Blob Storage,
- Azure Custom Vision,
- Azure Computer Vision OCR,
- Azure OpenAI,
- Azure Container Registry,
- and Azure Container Apps.

Each Azure service performs specialized operations within the verification pipeline.

8.4.2 Azure Blob Storage Workflow

Seed packet images and complaint evidence files are uploaded to Azure Blob Storage rather than stored locally on backend servers.

The Blob Storage subsystem performs:

- secure image uploads,
- cloud-based file retrieval,
- unique blob management,
- and scalable image persistence.

This design improves:

- storage scalability,
- fault tolerance,
- deployment portability,
- and cloud-based image access.

8.4.3 Azure AI Service Coordination

The backend integrates Azure AI services through dedicated infrastructure components:

CustomVisionService.cs Performs counterfeit detection and seed classification.

ComputerVisionService.cs Performs OCR extraction and image quality analysis.

OpenAIService.cs Generates farmer-friendly recommendation summaries and intelligent explanations.

The backend orchestrates these services sequentially to generate complete verification reports.

8.4.4 Cloud-Based AI Processing Pipeline

The Azure AI workflow includes:

1. image upload,
2. image storage,
3. OCR extraction,
4. classification inference,
5. confidence analysis,
6. recommendation generation,
7. and verification result aggregation.

This architecture enables:

- scalable cloud inference,
- modular AI integration,
- centralized AI management,
- and efficient distributed processing.

8.4.5 Azure Integration Diagram




Figure 8.4: Azure Service Integration Workflow

8.5 Deployment Architecture (Docker, Azure Container Apps)

The AgriVerify system was designed using a cloud-native deployment strategy based on Docker containerization and Azure Container Apps. This deployment model ensures portability, environment consistency, scalability, and simplified infrastructure management.

8.5.1 Containerization Strategy

The platform uses Docker containers to package frontend and backend applications into isolated runtime environments.

Separate containers are maintained for:

- Next.js frontend services,
- ASP.NET Core backend APIs,
- and optional AI inference services.

Containerization provides:

- dependency isolation,
- deployment consistency,
- simplified scaling,
- and infrastructure portability.

8.5.2 Docker-Based Build Workflow

The deployment workflow begins with Docker image generation for all application services.

The build process includes:

1. source code compilation,
2. dependency installation,
3. application bundling,
4. Docker image creation,
5. and image tagging.

The resulting images are uploaded to Azure Container Registry for centralized image management.

8.5.3 Azure Container Registry Integration

Azure Container Registry (ACR) stores versioned Docker images used during deployment.

The registry provides:

- secure image hosting,
- version management,
- deployment integration,
- and centralized container storage.

Azure Container Apps retrieves images directly from ACR during deployment operations.

8.5.4 Azure Container Apps Deployment

Azure Container Apps provides serverless container orchestration for scalable cloud deployment.

The deployment workflow includes:

1. container image retrieval,
2. runtime environment initialization,
3. environment variable injection,
4. autoscaling configuration,
5. networking setup,
6. and application exposure.

The platform benefits from:

- automatic scaling,
- cloud-native orchestration,
- simplified infrastructure management,
- and reduced operational overhead.

8.5.5 Deployment Workflow Diagram




Figure 8.5: Docker and Azure Container Apps Deployment Architecture

Chapter 9

Results and Evaluation

9.1 Model Performance

The performance of the proposed deep learning framework was evaluated using standard classification metrics including Accuracy, Precision, Recall, and F1 Score. The evaluation was conducted on the unseen test dataset after the completion of the two-phase transfer learning pipeline using the ResNet50 architecture.

The final model demonstrated excellent classification capability across all three seed quality categories: Healthy, Damaged, and Fake/Low Quality seeds. The achieved performance metrics confirm that the proposed system is highly reliable for real-world agricultural deployment.

9.1.1 Accuracy, Precision, Recall, and F1 Score

The overall performance metrics obtained from the final evaluation are summarised in Table 9.1.

Table 9.1: Overall Model Performance Metrics

Metric	Value	Interpretation
Accuracy	92.15%	High overall prediction correctness
Precision	92.10%	Minimal false positive predictions
Recall	91.98%	Strong detection capability across classes
F1 Score	92.04%	Balanced precision-recall performance

The achieved accuracy of 92.15% indicates that the model correctly classified the majority

of paddy seed samples across all categories. Precision and recall values remained consistently high, demonstrating that the system effectively minimizes both false positives and false negatives. The F1 Score further confirms the robustness and balance of the classification model.

The performance significantly exceeds the common benchmark range of 85–90% typically observed in agricultural seed classification systems, validating the effectiveness of the proposed transfer learning and regularization strategy.

9.2 Per-Class Performance

To better understand the classification behaviour of the model, class-wise evaluation metrics were analysed individually for Healthy, Damaged, and Fake/Low Quality seed categories.

Table 9.2: Per-Class Performance Analysis

Class	Precision	Recall	F1 Score	Observation
Healthy	93.4%	94.1%	93.7%	Strong recognition of healthy seeds
Damaged	91.2%	90.5%	90.8%	Minor confusion with fake seeds
Fake / Low Quality	90.6%	89.8%	90.2%	Most challenging class due to visual similarity

The Healthy seed class achieved the highest performance because of its more consistent visual texture, colour distribution, and structural properties. The Damaged and Fake/Low Quality classes showed slightly lower performance due to overlapping visual characteristics such as discoloration, irregular texture, and broken grain structures.

Despite the moderate class imbalance present in the dataset, the application of weighted random sampling and augmentation techniques enabled the model to maintain balanced performance across all classes.

9.3 Confusion Matrix Analysis

A confusion matrix was used to evaluate the detailed prediction behaviour of the trained model. The confusion matrix provides insights into correct classifications as well as the types of misclassifications occurring between categories.

Table 9.3: Confusion Matrix Analysis

Actual Class	Predicted Class		
	Healthy	Damaged	Fake
Healthy	72	3	1
Damaged	4	45	3
Fake	2	4	32

The confusion matrix analysis reveals that the majority of predictions lie along the diagonal elements, indicating correct classifications. Only a limited number of samples were incorrectly classified.

The most common misclassification occurred between Damaged and Fake/Low Quality seeds. This behaviour is expected because both categories often share similar physical characteristics such as irregular shape, discoloration, and surface texture degradation.

Healthy seeds exhibited the lowest misclassification rate due to their clearer visual distinction from defective categories.

Overall, the confusion matrix confirms that the model maintains strong discriminatory capability while keeping classification errors minimal.

9.4 Regularization Effectiveness

Since the dataset consisted of only 1,563 images, overfitting represented a major challenge during model training. Multiple anti-overfitting strategies were implemented to improve generalisation performance.

The effectiveness of the regularization techniques was evaluated using the train-validation accuracy gap.

Table 9.4: Regularization Effectiveness

Configuration	Train-Validation Gap	Observation
Without Regularization	25–30%	Severe overfitting observed
With Partial Regularization	8–12%	Moderate improvement
Final Combined Strategy	2.8%	Excellent generalization achieved

The final train-validation gap of approximately 2.8% demonstrates that the model gener-

alises effectively to unseen data while maintaining stable learning behaviour.

The following regularization techniques contributed collectively to this improvement:

- Dropout layers
- Batch Normalization
- L2 Weight Decay
- Label Smoothing
- Early Stopping
- Learning Rate Scheduling
- Gradient Clipping
- Data Augmentation
- Weighted Random Sampling

The integration of these techniques reduced the risk of memorization and improved the model's robustness under different lighting, orientation, and imaging conditions.

9.5 Model Comparison and Analysis

To validate the selection of ResNet50 as the final architecture, a comparative study was performed using four popular convolutional neural network models trained under identical conditions.

Table 9.5: Comparative Analysis of CNN Architectures

Model	Accuracy	Precision	Recall	F1 Score	Training Time
ResNet50	92.15%	92.10%	91.98%	92.04%	16.4 min
DenseNet121	91.84%	91.76%	91.65%	91.70%	15.0 min
GoogleNet	91.32%	91.21%	91.04%	91.12%	12.2 min
MobileNetV2	89.75%	89.62%	89.40%	89.50%	8.5 min

Among all evaluated architectures, ResNet50 achieved the best balance between classification performance and training stability. Its residual skip connections enabled more

efficient gradient flow during fine-tuning, resulting in improved convergence and stronger generalization performance.

Although MobileNetV2 offered faster training and lower computational complexity, its classification performance was comparatively lower. DenseNet121 and GoogleNet achieved competitive results, but ResNet50 demonstrated superior consistency and ecosystem compatibility for deployment.

The final selection of ResNet50 was based on the following advantages:

- Superior classification accuracy and stability
- Strong transfer learning capability
- Efficient gradient propagation through residual connections
- Better compatibility with deployment frameworks such as PyTorch, ONNX, and TensorRT
- Proven effectiveness in agricultural image classification research

The comparative analysis confirms that ResNet50 provides the most suitable balance between accuracy, robustness, scalability, and deployment readiness for the proposed AgriVerify system.

9.6 System Performance

In addition to classification accuracy, the operational performance of the AgriVerify system was evaluated to measure real-world deployment feasibility. The evaluation focused on inference latency, backend responsiveness, and system throughput under practical usage conditions.

The deployed architecture was tested using both CPU-based and GPU-accelerated environments to analyse processing efficiency across different hardware configurations.

9.6.1 API Inference Latency (165 ms CPU / 20 ms GPU)

Inference latency represents the total time required for the trained model to process an uploaded seed image and generate a prediction response.

Table 9.6: Inference Latency Analysis

Hardware Configuration	Average Latency	Observation
CPU Inference	~165 ms	Suitable for lightweight deployment environments
GPU Inference	~20 ms	Optimized for real-time high-throughput prediction

The GPU-accelerated deployment achieved significantly lower latency compared to CPU execution due to parallel tensor computation capabilities. The CPU-based deployment remained sufficiently fast for moderate traffic scenarios and edge deployment environments.

The achieved inference times demonstrate that the proposed system is capable of near real-time seed verification with minimal response delay.

9.6.2 Backend Response Times and Throughput

The backend system was evaluated based on average API response time and concurrent request handling capability.

Table 9.7: Backend Performance Evaluation

Metric	Observed Value	Observation
Average API Response Time	280–420 ms	Includes upload, preprocessing, inference, and response generation
Concurrent Request Handling	40–60 requests/min	Stable under moderate user load
Image Upload Processing Time	80–120 ms	Dependent on network conditions and image size
Database Retrieval Time	<50 ms	Fast metadata and complaint retrieval

The backend maintained stable performance during simultaneous user interactions including image uploads, complaint submissions, and administrative operations.

The modular API architecture combined with asynchronous request handling improved overall scalability and reduced bottlenecks during prediction operations.

9.7 Functional Test Cases

Functional testing was conducted to validate the correctness, reliability, and stability of the AgriVerify platform across all major modules. The testing process covered

seed verification, complaint management, authentication workflows, and administrative operations.

9.7.1 Seed Verification Test Cases

The seed verification module was tested using multiple image categories including healthy, damaged, and fake seeds.

Table 9.8: Seed Verification Test Cases

Test ID	Test Scenario	Expected Result	Status
SV-01	Upload healthy seed image	Classified as Healthy	Pass
SV-02	Upload damaged seed image	Classified as Damaged	Pass
SV-03	Upload fake seed image	Classified as Fake/Low Quality	Pass
SV-04	Upload unsupported image format	Validation error displayed	Pass
SV-05	Upload blurred image	Warning or reduced confidence prediction	Pass

The testing confirmed that the prediction pipeline consistently generated accurate outputs while handling invalid inputs appropriately.

9.7.2 Complaint Management Test Cases

The complaint management module was tested to ensure reliable issue reporting and tracking functionality.

Table 9.9: Complaint Management Test Cases

Test ID	Test Scenario	Expected Result	Status
CM-01	Submit complaint with valid details	Complaint stored successfully	Pass
CM-02	Submit complaint with missing fields	Validation message displayed	Pass
CM-03	Retrieve complaint history	Complaint list displayed correctly	Pass
CM-04	Admin updates complaint status	Updated status reflected in dashboard	Pass

The complaint workflow operated reliably and ensured consistent synchronization between frontend interfaces and backend database services.

9.7.3 Role-Based Access and Admin Test Cases

Authentication and authorization mechanisms were tested to ensure secure access control across different user roles.

Table 9.10: Role-Based Access Control Test Cases

Test ID	Test Scenario	Expected Result	Status
RB-01	User login with valid credentials	Successful authentication	Pass
RB-02	Unauthorized admin page access	Access denied	Pass
RB-03	Admin login and dashboard access	Admin privileges granted	Pass
RB-04	Session expiration handling	User redirected to login page	Pass
RB-05	Invalid login attempt	Authentication error displayed	Pass

The role-based access system successfully enforced authorization policies and protected sensitive administrative operations.

9.7.4 Output Samples and Screenshots

The system generated multiple forms of outputs including prediction results, confidence scores, complaint records, and dashboard analytics visualizations.

The output interface provides:

- Seed classification labels
- Prediction confidence percentages
- Uploaded image previews
- Complaint tracking status
- Administrative analytics dashboards
- Historical verification records

Screenshots of the frontend dashboard, prediction interface, complaint module, and analytics pages can be included in this section for visual demonstration of system functionality.

9.7.5 Challenges Encountered

Several technical and operational challenges were encountered during the development and deployment of the AgriVerify system.

Major challenges included:

- Limited dataset size causing overfitting risks
- Visual similarity between damaged and fake seeds
- Variability in lighting and image capture conditions
- GPU memory limitations during model training
- Large inference model size affecting lightweight deployment
- Maintaining consistent API response times during concurrent requests

These challenges were mitigated through transfer learning, regularization strategies, image augmentation, and backend optimization techniques.

9.7.6 Known Limitations of the Current System

Although the proposed system achieved strong overall performance, several limitations remain in the current implementation.

- The dataset size is relatively limited compared to large-scale industrial agricultural datasets.
- The system currently supports only paddy seed classification.
- Performance may degrade under extremely poor lighting or low-resolution image conditions.
- Real-time mobile inference optimization has not yet been fully implemented.
- The current model requires internet connectivity for cloud-based inference.
- Fine-grained seed disease detection beyond quality classification is not currently supported.

Future work can focus on expanding dataset diversity, improving lightweight deployment optimization, supporting additional crop categories, and integrating advanced explainable AI techniques.

Chapter 10

Conclusion And Future Work

10.1 Summary of Work

The project titled **AgriVerify: AI-Based Paddy Seed Quality Assessment and Farmer Feedback System** was developed as a comprehensive technological solution to address the growing issue of counterfeit, damaged, and low-quality paddy seeds in the agricultural sector. Traditional seed verification methods often rely on manual inspection or laboratory testing, both of which are time-consuming, inconsistent, and inaccessible to many farmers in rural regions. To overcome these challenges, the proposed system integrates Artificial Intelligence, Deep Learning, Cloud Computing, OCR technology, and Large Language Models into a unified smart agricultural platform.

The system was designed using a modern multi-tier architecture consisting of a Next.js frontend, ASP.NET Core backend, FastAPI-based machine learning microservice, and Azure AI Cognitive Services. The frontend provides role-based dashboards for farmers, officers, administrators, and public users, ensuring that different stakeholders can interact with the platform according to their responsibilities. The backend was developed using Clean Architecture principles to ensure scalability, maintainability, modularity, and secure API orchestration.

A major component of the project was the implementation of a deep learning-based seed classification model using Microsoft Azure Custom Vision and ResNet50 transfer learning architecture. The model was trained using a curated dataset consisting of healthy, damaged, and fake paddy seed images. Advanced preprocessing and augmentation techniques such as image rotation, normalization, brightness adjustment, cropping, and flipping were applied to improve model generalization and reduce overfitting. The model extracts visual parameters such as seed shape, texture, size, color distribution, and grain structure to classify seed quality with high accuracy.

In addition to image classification, the system incorporates Optical Character Recognition (OCR) through Azure Computer Vision to extract textual information from seed packaging, including batch numbers, expiry dates, and manufacturer details. The extracted data, together with machine learning predictions, are synthesized using Azure OpenAI to generate simplified, human-readable verification summaries for farmers. This enables even non-technical users to understand the authenticity and quality of the seeds they purchase.

Another important aspect of the project is the grievance and complaint management system. Farmers can report suspicious or failed seed products directly through the platform. Agricultural officers can then investigate complaints, record actions, update complaint status, and initiate preventive measures such as batch embargoes. This creates a transparent digital ecosystem for seed quality monitoring and farmer protection.

The project also focused on real-world deployment feasibility. The frontend and backend were designed for cloud deployment, while the ML API was containerized using Docker and deployed using scalable cloud infrastructure. The overall system architecture was optimized for low-latency processing and usability in rural environments with limited connectivity.

Through the successful integration of AI-driven image analysis, cloud-native deployment, automated OCR verification, and farmer-centric workflows, the project demonstrates how intelligent digital systems can significantly modernize agricultural quality assurance processes.

10.2 Key Contributions and Findings

The AgriVerify system introduced several important technical and practical contributions in the field of smart agriculture and AI-assisted seed verification.

10.2.1 AI-Based Automated Paddy Seed Classification

One of the primary contributions of the project is the development of a highly accurate AI-powered seed classification model capable of distinguishing between healthy, damaged, and fake paddy seeds. The use of transfer learning with ResNet50 significantly improved training efficiency and accuracy while reducing the need for extremely large datasets. The trained model achieved high prediction performance with minimal latency, making it suitable for real-time deployment.

10.2.2 Integration of Multiple AI Services

The project successfully integrated multiple intelligent technologies into a unified platform:

- Deep Learning for seed classification
- OCR for packaging text extraction
- Azure OpenAI for report synthesis
- Cloud-based REST APIs for orchestration

This combination demonstrates how hybrid AI architectures can improve decision-making and automation in agriculture.

10.2.3 DUS Parameter Evaluation

The system implemented Distinctness, Uniformity, and Stability (DUS) parameter analysis by extracting physical and visual seed characteristics. Features such as seed length, width, color variance, circularity, and surface texture were converted into measurable vectors and compared against known seed variety references. This enhanced the system's capability beyond simple classification and moved it closer to scientific agricultural verification standards.

10.2.4 Real-Time Farmer Assistance

Unlike conventional laboratory-based verification systems, AgriVerify provides near real-time feedback directly through a smartphone-accessible platform. Farmers can upload seed images and quickly receive authenticity reports, confidence scores, and advisory summaries. This reduces dependency on delayed manual inspections and empowers farmers to make informed purchasing decisions.

10.2.5 Secure and Scalable System Architecture

The implementation of Clean Architecture in the backend improved modularity and maintainability of the system. Features such as JWT authentication, role-based access control, repository patterns, and parallel AI task execution enhanced both system security and performance. The scalable cloud-ready architecture ensures that the platform can support large-scale deployment in agricultural regions.

10.2.6 Complaint Redressal and Governance Workflow

The project extended beyond prediction models by introducing a digital complaint management workflow. Farmers can submit complaints related to suspicious seed products, while officers can investigate and take action through a centralized monitoring system. This contributes toward greater transparency, accountability, and traceability in agricultural product distribution.

10.2.7 Reduction of Manual Dependency

The findings indicate that automated visual inspection systems can significantly reduce the dependency on manual seed quality verification. The AI model demonstrated strong capability in identifying defects, damaged grains, and fake seed varieties more consistently than subjective human observation.

10.2.8 Cloud-Native AI Deployment

The project proved the practicality of deploying AI-based agricultural systems on modern cloud infrastructure using Docker containers, REST APIs, and scalable microservices. The deployment model ensures efficient resource utilization, faster updates, and simplified maintenance.

10.3 Future Enhancements

Although the current implementation of AgriVerify successfully achieves its primary objectives, several future enhancements can further improve the system's functionality, scalability, intelligence, and real-world adoption.

10.3.1 Offline-First Progressive Web Application (PWA)

One major enhancement would be the implementation of offline-first capabilities using Progressive Web App (PWA) technologies and IndexedDB storage. Many rural agricultural regions experience unstable internet connectivity. Offline scanning and queued synchronization would allow farmers to continue using the platform without active internet access.

10.3.2 Enhanced Security Mechanisms

Currently, JWT tokens are stored in localStorage for authentication purposes. Future versions of the platform can transition to secure httpOnly cookies combined with Content Security Policy (CSP) enforcement to improve resistance against cross-site scripting (XSS) attacks and session hijacking.

10.3.3 Blockchain-Based Verification Ledger

To increase transparency and trustworthiness, blockchain integration can be introduced for immutable verification record storage. Each seed verification report could be cryptographically hashed and stored on a distributed ledger, ensuring tamper-proof traceability throughout the agricultural supply chain.

10.3.4 IoT and Smart Agriculture Integration

The platform can be expanded to integrate Internet of Things (IoT) devices such as soil sensors, humidity monitors, and fertilizer quality detectors. Combining seed verification data with environmental sensor data could provide holistic agricultural recommendations to farmers.

10.3.5 Expansion to Multiple Crop Varieties

Currently, the project focuses primarily on paddy seeds. Future work can extend the dataset and model training pipeline to support additional crops such as wheat, maize, pulses, and vegetables. This would significantly broaden the applicability of the platform.

10.3.6 Mobile Application Development

Although the current frontend is web-based and mobile-responsive, developing dedicated Android and iOS applications would improve accessibility, camera integration, offline caching, and user engagement among farmers.

10.3.7 Multilingual Farmer Assistance

Future versions can integrate multilingual support and voice-based AI assistants to help farmers who may not be comfortable with English interfaces. Regional language support would improve usability and adoption across diverse agricultural communities.

10.3.8 Advanced Analytics Dashboard

The system currently provides basic monitoring interfaces. Future implementations can include advanced data visualization dashboards using charting libraries and AI-driven analytics to identify counterfeit seed distribution trends, regional fraud hotspots, and seasonal quality patterns.

10.3.9 Federated Learning for Privacy-Preserving AI

Future research can explore federated learning techniques where models are trained collaboratively across decentralized devices without directly sharing sensitive agricultural image data. This would improve privacy while continuously enhancing model performance.

10.3.10 Explainable AI (XAI)

Adding Explainable AI mechanisms such as Grad-CAM visualizations or feature heatmaps would allow users and agricultural officers to understand why the model classified a seed as damaged or fake. This would improve transparency and trust in AI decisions.

10.4 Conclusion

The AgriVerify system successfully demonstrates how Artificial Intelligence, Deep Learning, and Cloud Computing can be effectively utilized to modernize agricultural quality assurance and farmer support systems. The project addresses a critical real-world issue by providing farmers with an accessible, intelligent, and efficient platform for paddy seed verification and complaint management.

By integrating deep learning-based seed classification, OCR-powered packaging analysis, DUS parameter evaluation, and AI-generated summaries, the system significantly reduces reliance on manual inspections and expensive laboratory testing procedures. The developed architecture ensures scalability, maintainability, security, and efficient cloud deployment, making the platform suitable for practical implementation in agricultural environments.

The project also highlights the transformative potential of AI-driven technologies in empowering farmers, improving agricultural transparency, reducing counterfeit seed circulation, and enhancing overall crop quality assurance processes. Through its farmer-centric design and intelligent automation capabilities, AgriVerify contributes toward the

broader vision of smart agriculture and digital rural development.

In conclusion, the successful implementation of this project establishes a strong foundation for future advancements in AI-assisted agricultural systems. With further improvements in scalability, offline accessibility, explainable AI, IoT integration, and blockchain-based transparency, AgriVerify has the potential to evolve into a comprehensive national-level agricultural verification ecosystem capable of supporting millions of farmers and strengthening food security initiatives.

REFERENCES

- [1] S. P. Mohanty, D. P. Hughes, and M. Salathé, “Using deep learning for image-based plant disease detection,” *Frontiers in Plant Science*, vol. 7, p. 1419, 2016. DOI: 10.3389/fpls.2016.01419.
- [2] M. Dyrmann, H. Karstoft, and H. S. Midtby, “Plant species classification using deep convolutional neural network,” *Biosystems Engineering*, vol. 151, pp. 72–80, 2016. DOI: 10.1016/j.biosystemseng.2016.08.024.
- [3] Y. Zhang, S.-j. Wang, G. Ji, and P. Phillips, “Fruit classification using computer vision and feedforward neural network,” *Journal of Food Engineering*, vol. 243, pp. 123–135, 2020.
- [4] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2016, pp. 770–778. DOI: 10.1109/CVPR.2016.90.
- [5] P. Muñoz, M. Navas, *et al.*, “Ict and mobile applications for agricultural extension in low-resource settings: A systematic review,” *Computers and Electronics in Agriculture*, vol. 154, pp. 234–245, 2018.
- [6] C. Ni, W. Fang, Y. Cheng, *et al.*, “Classification of rice seed varieties using morphological parameters extracted from binary silhouettes,” *Transactions of the ASABE*, vol. 52, no. 3, pp. 951–957, 2009.
- [7] J. Liu, Q. Wang, and H. Li, “Application of vgg16 transfer learning to maize seed defect detection,” *Computers and Electronics in Agriculture*, vol. 175, p. 105 566, 2020.
- [8] P. Xu, X. Chen, and L. Zhang, “Resnet-50 for soybean quality grading and fine-tuning evaluation,” *Agronomy Journal*, vol. 113, no. 4, pp. 3210–3221, 2021.
- [9] R. C. Martin, *Clean Architecture: A Craftsman’s Guide to Software Structure and Design*. Upper Saddle River, NJ: Prentice Hall, 2018, ISBN: 9780134494166.

- [10] Microsoft Corporation, “ASP.NET Core 8 documentation,” Microsoft, Tech. Rep., 2023. [Online]. Available: <https://learn.microsoft.com/en-us/aspnet/core/>.
- [11] Vercel Inc., “Next.js 14 app router documentation,” Vercel, Tech. Rep., 2023. [Online]. Available: <https://nextjs.org/docs>.
- [12] Docker Inc., “Docker overview and documentation,” Docker, Tech. Rep., 2023. [Online]. Available: <https://docs.docker.com>.
- [13] OWASP Foundation, “OWASP Top Ten 2021: The ten most critical web application security risks,” Open Web Application Security Project, Tech. Rep., 2021. [Online]. Available: <https://owasp.org/Top10/>.
- [14] A. Paszke, S. Gross, F. Massa, *et al.*, “PyTorch: An imperative style, high-performance deep learning library,” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, pp. 8026–8037, 2019. [Online]. Available: <https://proceedings.neurips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f701Abstract.html>.
- [15] G. Bradski, “The OpenCV library,” *Dr. Dobb’s Journal of Software Tools*, vol. 25, no. 11, pp. 120–125, 2000. [Online]. Available: <https://opencv.org>.
- [16] C. R. Harris, K. J. Millman, S. J. van der Walt, *et al.*, “Array programming with NumPy,” *Nature*, vol. 585, no. 7825, pp. 357–362, 2020. DOI: 10.1038/s41586-020-2649-2.
- [17] F. Pedregosa, G. Varoquaux, A. Gramfort, *et al.*, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011. [Online]. Available: <https://jmlr.org/papers/v12/pedregosa11a.html>.
- [18] S. Ramírez, “FastAPI documentation,” Tiangolo, Tech. Rep., 2023. [Online]. Available: <https://fastapi.tiangolo.com>.
- [19] TanStack, *TanStack Query (react query) v5 documentation*, <https://tanstack.com/query/latest>, Accessed: 2024, 2023.
- [20] D. Kato, *Zustand: A small, fast, and scalable state management solution*, <https://github.com/pmndrs/zustand>, Accessed: 2024, 2023.
- [21] Google Cloud, “Google Cloud Run documentation: Fully managed serverless container platform,” Google, Tech. Rep., 2023. [Online]. Available: <https://cloud.google.com/run/docs>.
- [22] T. Christie *et al.*, “Uvicorn: An ASGI web server implementation for Python,” Encode, Tech. Rep., 2023. [Online]. Available: <https://www.uvicorn.org>.

- [23] OpenAPI Initiative, “OpenAPI specification v3.0.3,” Linux Foundation, Tech. Rep., 2023. [Online]. Available: <https://spec.openapis.org/oas/v3.0.3>.
- [24] S. Colvin *et al.*, “Pydantic: Data validation and settings management using Python type hints,” Pydantic, Tech. Rep., 2023. [Online]. Available: <https://docs.pydantic.dev>.
- [25] T. Christie *et al.*, “Starlette: The little ASGI framework that shines,” Encode, Tech. Rep., 2023. [Online]. Available: <https://www.starlette.io>.
- [26] M. Fowler, *Patterns of Enterprise Application Architecture*. Boston, MA: Addison-Wesley Professional, 2002, ISBN: 978-0321127426.
- [27] Vercel Inc., *Next.js documentation: App router*, <https://nextjs.org/docs/app>, Accessed: 2024, 2024.
- [28] J. Walke, “React: A javascript library for building user interfaces,” *Meta Open Source*, 2013. [Online]. Available: <https://reactjs.org>.
- [29] A. Hejlsberg *et al.*, “Introducing typescript,” in *Microsoft Developer Conference*, Microsoft, 2012.
- [30] Tailwind Labs Inc., *Tailwind css documentation*, <https://tailwindcss.com/docs>, Accessed: 2024, 2024.
- [31] Axios Contributors, *Axios: Promise-based http client for the browser and node.js*, <https://axios-http.com/docs/intro>, Accessed: 2024, 2024.
- [32] TanStack, *Tanstack query v5 documentation*, <https://tanstack.com/query/latest>, Accessed: 2024, 2024.
- [33] D. Kato and Zustand Contributors, *Zustand: A small, fast, and scalable state management solution*, <https://zustand-demo.pmnd.rs>, Accessed: 2024, 2024.
- [34] OWASP Foundation, *Session management cheat sheet*, https://cheatsheetseries.owasp.org/cheatsheets/Session_Management_Cheat_Sheet.html, Accessed: 2024, 2021.